

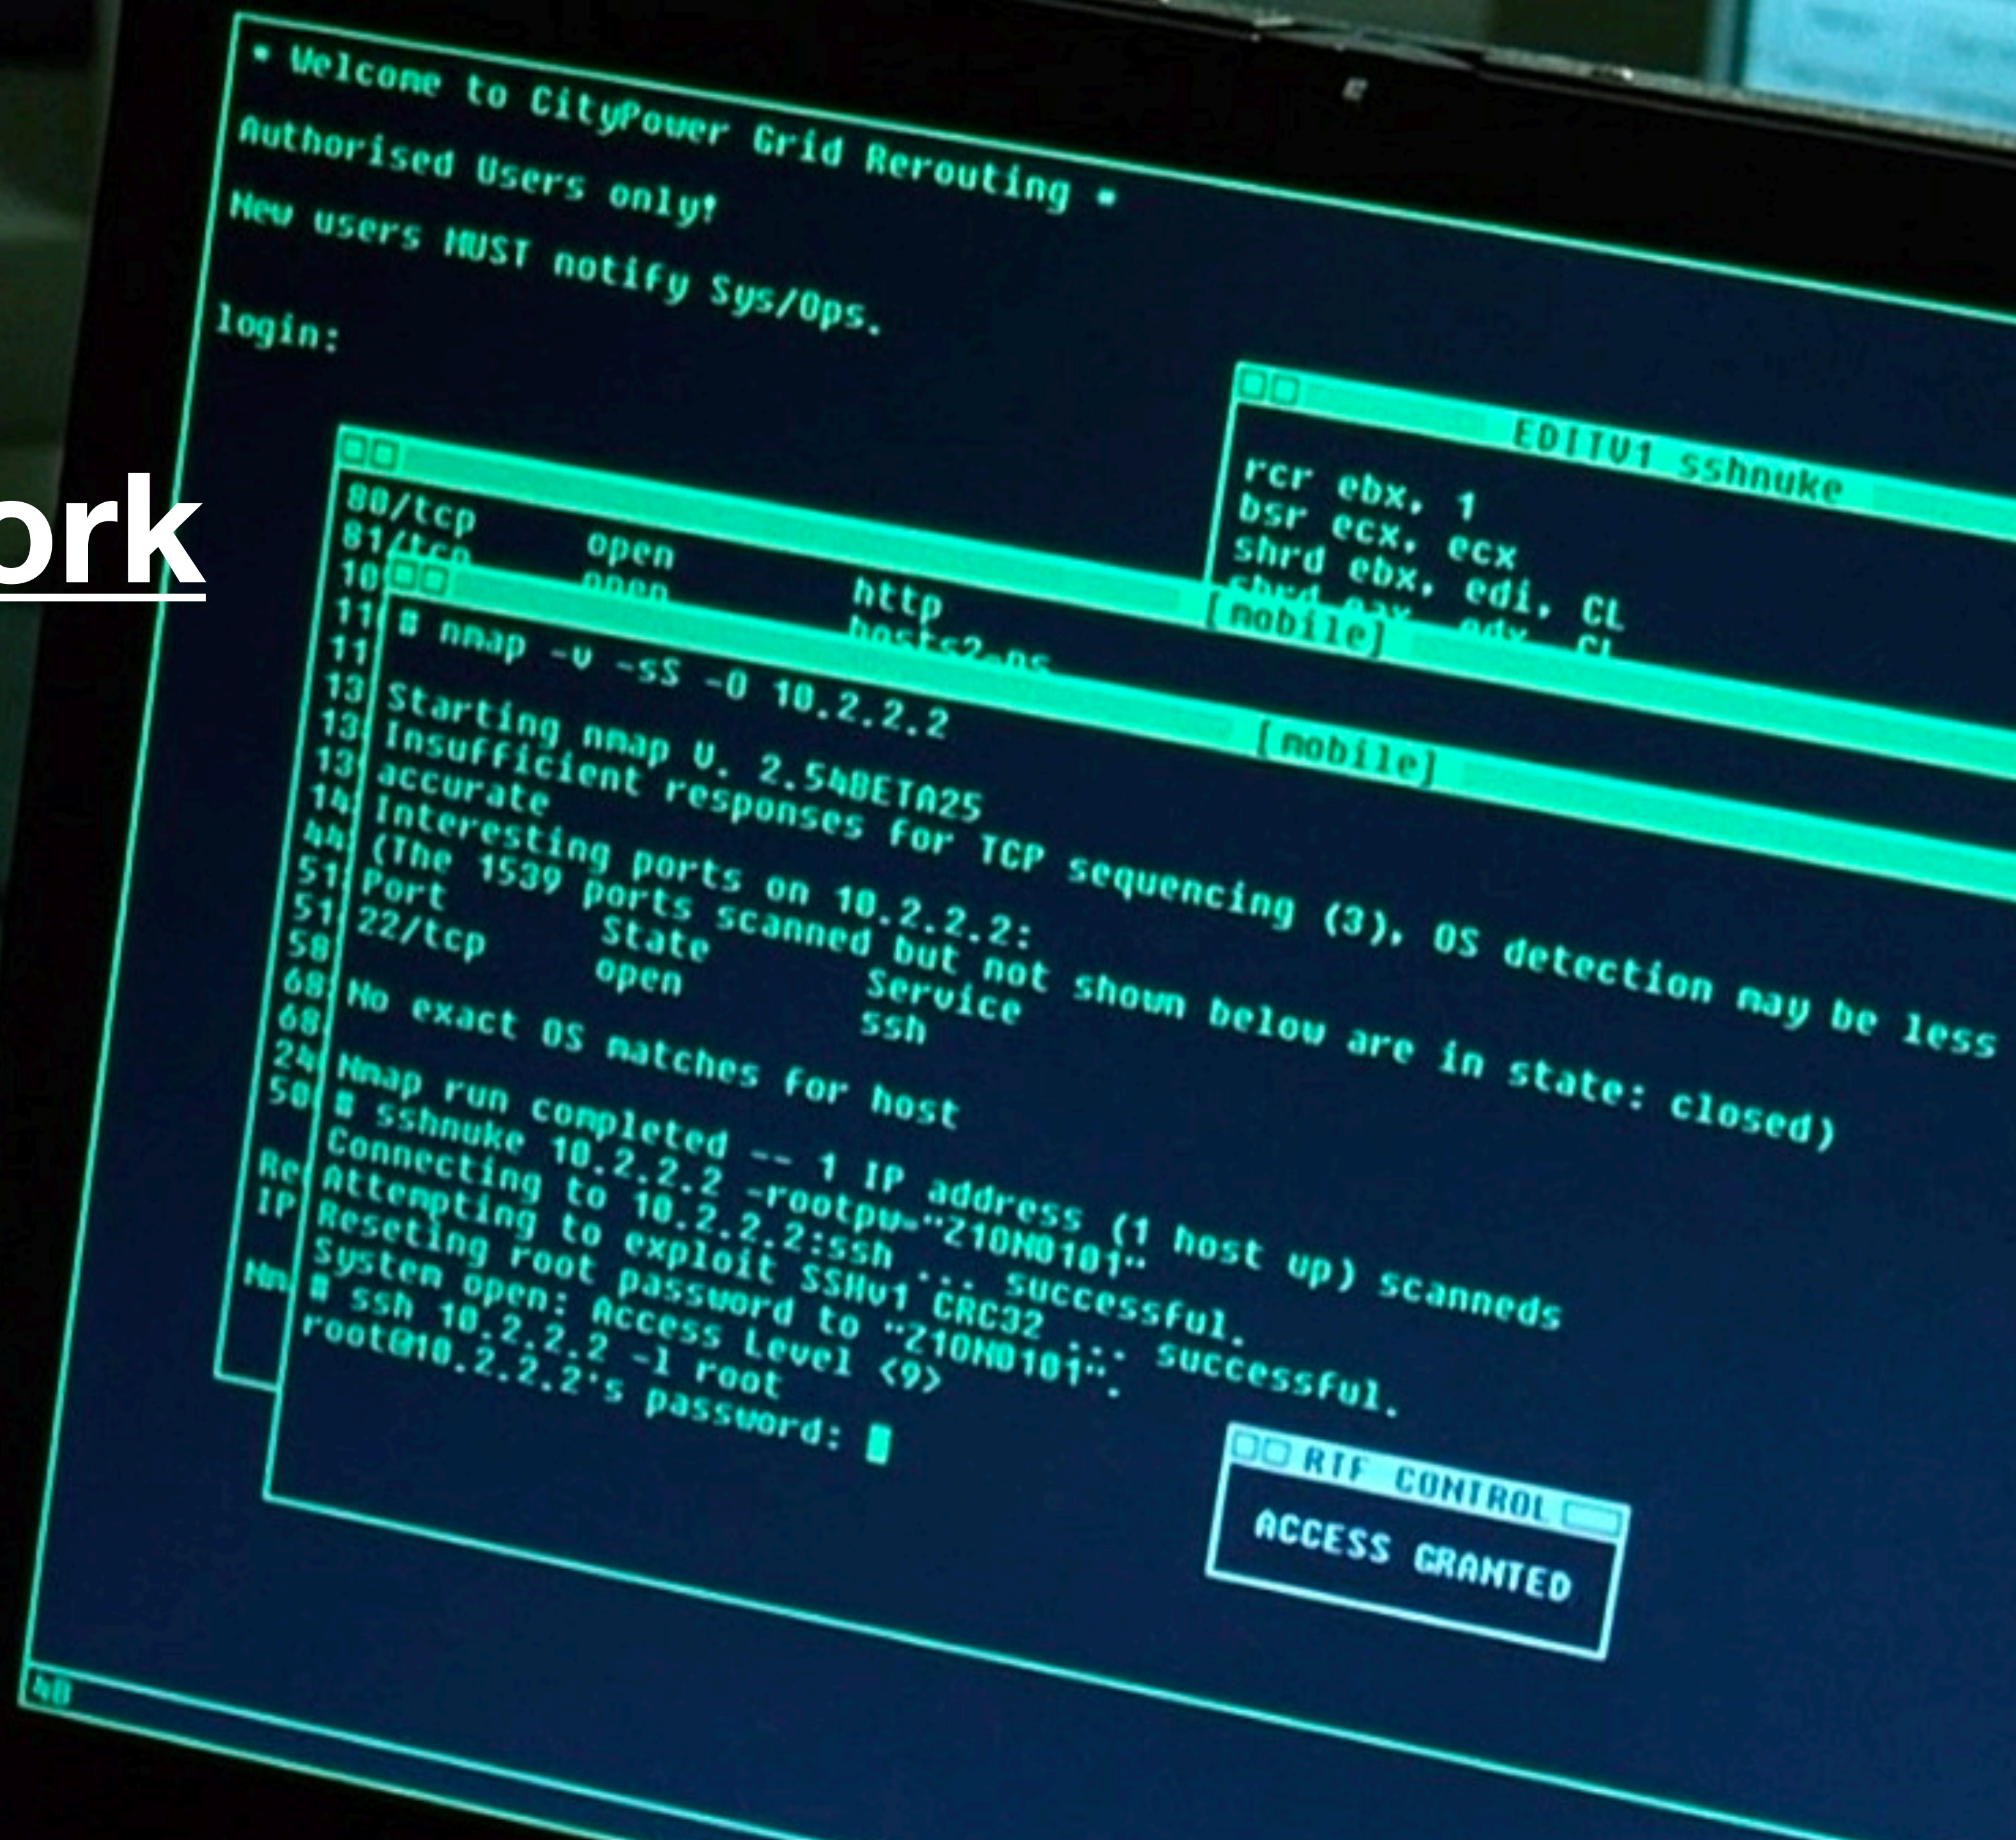
Computer Network

Security

ECE 4112/6612

CS 4262/6262

Prof. Frank Li




Logistics

Quiz 1 scores released. Regrade requests open until next **Tuesday, Oct 3, midnight.**

Project proposals due this **Friday, midnight.**

Continuing with DNS

Distributed, Hierarchical Database

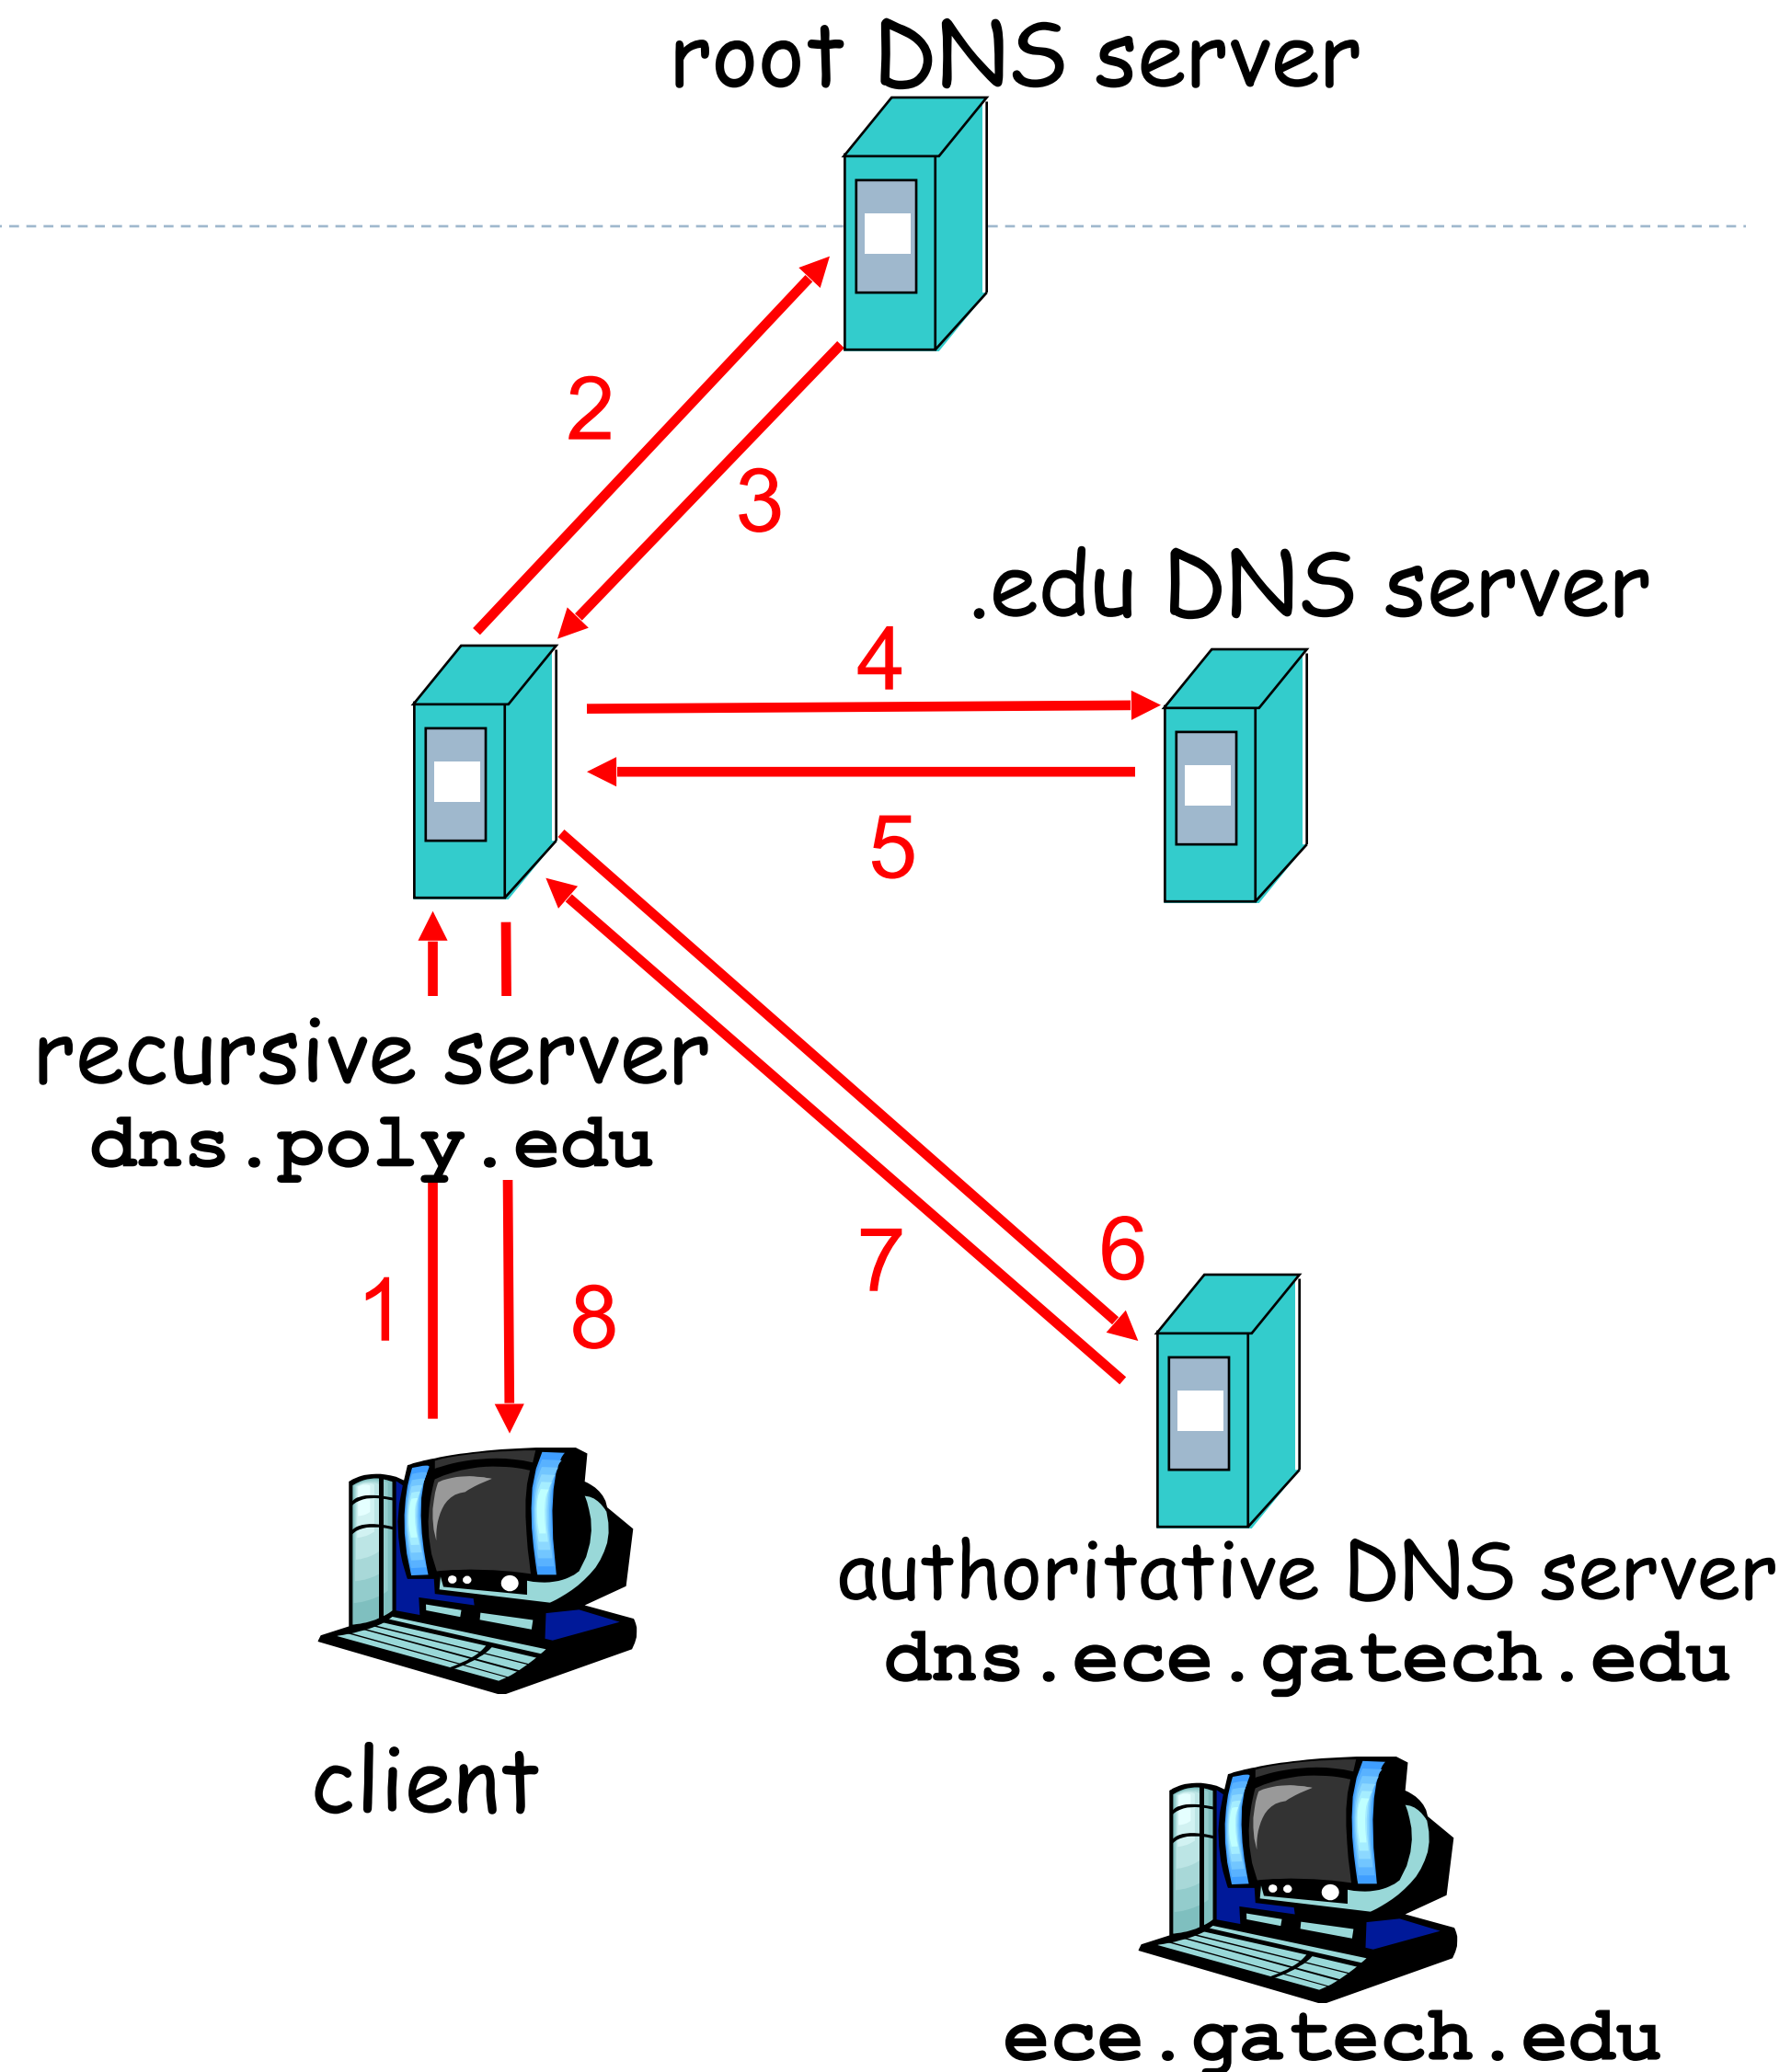


Client wants IP for www.amazon.com; 1st approx:

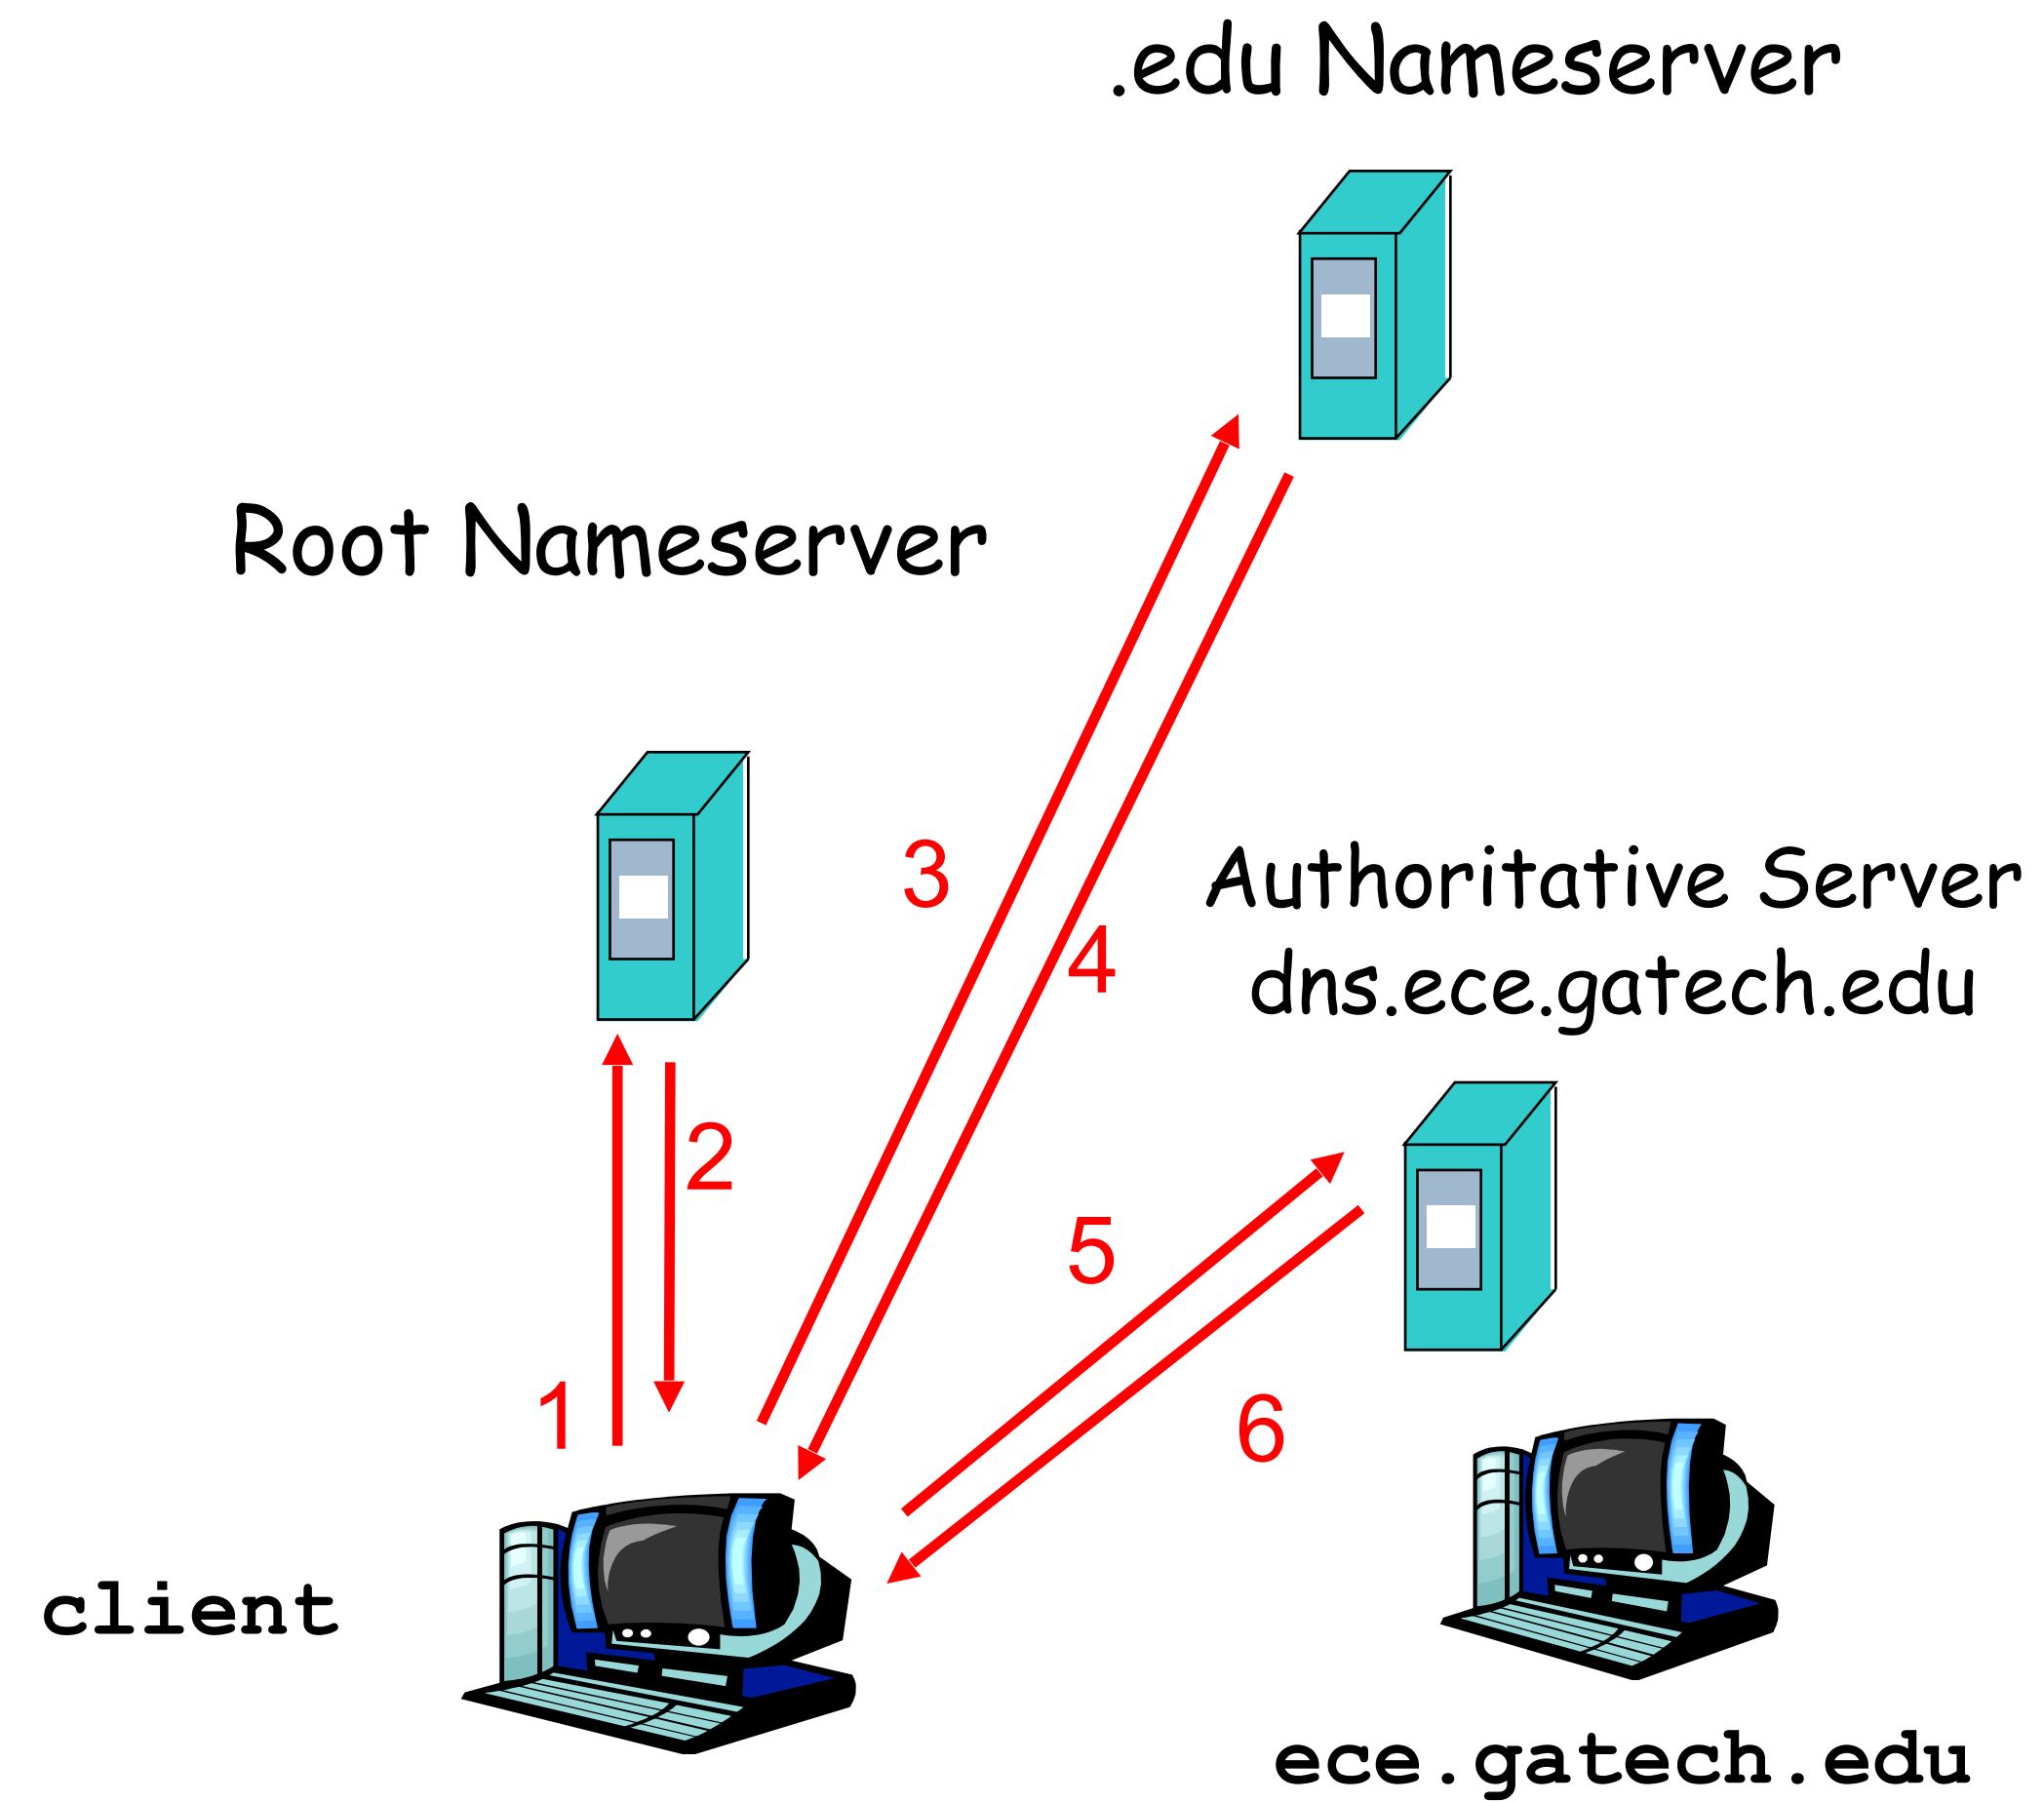
- ▶ Client queries a root name server to find "com" name server
- ▶ Client queries "com" name server to get amazon.com name server
- ▶ Client queries amazon.com name server to get IP address for www.amazon.com

Recursive Query

- ▶ Host at ex.poly.edu wants IP address for ece.gatech.edu
- ▶ Host sends a "**recursion-requested**" query request to dns.poly.edu.
- ▶ Local DNS server does a "**recursive**" search. This requires contacting several other DNS servers before the final answer is given to host.



Iterative Query



DNS: caching and updating records

- ▶ once (any) name server learns mapping, it **caches** mapping
- ▶ cache entries timeout (disappear) after some time. Sometimes based on ttl but *not always*
- ▶ TLD servers typically cached in local name servers
 - ▶ Thus root name servers not often visited

Inserting records into DNS

- ▶ Example: just created startup "Network Utopia"
- ▶ Register name networkutopia.com at a **registrar** (e.g., Network Solutions)
 - ▶ Need to provide registrar with names and IP addresses of your authoritative name server
 - ▶ Registrar inserts two RRs into the com TLD server:

```
(networkutopia.com, dns1.networkutopia.com, NS)  
(dns1.networkutopia.com, 212.212.212.1, A)
```

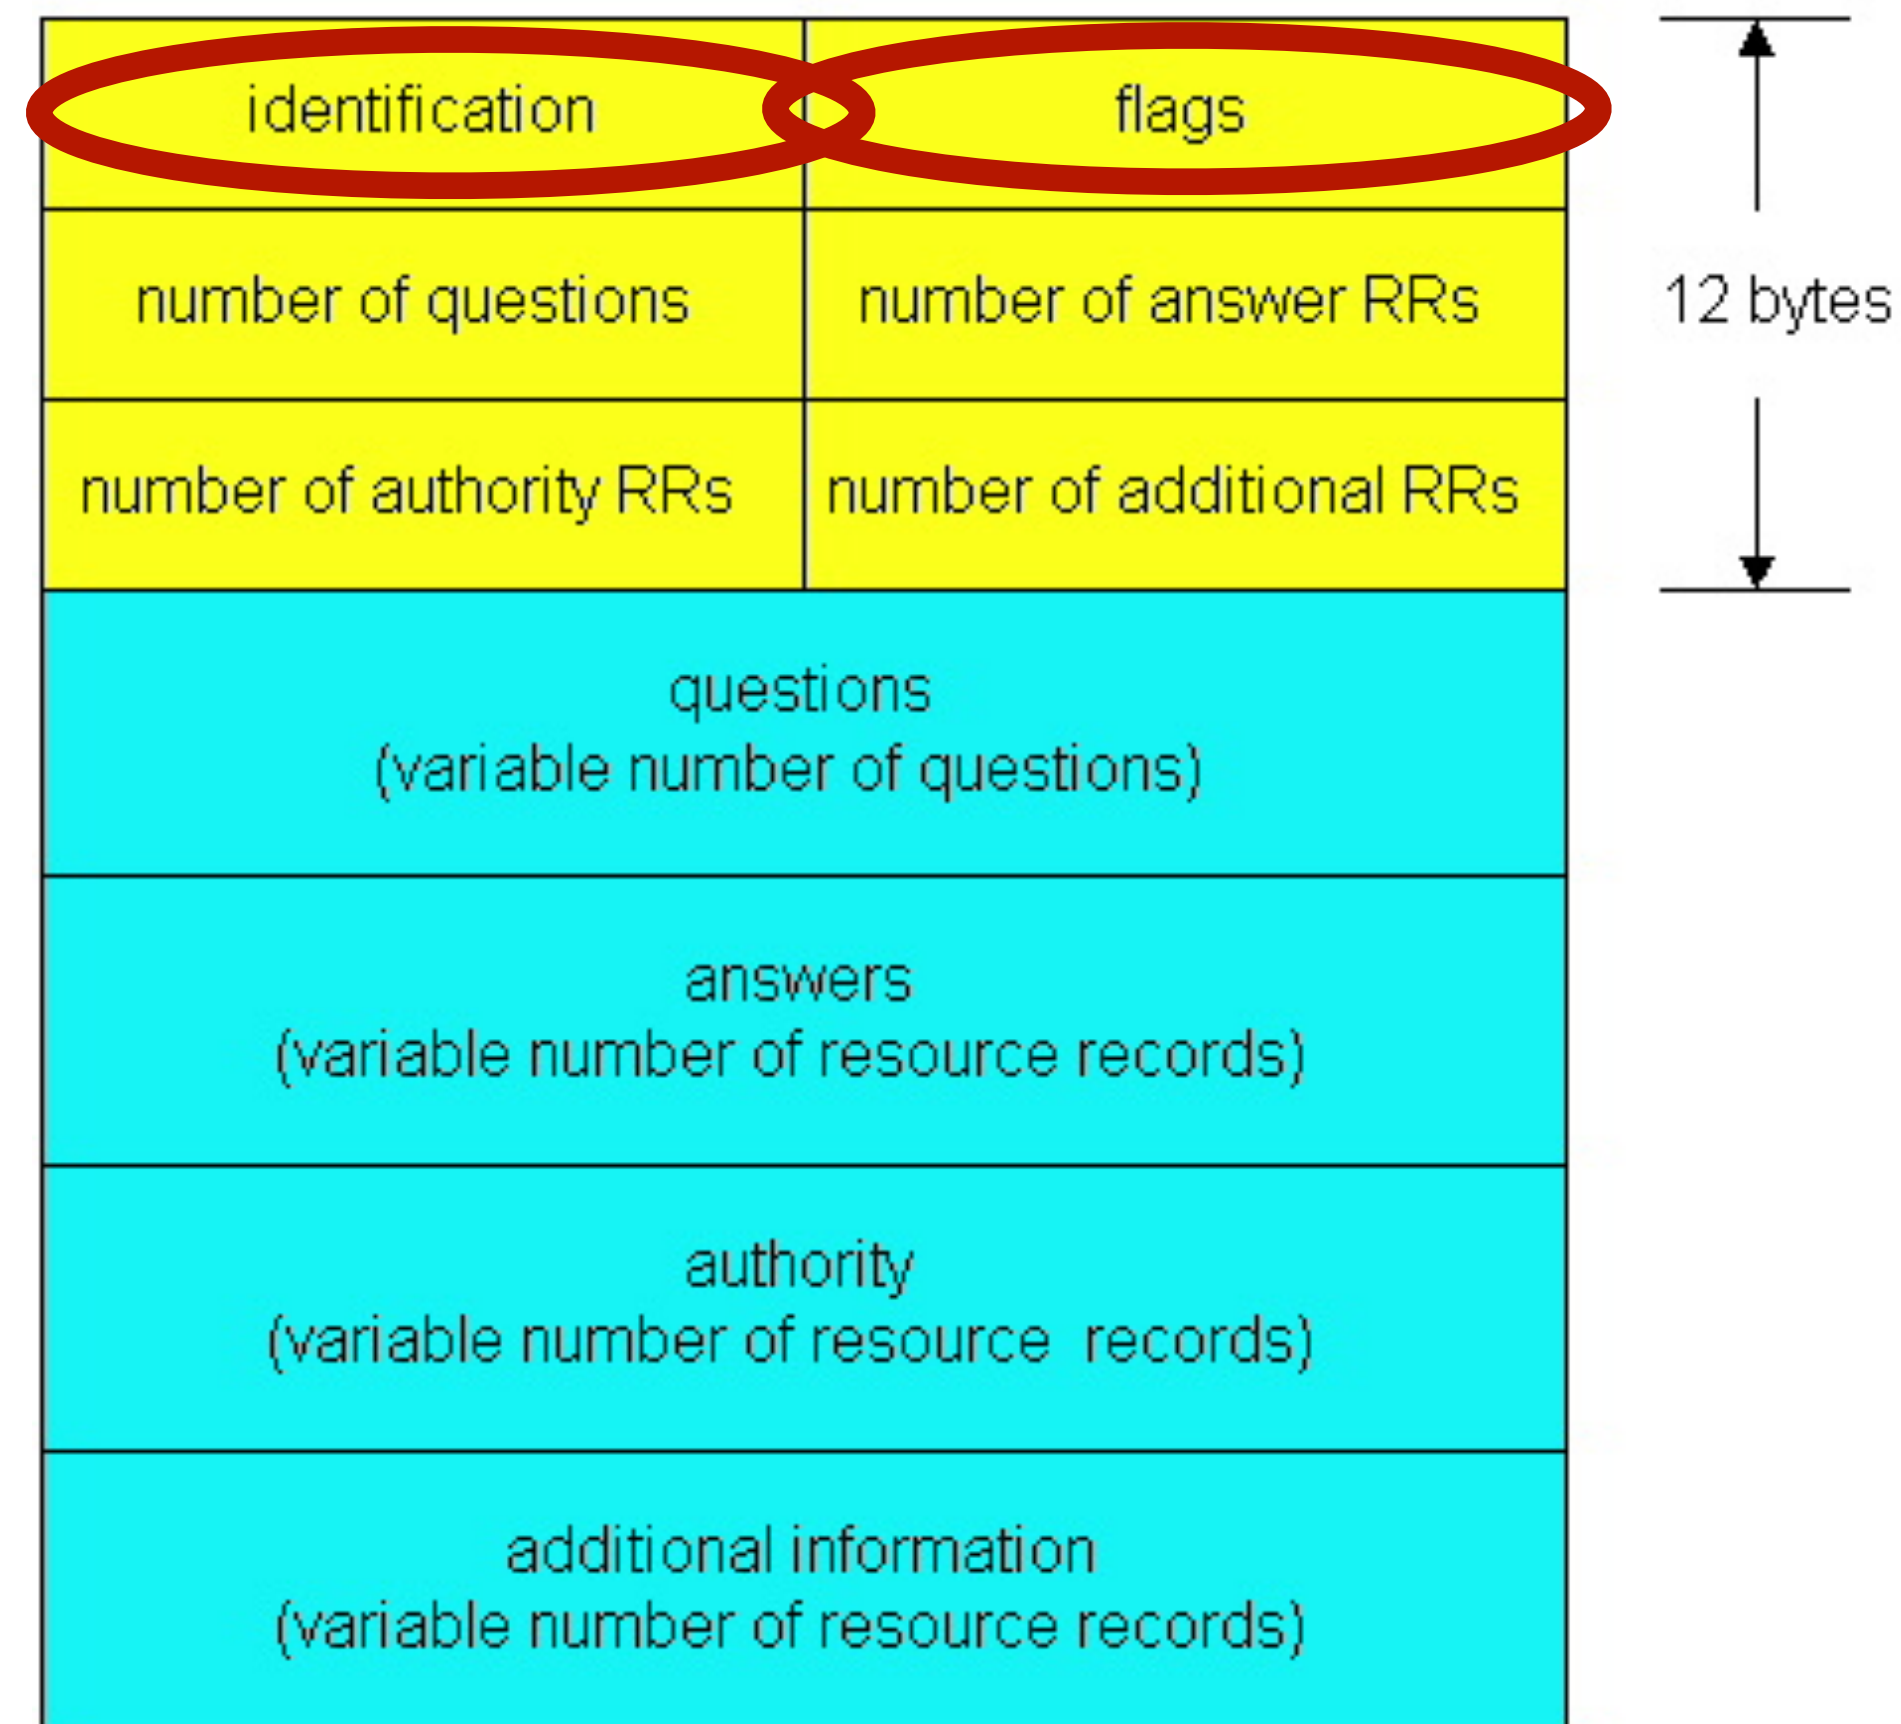
- ▶ At authoritative server, can configure Type A record for www.networkutopia.com and Type MX record for networkutopia.com

DNS protocol, messages

DNS protocol: UDP (port 53) **query** and **reply** messages, both with same **message format**

msg header

- ❑ **identification**: 16 bit #
for query, reply to query
uses same #
- ❑ **flags**:
 - ❖ query or reply
 - ❖ recursion desired
 - ❖ recursion available
 - ❖ reply is authoritative



dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.                IN      A

;; ANSWER SECTION:
eecs.mit.edu.                21600   IN      A      18.62.1.6

;; AUTHORITY SECTION:
mit.edu.                    11088   IN      NS      BITSY.mit.edu.
mit.edu.                    11088   IN      NS      BITSY.mit.edu.
mit.edu.                    11088   IN      NS      BITSY.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.            166408  IN      A      18.72.0.3
BITSY.mit.edu.             166408  IN      A      18.72.0.3
W20NS.mit.edu.             126738  IN      A      18.70.0.160
```

In general, a single *Resource Record* (RR) like this includes, left-to-right, a DNS name, a *time-to-live*, a family (IN for our purposes - ignore), a type (A here, which stands for "Address"), and an associated value


dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: 0x00, flags: 0x00, QNAME: eecs.mit.edu, QTYPE: A, QCLASS: IN, QCOUNT: 1, ANCOUNT: 3, AUTHORITY: 3, ADDITIONAL: 3
```

“Answer” tells us the IP address associated with eecs.mit.edu is 18.62.1.6 and we can *cache* the result for 21,600 seconds

```
;; QUESTION SECTION:
;eecs.mit.edu.
```

```
;; ANSWER SECTION:
eecs.mit.edu.
```



```
;; AUTHORITY SECTION:
mit.edu. 11088 IN NS BITSY.mit.edu.
mit.edu. 11088 IN NS W20NS.mit.edu.
mit.edu. 11088 IN NS STRAWB.mit.edu.
```

```
;; ADDITIONAL SECTION:
STRAWB.mit.edu. 126738 IN A 18.71.0.151
BITSY.mit.edu. 166408 IN A 18.72.0.3
W20NS.mit.edu. 126738 IN A 18.70.0.160
```


dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-API
;; global options: +cr
;; Got answer:
;; ->>HEADER<<- opcode
;; flags: qr rd ra; QU
```

```
;; QUESTION SECTION:
;eecs.mit.edu.
```

```
;; ANSWER SECTION:
eecs.mit.edu.
```

```
;; AUTHORITY SECTION:
mit.edu.
mit.edu.
mit.edu.
```

```
;; ADDITIONAL SECTION:
STRAWB.mit.edu.
BITSY.mit.edu.
W20NS.mit.edu.
```

“Authority” tells us the *name servers* responsible for the answer. Each RR gives the *hostname* of a different name server (“NS”) for names in `mit.edu`. We should cache each record for 11,088 seconds.

If the **“Answer”** had been empty, then the resolver’s next step would be to send the original query to one of these name servers.

21600	IN	A	18.62.1.6
-------	----	---	-----------

11088	IN	NS
11088	IN	NS
11088	IN	NS

BITSY.mit.edu.
W20NS.mit.edu.
STRAWB.mit.edu.

126738	IN	A	18.71.0.151
166408	IN	A	18.72.0.3
126738	IN	A	18.70.0.160

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3
```

```
;; QUESTION SECTION:
;eecs.mit.edu.
```

```
;; ANSWER SECTION
eecs.mit.edu.
```

```
;; AUTHORITY SECTION
mit.edu.
mit.edu.
mit.edu.
```

```
;; ADDITIONAL SECTION:
STRAWB.mit.edu.
BITSY.mit.edu.
W20NS.mit.edu.
```

“Additional” provides extra information to save us from making separate lookups and helps with bootstrapping (**glue records**). Here, it tells us the IP addresses for the hostnames of the name servers. We add these to our cache.

11088	IN	NS	BITSY.mit.edu.
11088	IN	NS	W20NS.mit.edu.
11088	IN	NS	STRAWB.mit.edu.
126738	IN	A	18.71.0.151
166408	IN	A	18.72.0.3
126738	IN	A	18.70.0.160

DNS Security Threats

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1  ANSWER: 1  AUTHORITY: 3  ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.

;; ANSWER SECTION:
eecs.mit.edu.      210

;; AUTHORITY SECTION:
mit.edu.           11088      IN         NS         BITSY.mit.edu.
mit.edu.           11088      IN         NS         W20NS.mit.edu.
mit.edu.           11088      IN         NS         STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.    126738     IN         A          18.71.0.151
BITSY.mit.edu.     166408     IN         A          18.72.0.3
W20NS.mit.edu.     126738     IN         A          18.70.0.160
```

What if the mit.edu name server is untrustworthy?
Could its operator steal, say, all of our web surfing to Facebook?

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.

;; ANSWER SECTION:
eecs.mit.edu.      21000 IN      A      18.71.0.151

;; AUTHORITY SECTION:
mit.edu.           11088 IN      NS      BITSY.mit.edu.
mit.edu.           11088 IN      NS      W20NS.mit.edu.
mit.edu.           11088 IN      NS      STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.    126738 IN      A      18.71.0.151
BITSY.mit.edu.     166408 IN      A      18.72.0.3
W20NS.mit.edu.     126738 IN      A      18.70.0.160
```

Let's look at a flaw in the original DNS design (since fixed)

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3
```

```
;; QUESTION SECTION:
;eecs.mit.edu.
```

```
;; ANSWER SECTION:
eecs.mit.edu.
```

21600	IN	A	18.62.1.6
-------	----	---	-----------

```
;; AUTHORITY SECTION:
```

mit.edu.	11088	IN	NS	BITSY.mit.edu.
mit.edu.	11088	IN	NS	W20NS.mit.edu.
mit.edu.	11088	IN	NS	www.facebook.com.

```
;; ADDITIONAL SECTION:
```

www.facebook.com	11088	IN	A	18.6.6.6
BITSY.mit.edu.	166408	IN	A	18.72.0.3
W20NS.mit.edu.	126738	IN	A	18.70.0.160

What could happen if the mit.edu name server returns the following to us instead?

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3
```

```
;; QUESTION SECTION:
;eecs.mit.edu.
```

```
;; ANSWER SECTION:
eecs.mit.edu.
```

```
;; AUTHORITY SECTION:
mit.edu.
mit.edu.
mit.edu.
```

```
;; ADDITIONAL SECTION:
www.facebook.com
BITSY.mit.edu.
W20NS.mit.edu.
```

We'd dutifully store in our cache a mapping of `www.facebook.com` to an IP address under MIT's control. (It could have been any IP address they wanted, not just one of theirs.)

11088	IN	NS	BITSY.mit.edu.
11088	IN	NS	W20NS.mit.edu.
11088	IN	NS	www.facebook.com.
11088	IN	A	18.6.6.6
166408	IN	A	18.72.0.3
126738	IN	A	18.70.0.160

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3
```

```
;; QUESTION SECTION:
;eecs.mit.edu.
```

```
;; ANSWER SECTION:
eecs.mit.edu.
```

```
;; AUTHORITY SECTION:
mit.edu.
mit.edu.
mit.edu.
```

```
;; ADDITIONAL SECTION:
www.facebook.com
BITSY.mit.edu.
W20NS.mit.edu.
```

Next time one of our clients starts to connect to `www.facebook.com`, it will ask our resolver for the corresponding IP address. The resolver will find the answer in its cache and return **18.6.6.6** 😬

mit.edu.	11088	IN	NS	W20NS.mit.edu.
mit.edu.	11088	IN	NS	www.facebook.com.
www.facebook.com	11088	IN	A	18.6.6.6
BITSY.mit.edu.	166408	IN	A	18.72.0.3
W20NS.mit.edu.	126738	IN	A	18.70.0.160

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3
```

```
;; QUESTION SECTION:
```

```
;eecs.mit.edu.                IN      A
```

```
;; ANSWER SECTION:
```

```
eecs.mit.edu.                31088 IN      A
```

How do we fix such *cache poisoning*?

```
;; AUTHORITY SECTION:
```

mit.edu.	11088	IN	NS	BITSY.mit.edu.
mit.edu.	11088	IN	NS	W20NS.mit.edu.
mit.edu.	11088	IN	NS	www.facebook.com.

```
;; ADDITIONAL SECTION:
```

www.facebook.com	11088	IN	A	18.6.6.6
BITSY.mit.edu.	166408	IN	A	18.72.0.3
W20NS.mit.edu.	126738	IN	A	18.70.0.160

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-AP
;; global options: +c
;; Got answer:
;; ->>HEADER<<- opcode
;; flags: qr rd ra; Q

;; QUESTION SECTION:
;eecs.mit.edu.

;; ANSWER SECTION:
eecs.mit.edu.
```

```
;; AUTHORITY SECTION:
```

mit.edu.	11088	IN	NS	BITSY.mit.edu.
mit.edu.	11088	IN	NS	W20NS.mit.edu.
<u>mit.edu.</u>	<u>11088</u>	<u>IN</u>	<u>NS</u>	<u>www.facebook.com.</u>

```
;; ADDITIONAL SECTION:
```

<u>www.facebook.com</u>	<u>11088</u>	<u>IN</u>	<u>A</u>	<u>18.6.6.6</u>
BITSY.mit.edu.	166408	IN	A	18.72.0.3
W20NS.mit.edu.	126738	IN	A	18.70.0.160

Don't accept **Additional** records unless they're for the domain of the name server we queried

E.g., contacting a name server for mit.edu ⇒ only accept additional records from *.mit.edu

No extra risk in accepting these since server could return them to us directly in an **Answer** anyway.

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-AP
;; global options: +c
;; Got answer:
;; ->>HEADER<<- opcode
;; flags: qr rd ra; Q
```

```
;; QUESTION SECTION:
;eecs.mit.edu.
```

```
;; ANSWER SECTION:
eecs.mit.edu.
```

;; AUTHORITY SECTION:

mit.edu.

mit.edu.

mit.edu.

~~;; ADDITIONAL SECTION:~~

www.facebook.com

BITSY.mit.edu.

W20NS.mit.edu.

Don't accept **Additional** records unless they're for the domain of the name server we queried

E.g., contacting a name server for `mit.edu` \Rightarrow only accept additional records from `*.mit.edu`

No extra risk in accepting these since server could return them to us directly in an **Answer** anyway.

This is called “bailiwick checking”.

11000 TN NC DTTCV mit edu

11 bailiwick | 'bālə,wɪk |

noun

1 (one's bailiwick) one's sphere of operations or particular area of interest: *you never give the presentations—that's my bailiwick.*

Bailiwick Checking

Bailiwick rules (for querying for `www.google.com`):

1. Root servers can return any record (bailiwick is everything)

- Ex: `k.root-server.net` is a root server. It responds with `a.gtld-servers.net` as one of the `.com` name servers (in the authority + additional/glue section).

2. `.com` NS can return any record within `.com` (bailiwick is `.com`)

- Ex: `a.gtld-server.net` returns `ns1.google.com` as a `google.com` name server (in the authority and additional)

3. `google.com` NS can return any record within `google.com` (bailiwick is `google.com`)

- Ex: Returns answer for `www.google.com`

Modern DNS software enforces bailiwick checking (not in original DNS design though)

DNS Threats: Spoofing

- If an attacker observes the DNS identification number, they can spoof a DNS response to a victim, giving them false DNS answers (e.g., respond to a DNS query for `mail.google.com` with an A record pointing to the attacker server)
- Has been used in real attacks, including censorship.

16 bits	16 bits
SRC=53	DST=53
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

DNS Threats: Spoofing

What about *blind spoofing*?

- Say we look up `mail.google.com`; how can an **off-path** attacker feed us a **bogus A answer** before the legitimate server replies?
- How can such a **remote** attacker even know we are looking up `mail.google.com`?

Suppose, e.g., we visit a web page under their control:

```
... ...
```

16 bits	16 bits
SRC=53	DST=53
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

DNS Threats: Spoofing

What about *blind spoofing*?

- Say we look up mail.google.com; how can an **off-path** attacker feed us a **bogus A answer** before the legitimate

- How can we cause our browser to try to fetch an image from mail.google.com. To do that, our browser first has to look up the IP address associated with that name.

Suppose, e.g., we visit a web page under their control:

... ...

16 bits	16 bits
SRC=53	DST=53
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (resource records)	
Answers (resource records)	
Authority (resource records)	
Additional information (variable # of resource records)	

DNS Blind Spoofing, con't

Fix?

Once they know we're looking it up, they just have to guess the Identification field, and reply before legit server.

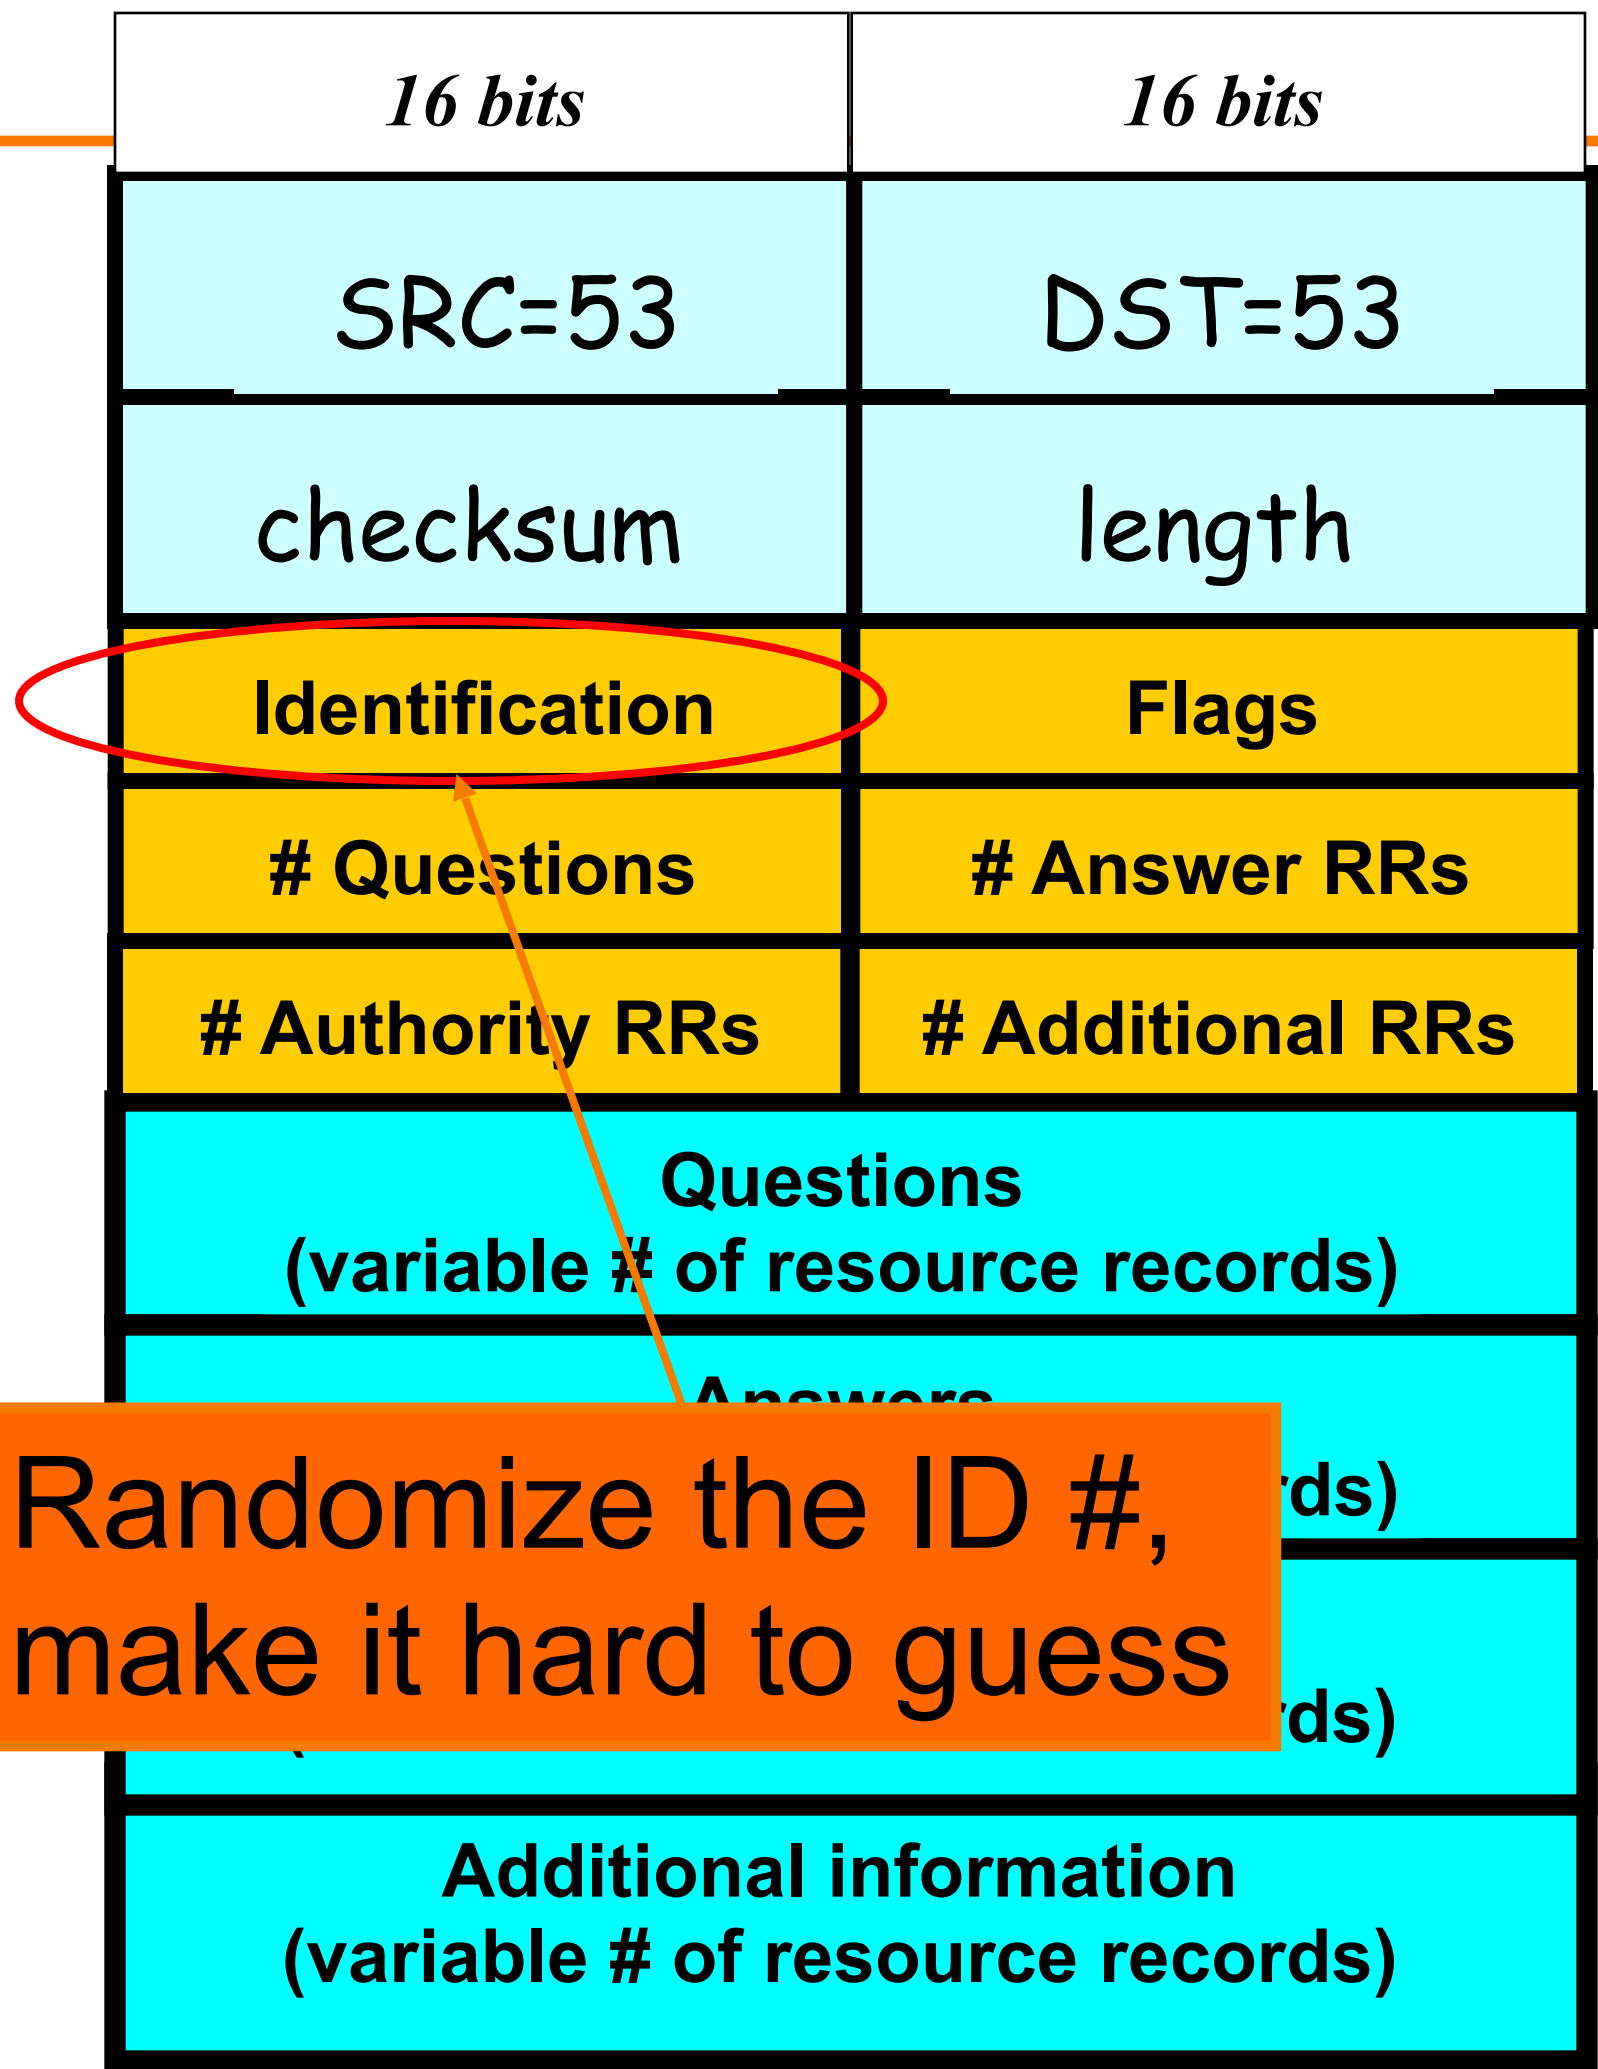
How hard is that?

Originally, identification field incremented by 1 for each request. How does attacker guess it?

(Assuming attacker controls their domain's name server)

They observe ID k here

So this will be k+1



DNS Blind Spoofing, con't

Once we **randomize** the Identification, attacker has a 1/65536 chance of guessing it correctly.

Are we pretty much safe?

Attacker can send *lots* of replies, not just one ...

However: once a reply from legit server arrives (with correct Identification), it's **cached** and no more opportunity to poison it. Victim is inoculated!

16 bits	16 bits
SRC=53	DST=53
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

Unless attacker can send 1000s of replies before legit arrives, we're likely safe -
pew! ?

DNS Blind Spoofing (Kaminsky 2008)

- Two key ideas:
 - Spoof uses **Additional** field (rather than **Answer**)
 - Attacker can get around caching of legit replies by generating a **series** of *different* name lookups:

```
  
  
  
...  

```

Kaminsky Blind Spoofing, con't

For each lookup of *randomk.google.com*, attacker **spoofs** a **bunch** of records like this, each with a different Identifier

;; QUESTION SECTION:
;randomk.google.com.

IN A

;; ANSWER SECTION:
randomk.google.com

21600 IN A *doesn't matter*

;; AUTHORITY SECTION:
google.com.

11088 IN NS mail.google.com

;; ~~ADDITIONAL SECTION:~~

mail.google.com 126738 IN A 6.6.6.6

Once they win the race, not only have they poisoned mail.google.com ...

Kaminsky Blind Spoofing, con't

For each lookup of *randomk.google.com*, attacker **spoofs** a **bunch** of records like this, each with a different Identifier

;; QUESTION SECTION:
;randomk.google.com.

IN A

;; ANSWER SECTION:

randomk.google.com 21600 IN A *doesn't matter*

;; AUTHORITY SECTION:

google.com. 11088 IN NS mail.google.com

;; ADDITIONAL SECTION:

mail.google.com 126738 IN A 6.6.6.6

Once they win the race, not only have they poisoned *mail.google.com* ... but also the cached NS record for *google.com*'s name server - so any **future** *X.google.com* lookups *go through the attacker's machine*

Defending Against Blind Spoofing

Central problem: all that tells a client they should accept a response is that it matches the **Identification** field.

With only **16 bits**, it lacks sufficient **entropy**: even if truly random, the *search space* an attacker must *brute force* is too small.

Where can we get more entropy?

16 bits	16 bits
SRC=53	DST=53
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

Defending Against Blind Spoofing

Central problem: all that tells a client they should accept a response is that it matches the Identification field.

With only 16 bits, it lacks sufficient entropy: even if truly random, the *search space* an attacker must *brute force* is too small.

Where can we get more entropy? (*Without* requiring a protocol change.)

16 bits	16 bits
SRC=53	DST=53
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

Defending Against Blind Spoofing

Total entropy: 16 bits

For requestor to receive DNS reply, needs both correct **Identification** and correct **ports**.

On a request, DST port = 53.
SRC port usually also 53 - but not fundamental, just **convenient**.

16 bits	16 bits
SRC=53	DST=53
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

Defending Against Blind Spoofing

“Fix”: client uses **random** source port \Rightarrow attacker doesn’t know correct dest. port to use in reply

Total *entropy*: ? bits

16 bits	16 bits
SRC=53	DST=rnd
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

Defending Against Blind Spoofing

Total entropy: 32 bits

“Fix”: client uses random source port \Rightarrow attacker doesn’t know correct dest. port to use in reply

32 bits of entropy makes it **orders of magnitude** harder for attacker to guess all the necessary fields and dupe victim into accepting spoof response.

16 bits	16 bits
SRC=53	DST=rnd
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

Defending Against Blind Spoofing

Total entropy: 32 bits

“Fix”: client uses random source port \Rightarrow attacker doesn’t know correct dest. port to use in reply

32 bits of entropy makes it **orders of magnitude** harder for attacker to guess all the necessary fields and dupe victim into accepting spoof response.

This is what primarily “secures” DNS against blind spoofing today. (Note: not all resolvers have implemented random source ports!)

16 bits	16 bits
SRC=53	DST=rnd
checksum	length
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	

DNS Poisoning

Malicious DNS name server could provide invalid authority/additional section for domains outside of its zone, poisoning the DNS cache.

- Prevent this with DNS bailiwick checking (only accept DNS records from a name server that are within its zone)

DNS spoofing requires the attacker to match the DNS ID #.

- Need to make ID # random to make this unpredictable.
- Attacker can guess ID # blindly, but must race against real DNS response. Kaminsky attack makes attack effective (try random subdomains, only need to win race once).
 - Made hard by randomizing the request source port also, adding more bits of entropy that need to be guessed correctly.
- Still easy if attacker can observe the DNS request.

Securing DNS Lookups

- How can we ensure when clients look up names with DNS, they can **trust** answers they receive?
- Idea #1: do DNS lookups over TLS
 - (assuming either we run DNS over TCP, or we use “Datagram TLS”)

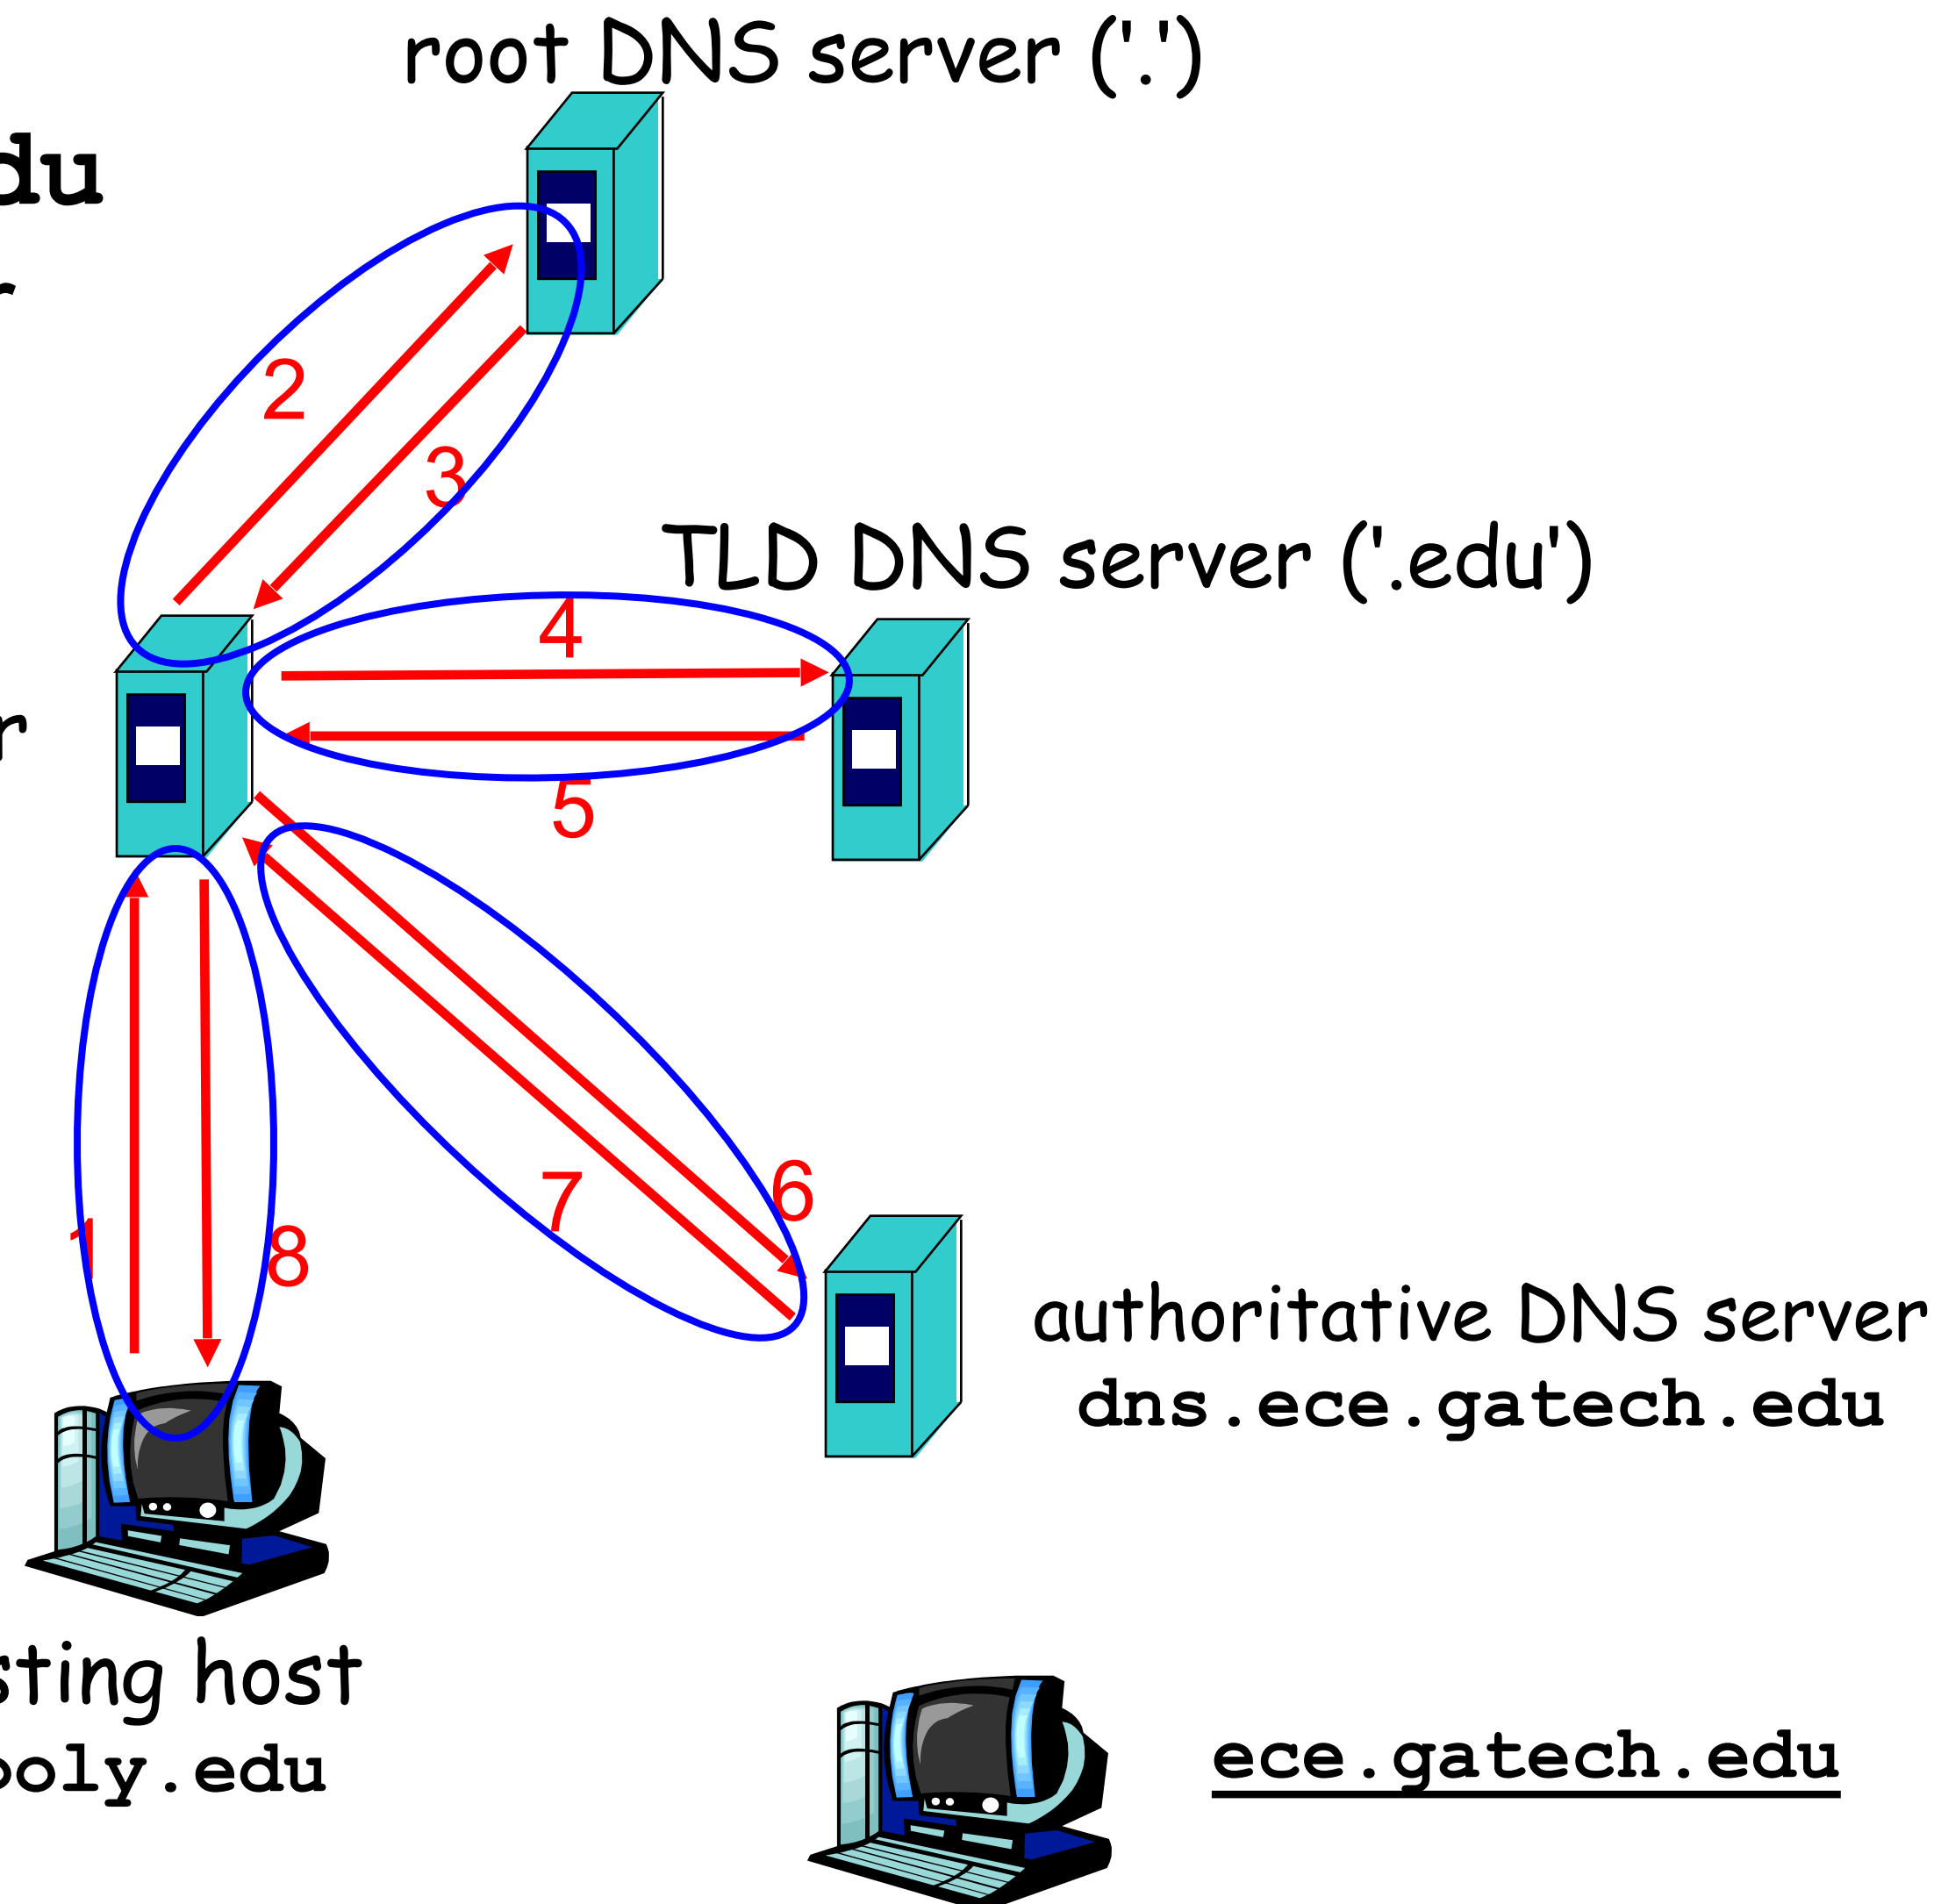
Securing DNS using SSL / TLS?

Host at **xyz.poly.edu**
wants IP address for
ece.gatech.edu

local DNS server
(resolver)
128.238.1.68

Idea: connections
{1,8}, {2,3}, {4,5}
and {6,7} all run
over SSL / TLS

requesting host
xyz.poly.edu



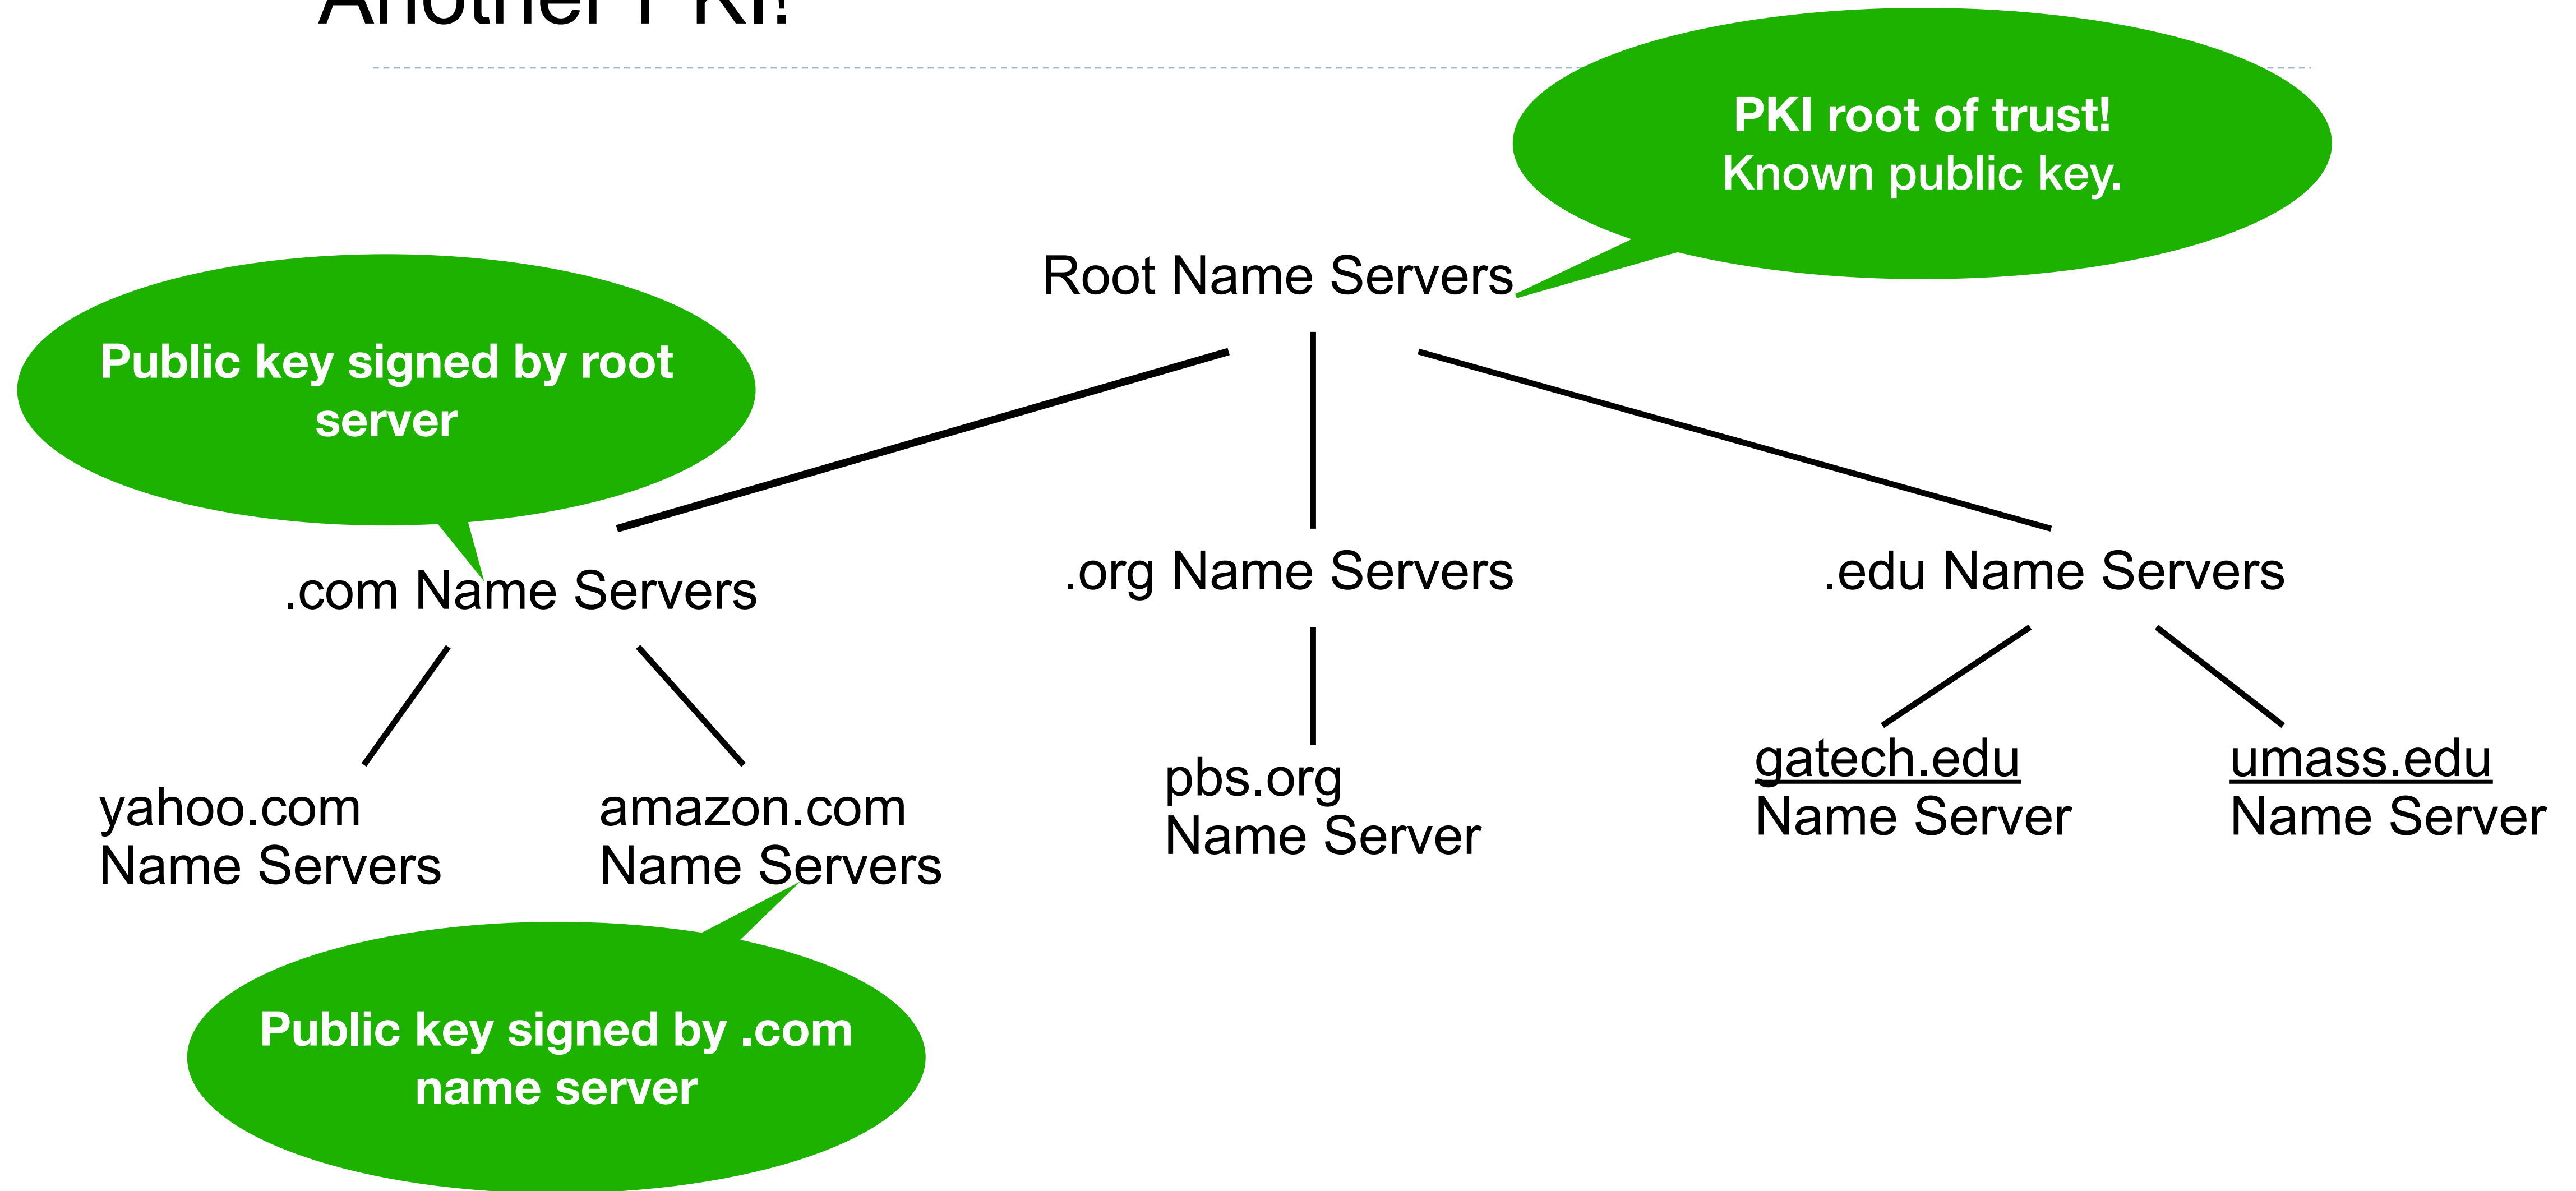
Securing DNS Lookups

- How can we ensure when clients look up names with DNS, they can trust answers they receive?
- Idea #1: do DNS lookups over TLS
 - (assuming either we run DNS over TCP, or we use “Datagram TLS”)
 - Issues?
 - **Performance**: DNS is very lightweight. TLS is not.
 - **Caching**: crucial for DNS scaling. But then how do we keep authentication assurances?
 - ***Object** security vs. **Channel** security*

Securing DNS Lookups

- How can we ensure when clients look up names with DNS, they can trust answers they receive?
- Idea #1: do DNS lookups over TLS
 - (assuming either we run DNS over TCP, or we use “Datagram TLS”)
 - Issues?
 - Performance: DNS is very lightweight. TLS is not.
 - Caching: crucial for DNS scaling. But then how do we keep authentication assurances?
 - *Object security vs. Channel security*
- Idea #2: make DNS results like *certs*
 - I.e., a **verifiable signature** that guarantees who generated a piece of data; signing happens **off-line**

Another PKI!




Operation of DNSSEC

- DNSSEC = standardized DNS security extensions currently being deployed
- As a resolver works its way from DNS root down to final name server for a name, at each level it gets a signed statement regarding the key(s) used by the next level
 - This builds up a chain of trusted keys
 - Resolver has root's key **wired into it**
- The final answer that the resolver receives is signed by that level's key
 - Resolver can trust it's the right key because of chain of support from higher levels
- *All keys as well as signed results are **cacheable***

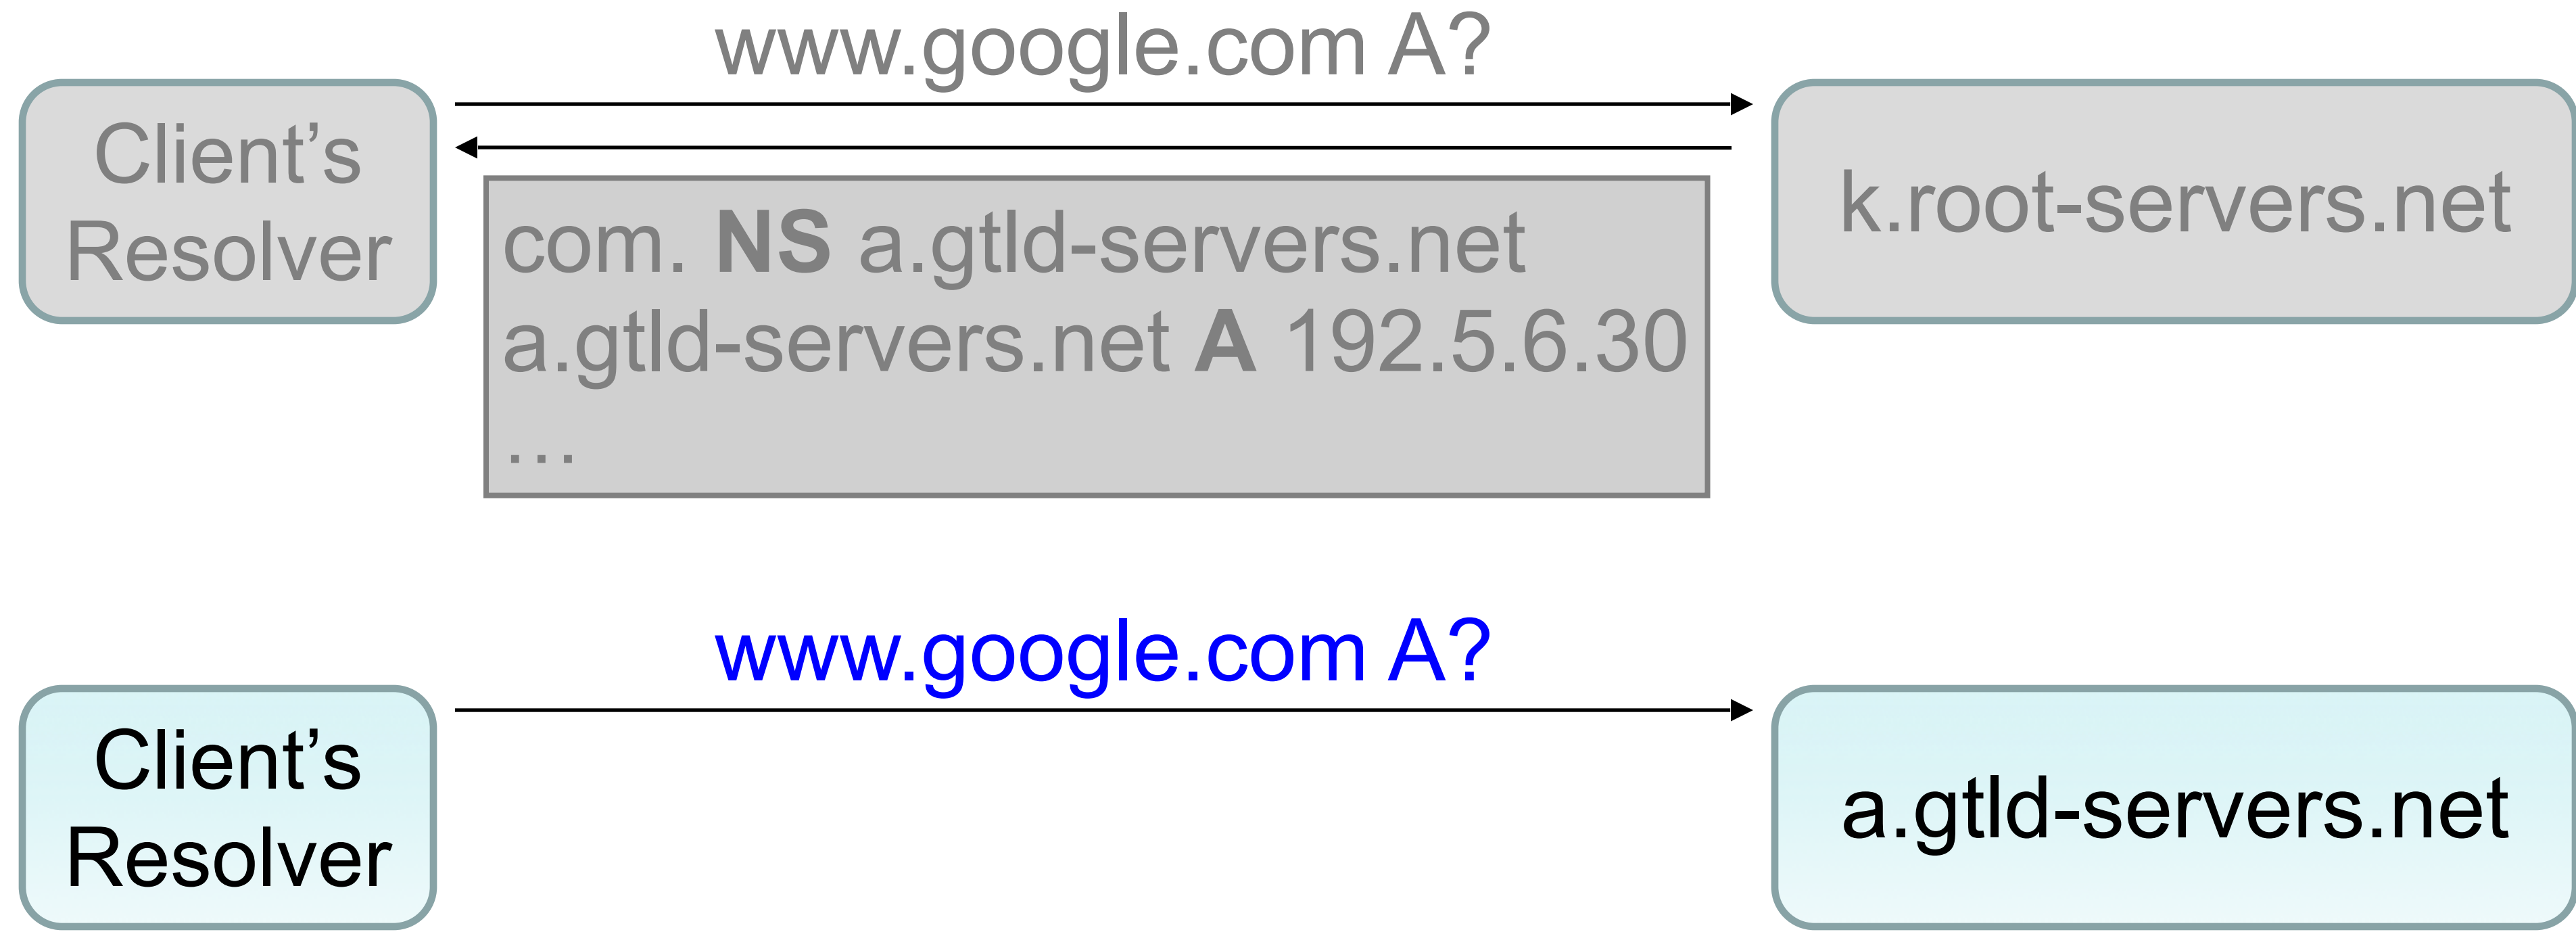
Ordinary DNS:



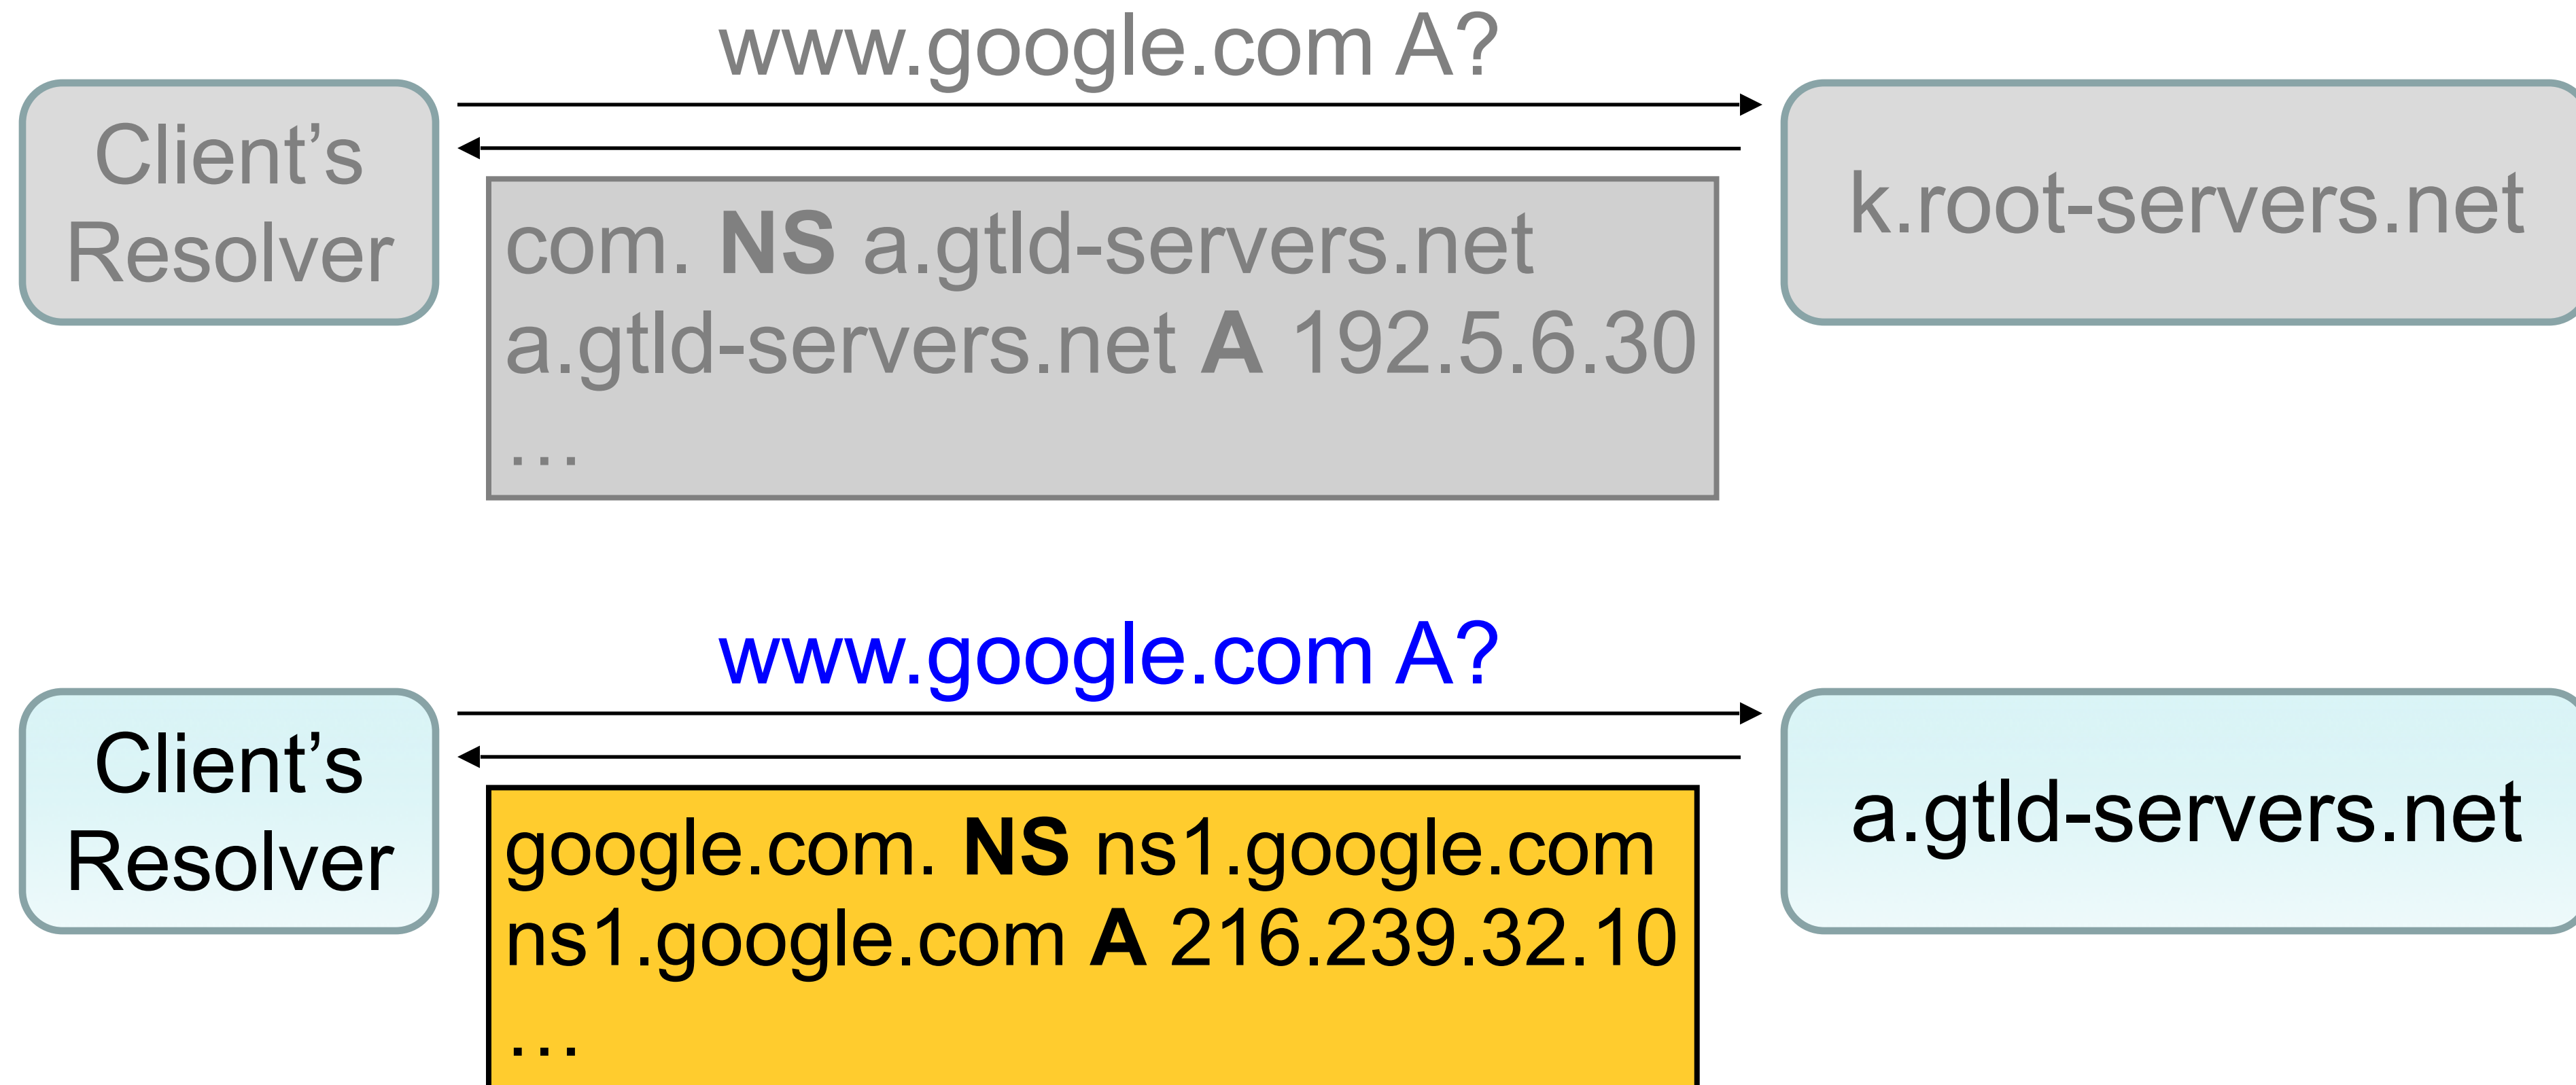
Ordinary DNS:



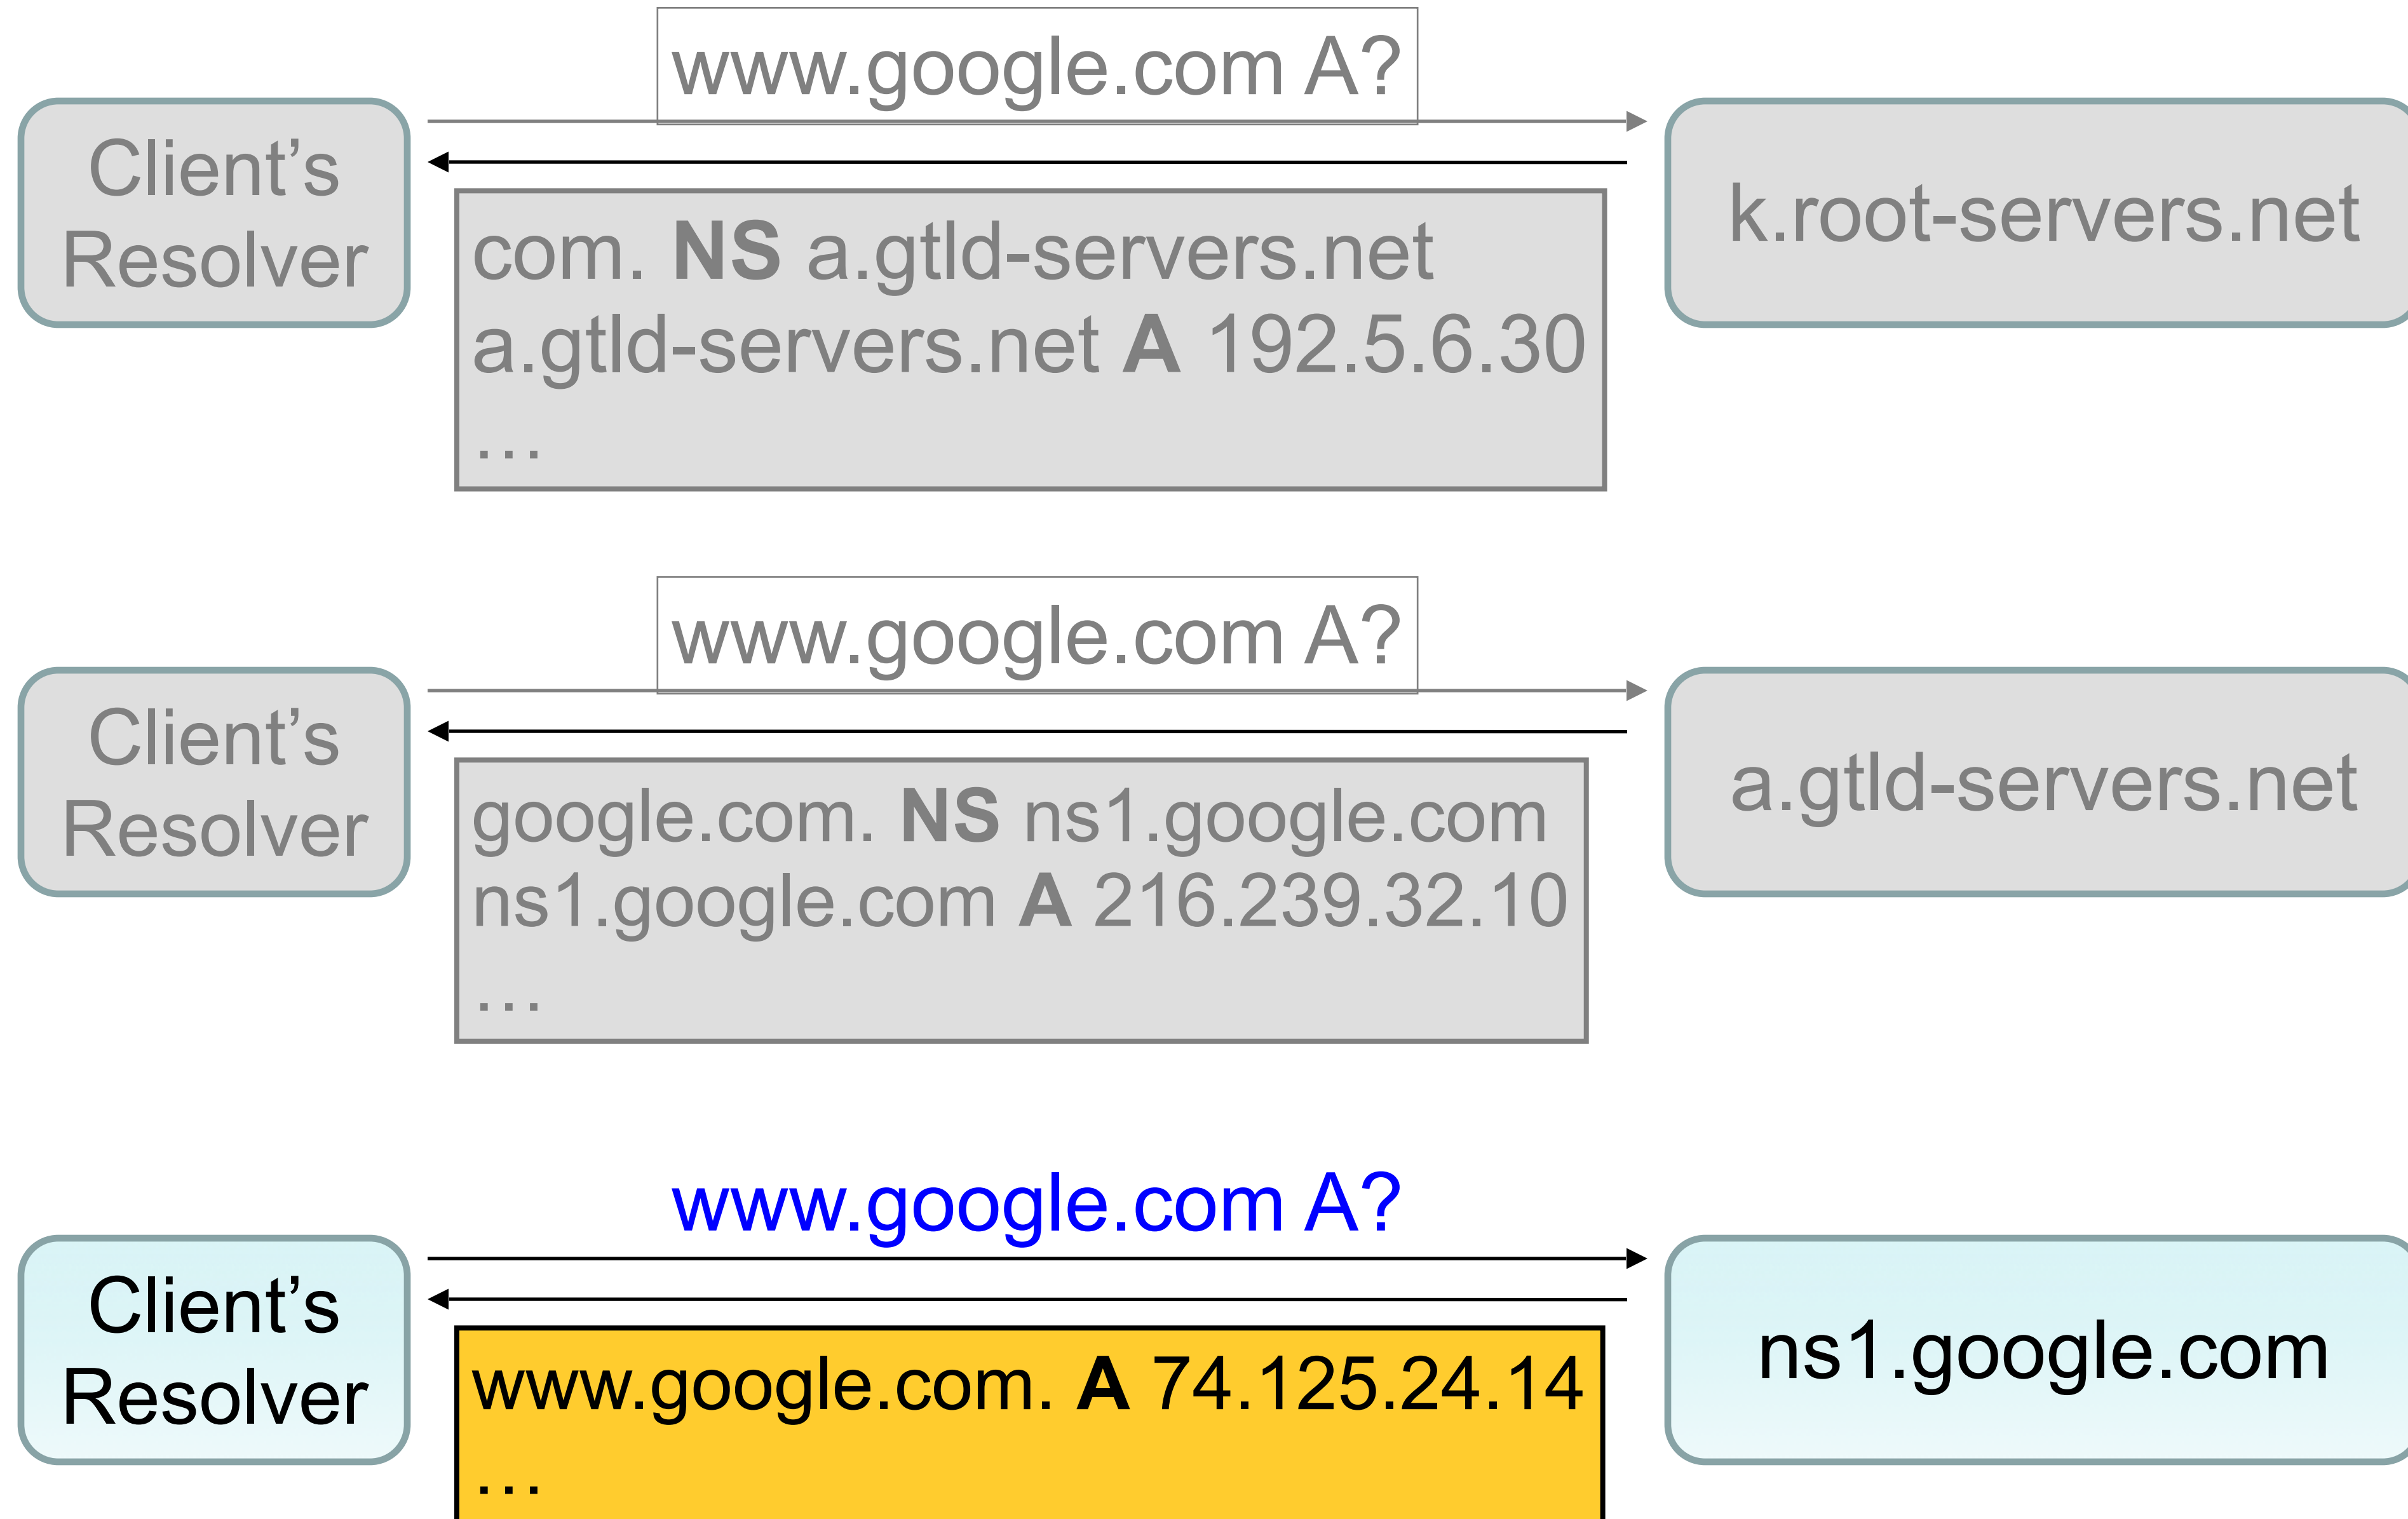
Ordinary DNS:




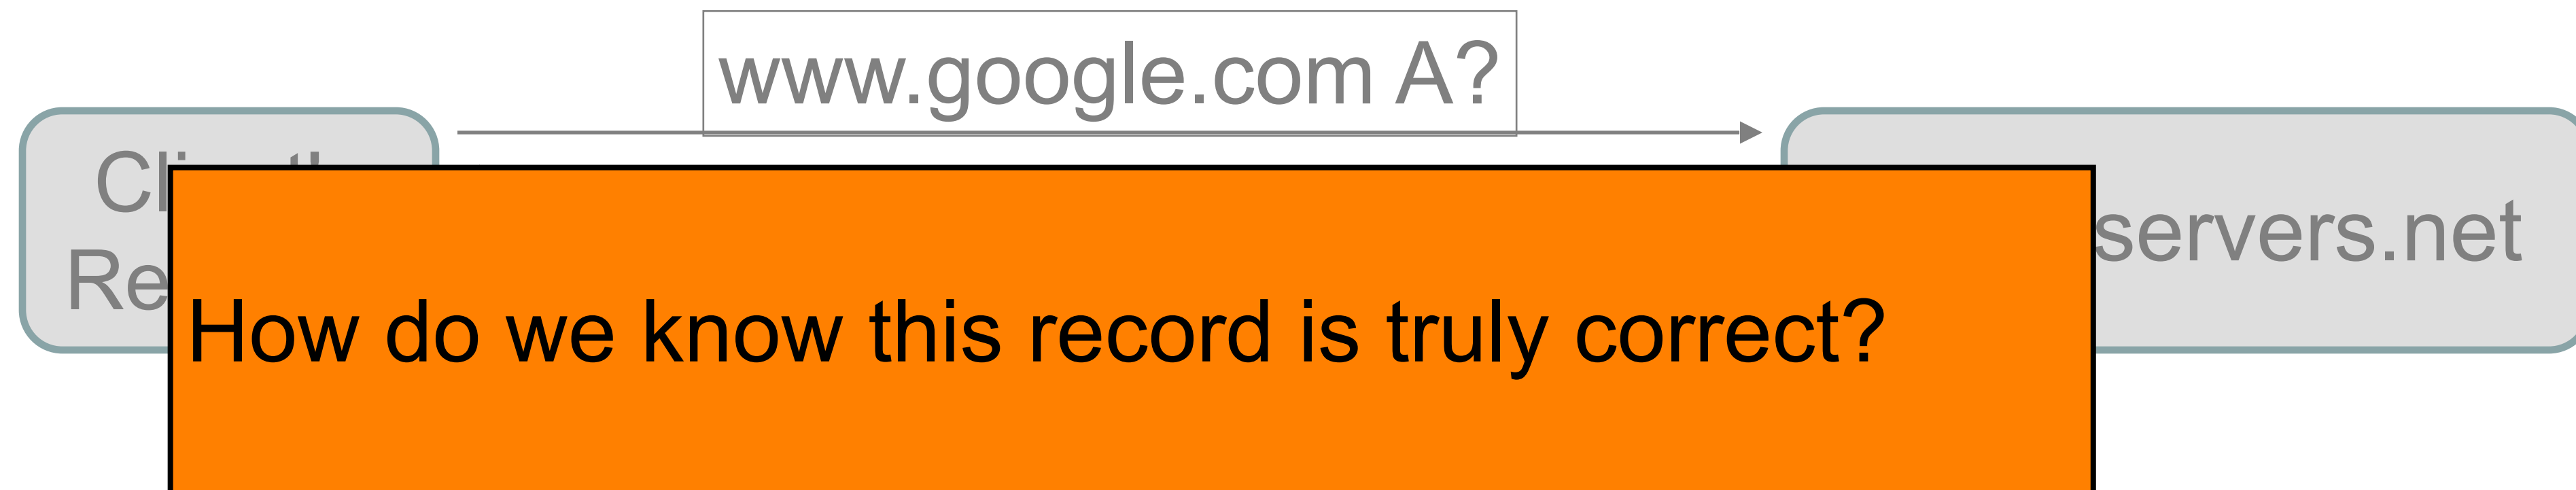
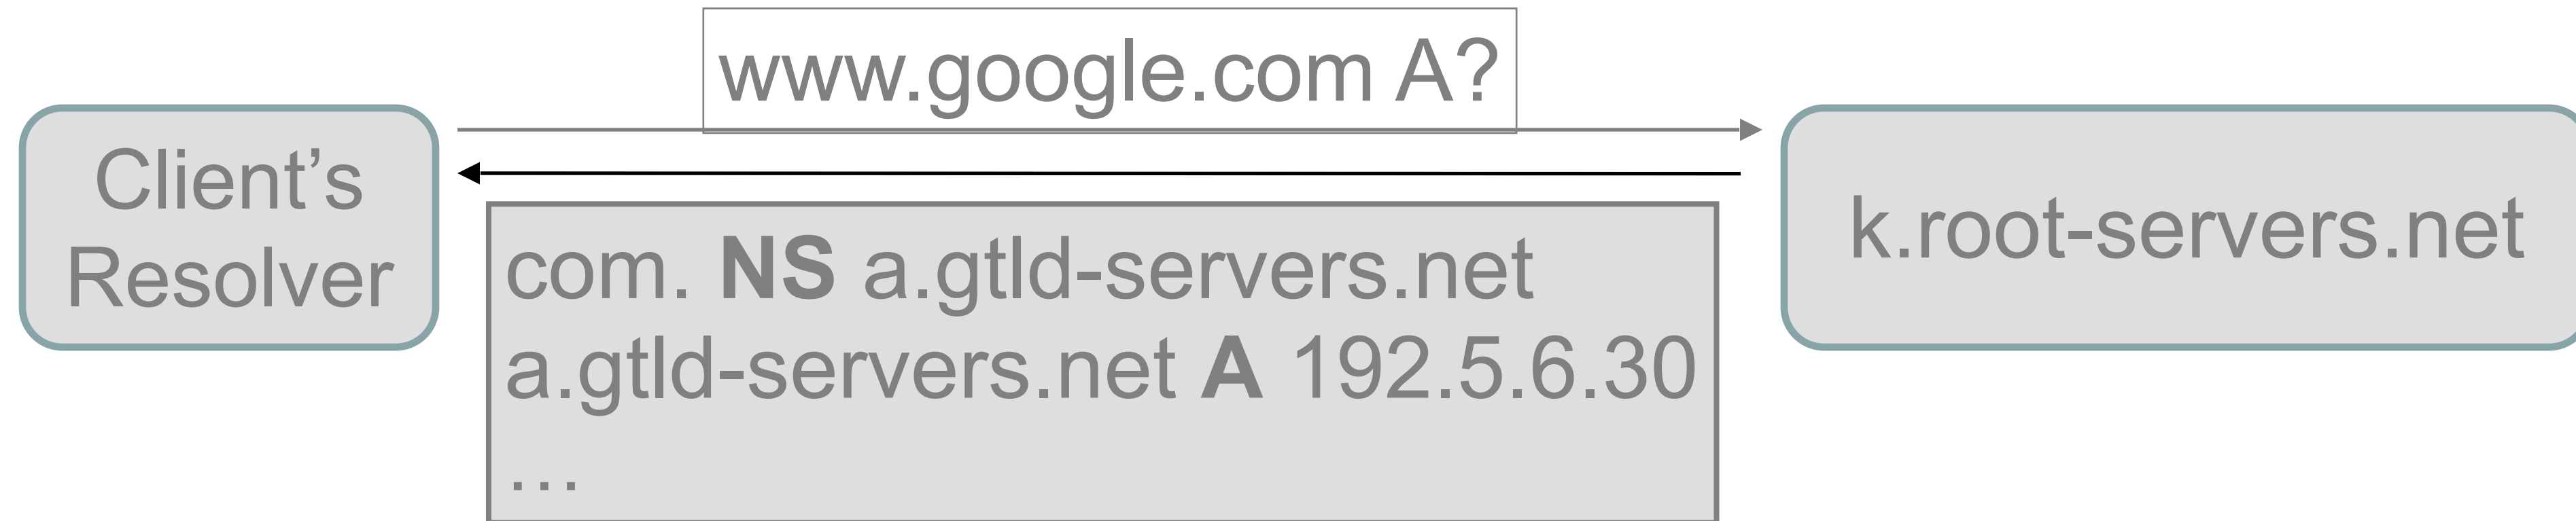
Ordinary DNS:



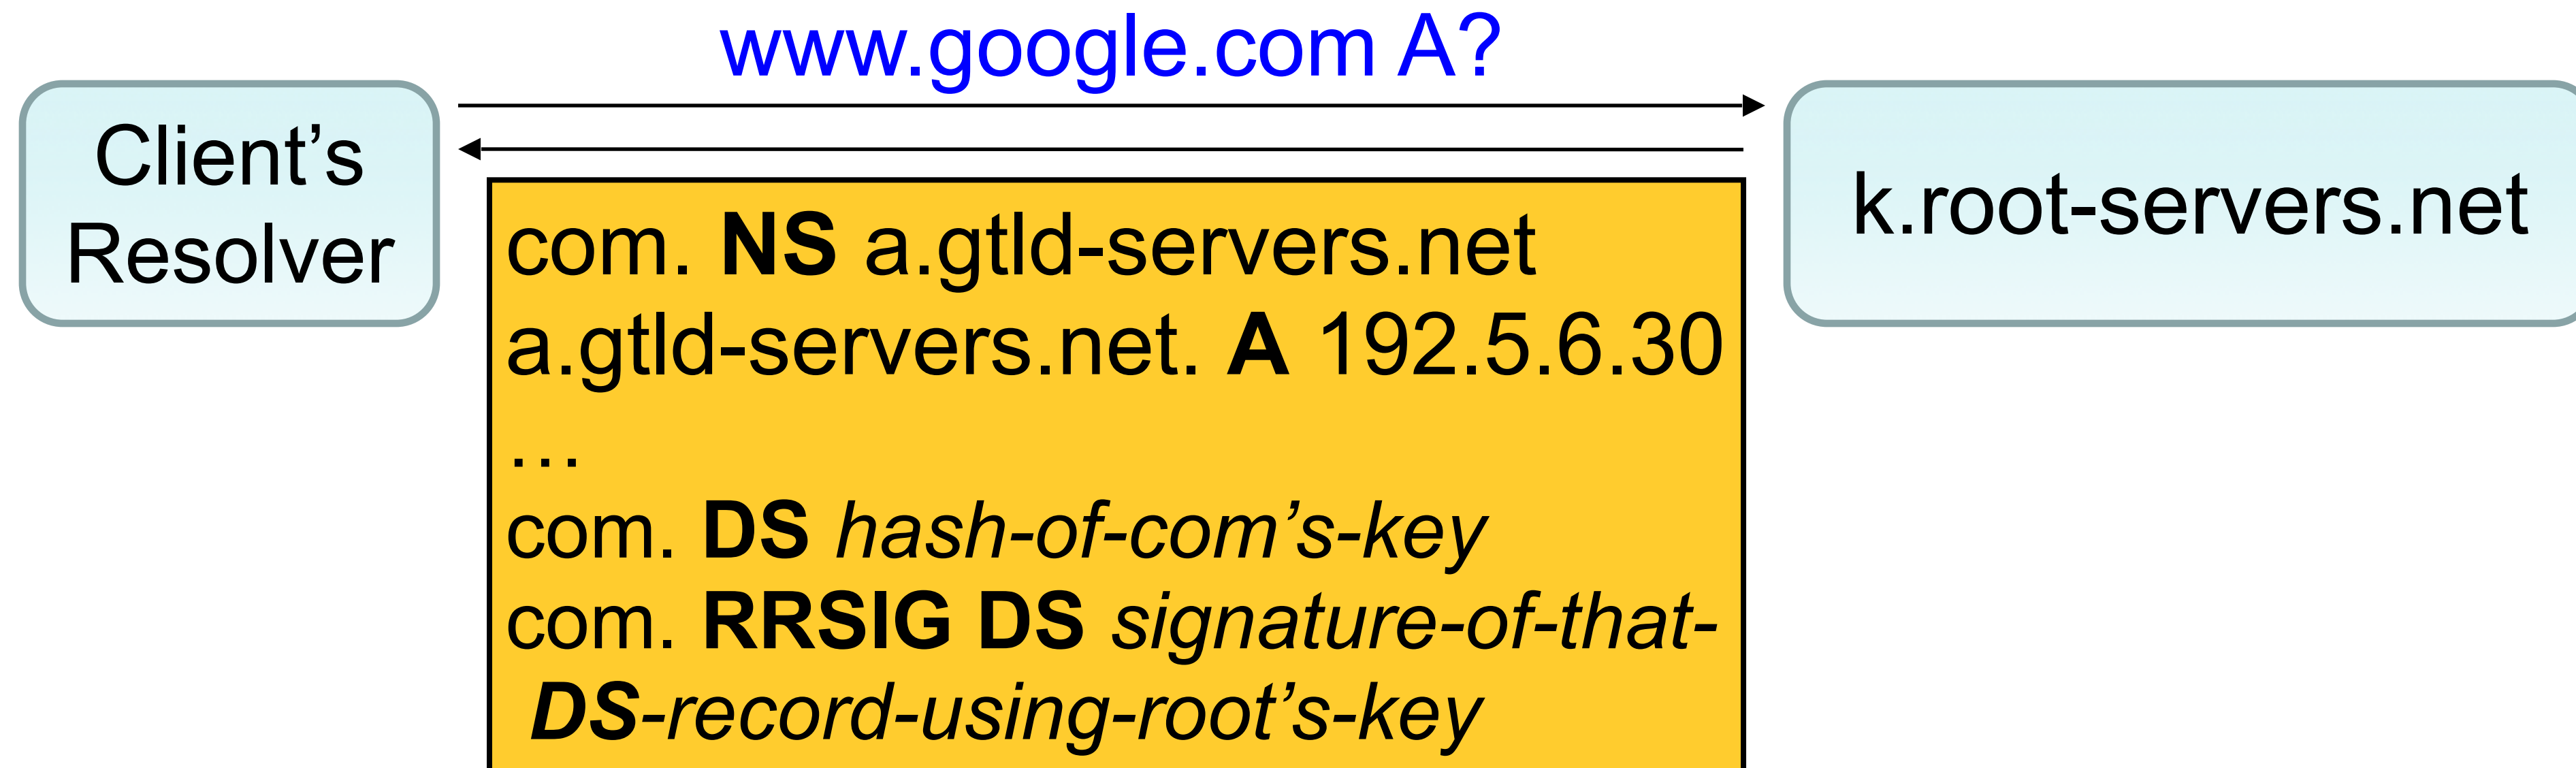
Ordinary DNS:



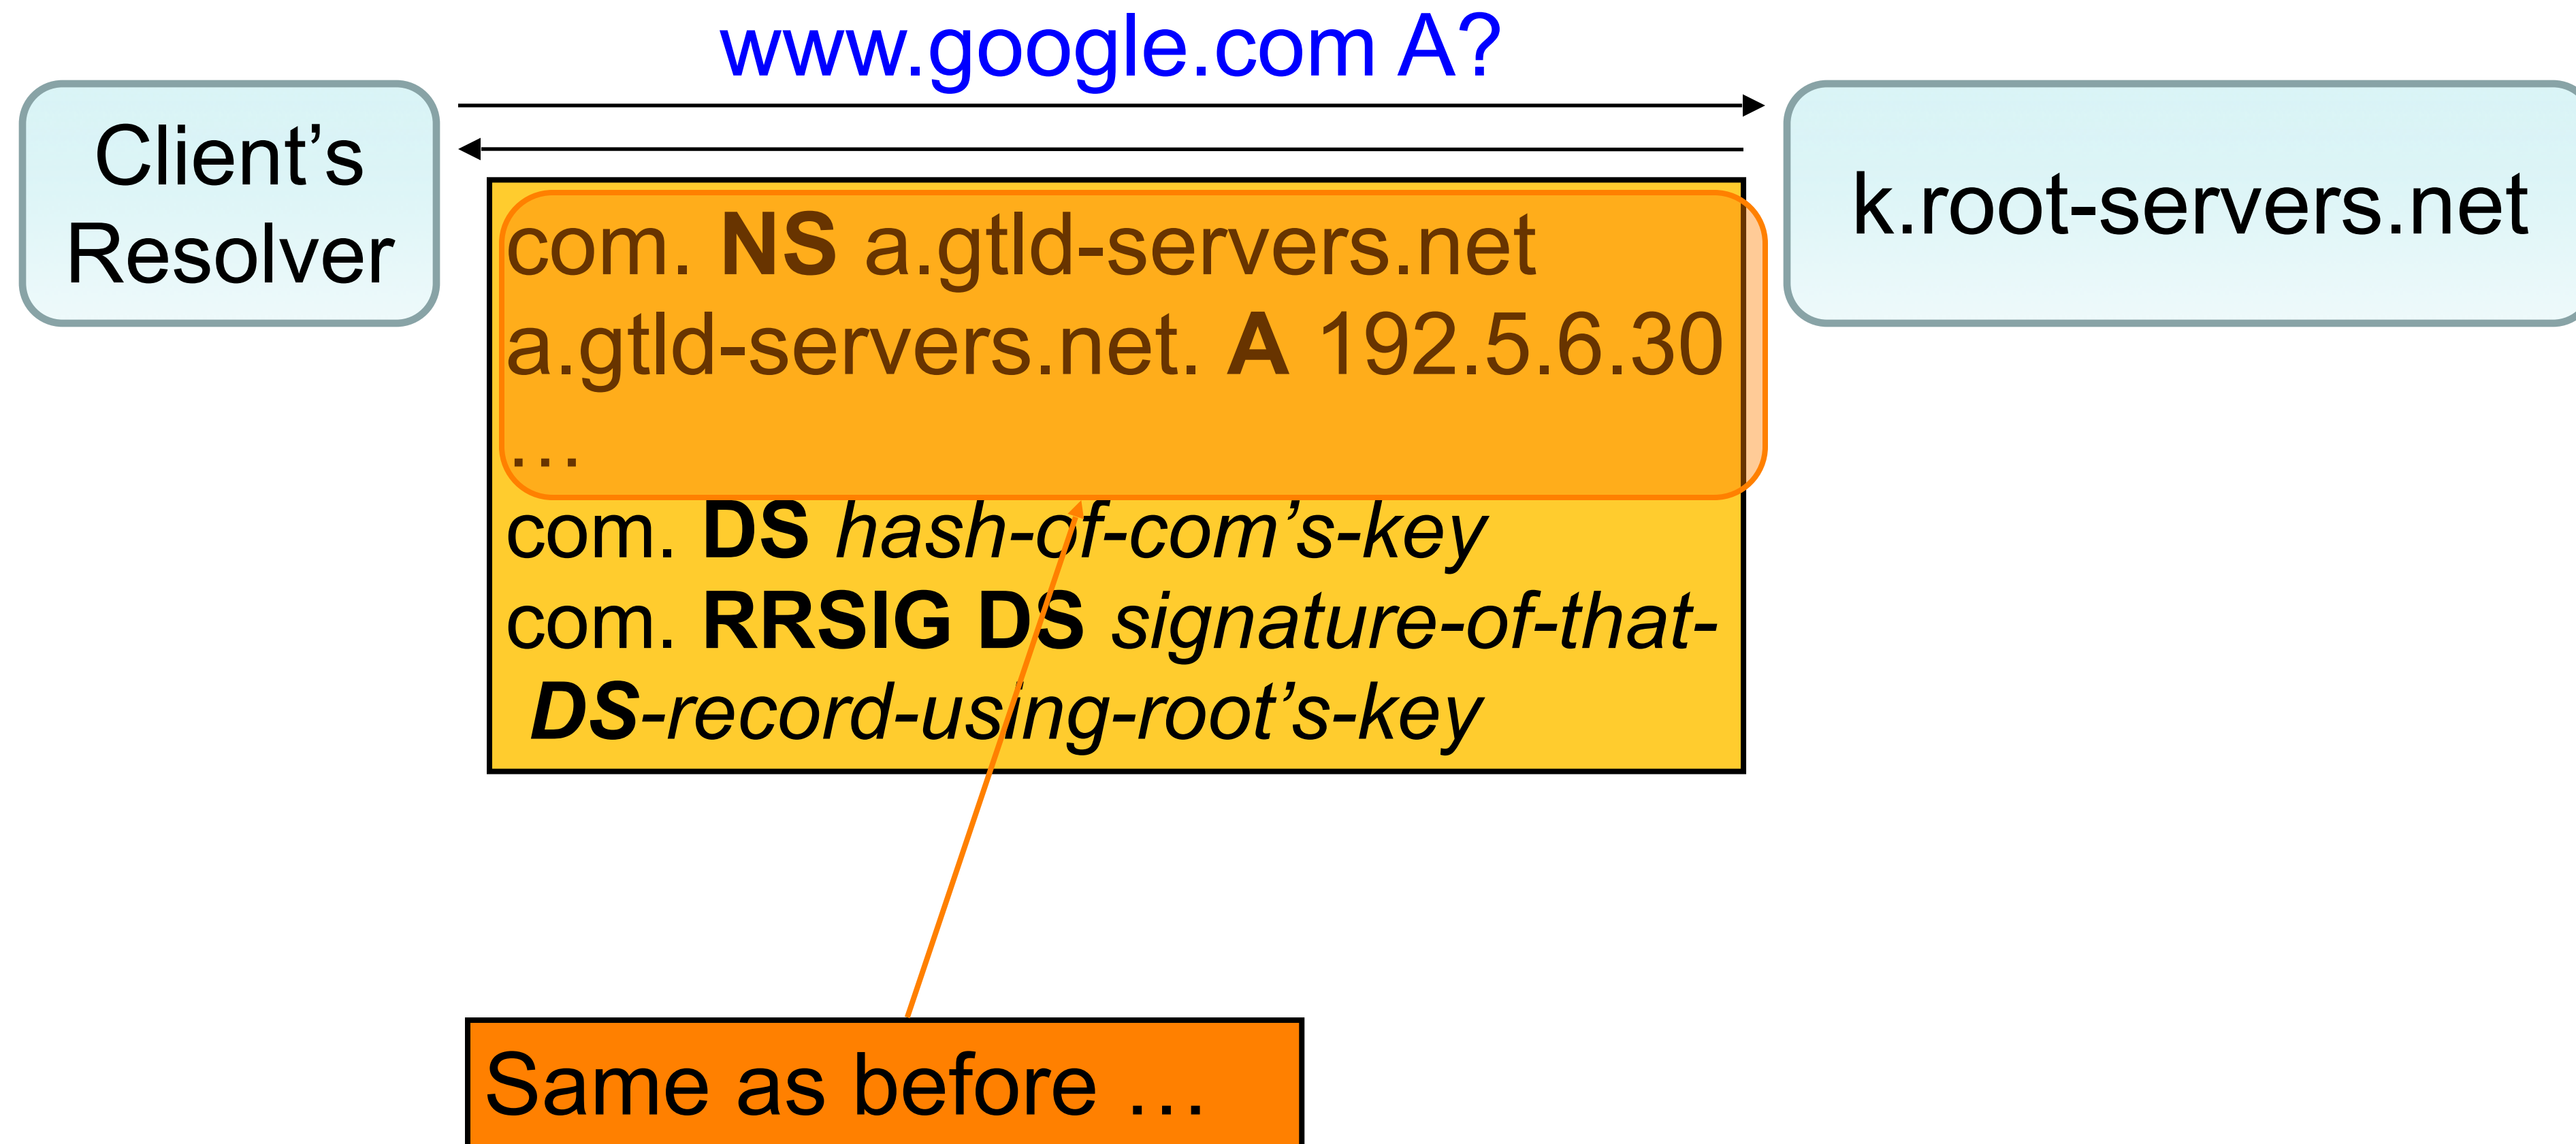
Ordinary DNS:



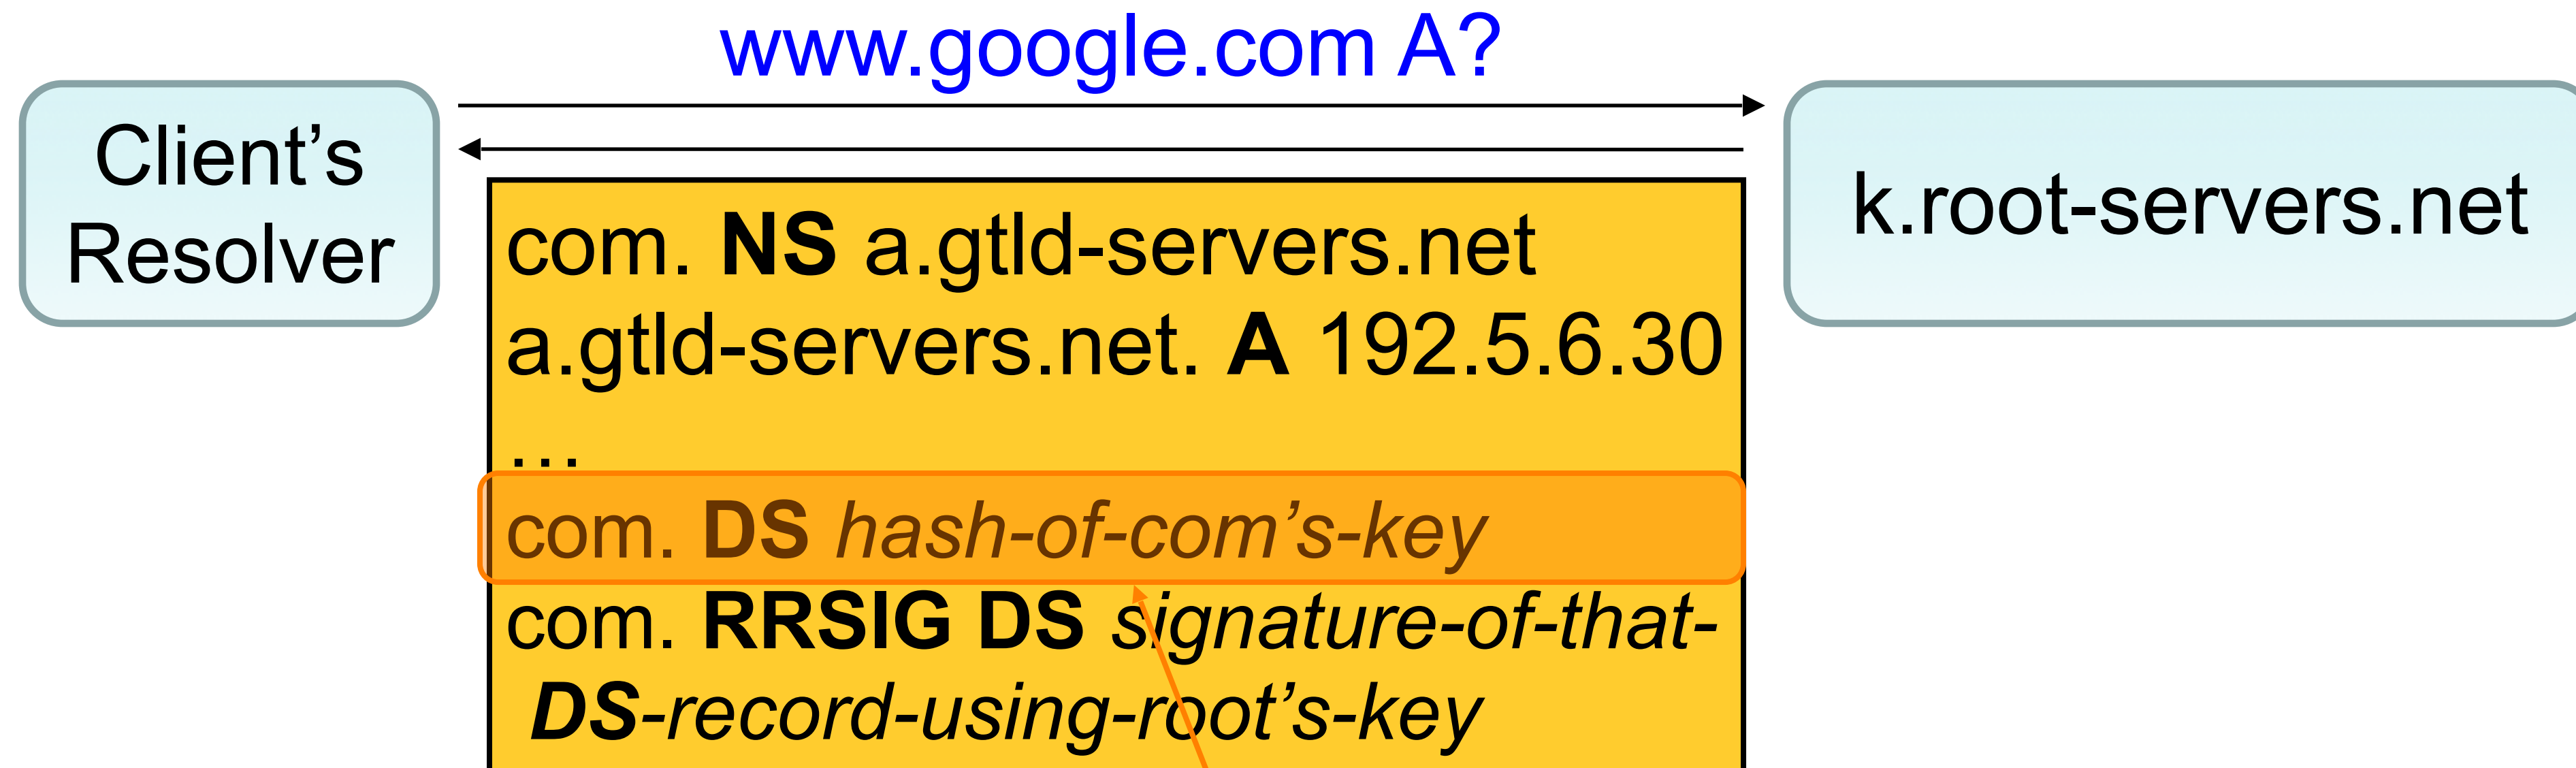
DNSSEC (with simplifications):



DNSSEC (with simplifications):

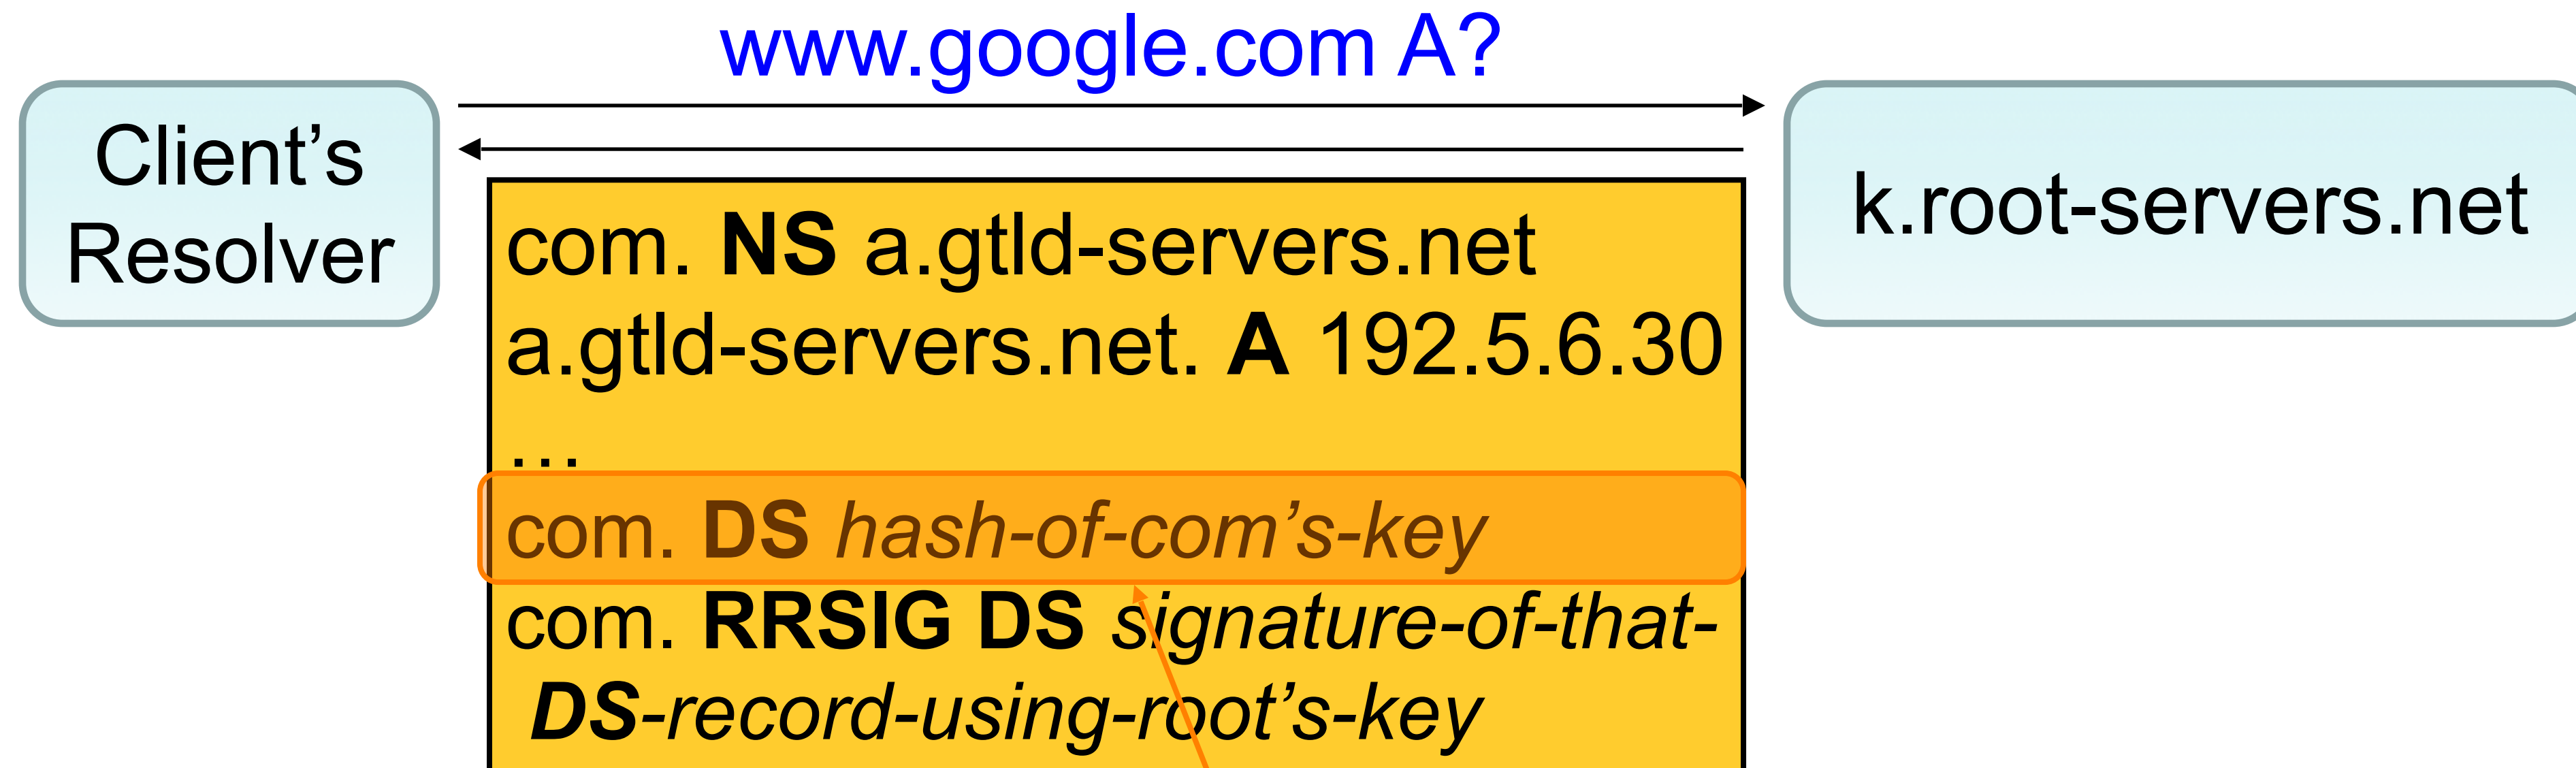


DNSSEC (with simplifications):



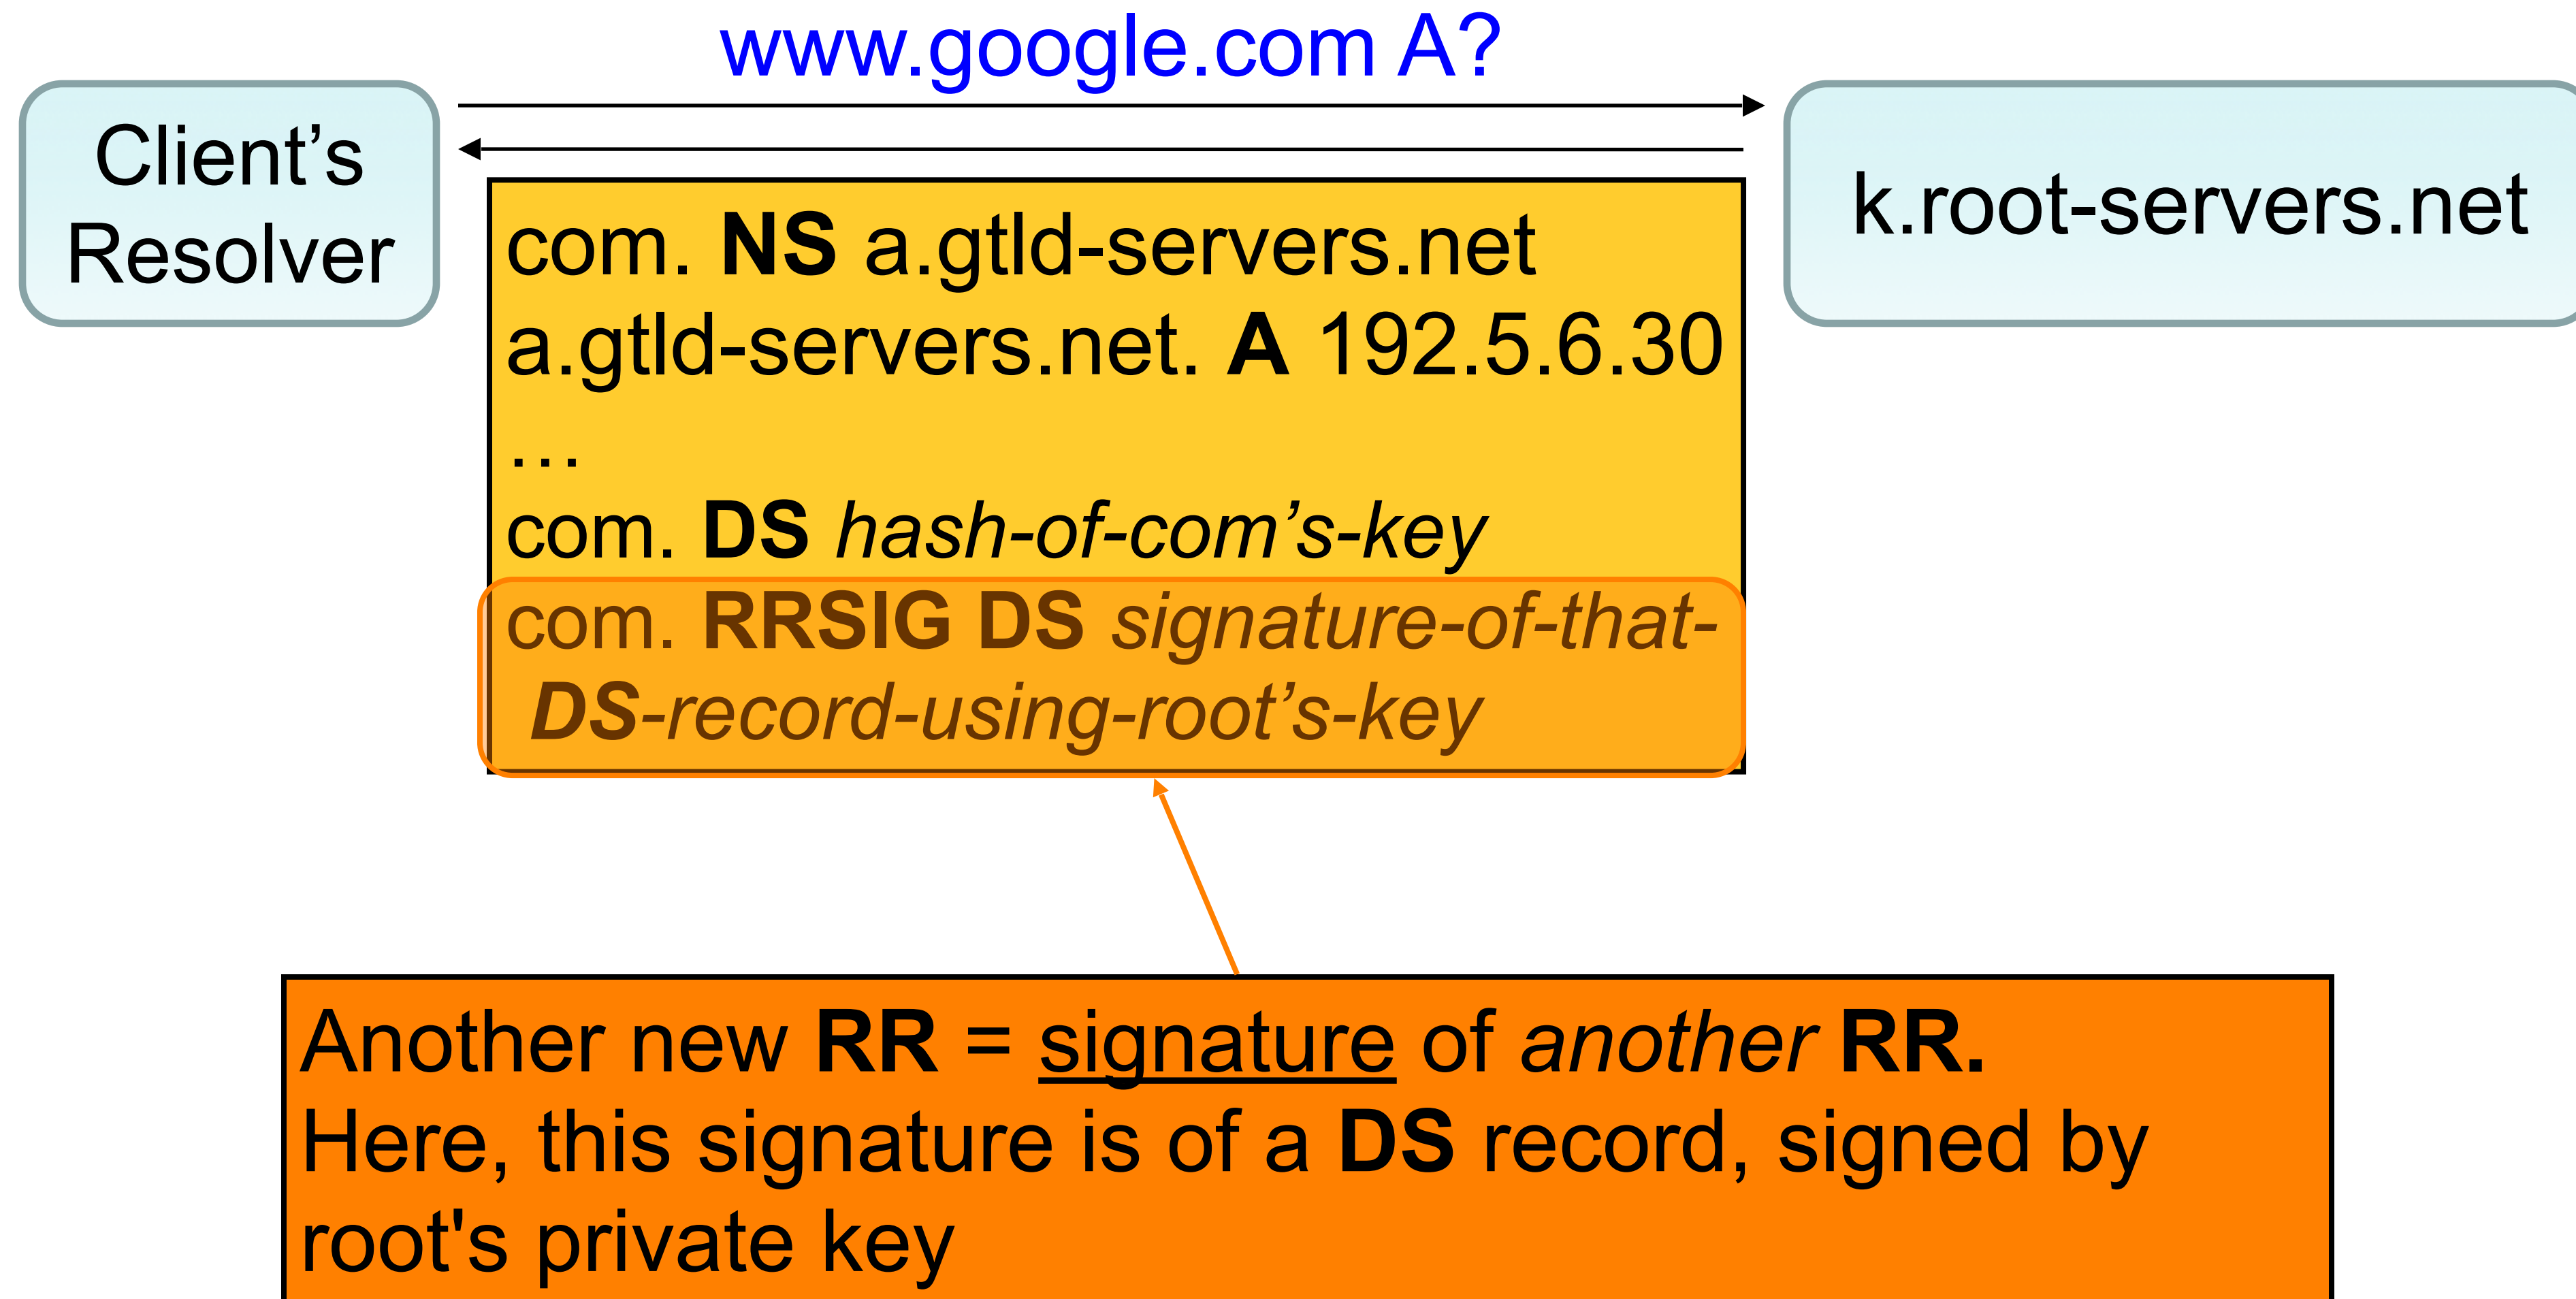
New **RR** ("Delegation Signer") tells us if we have correct copy of .com's public key (by comparing hash values)

DNSSEC (with simplifications):

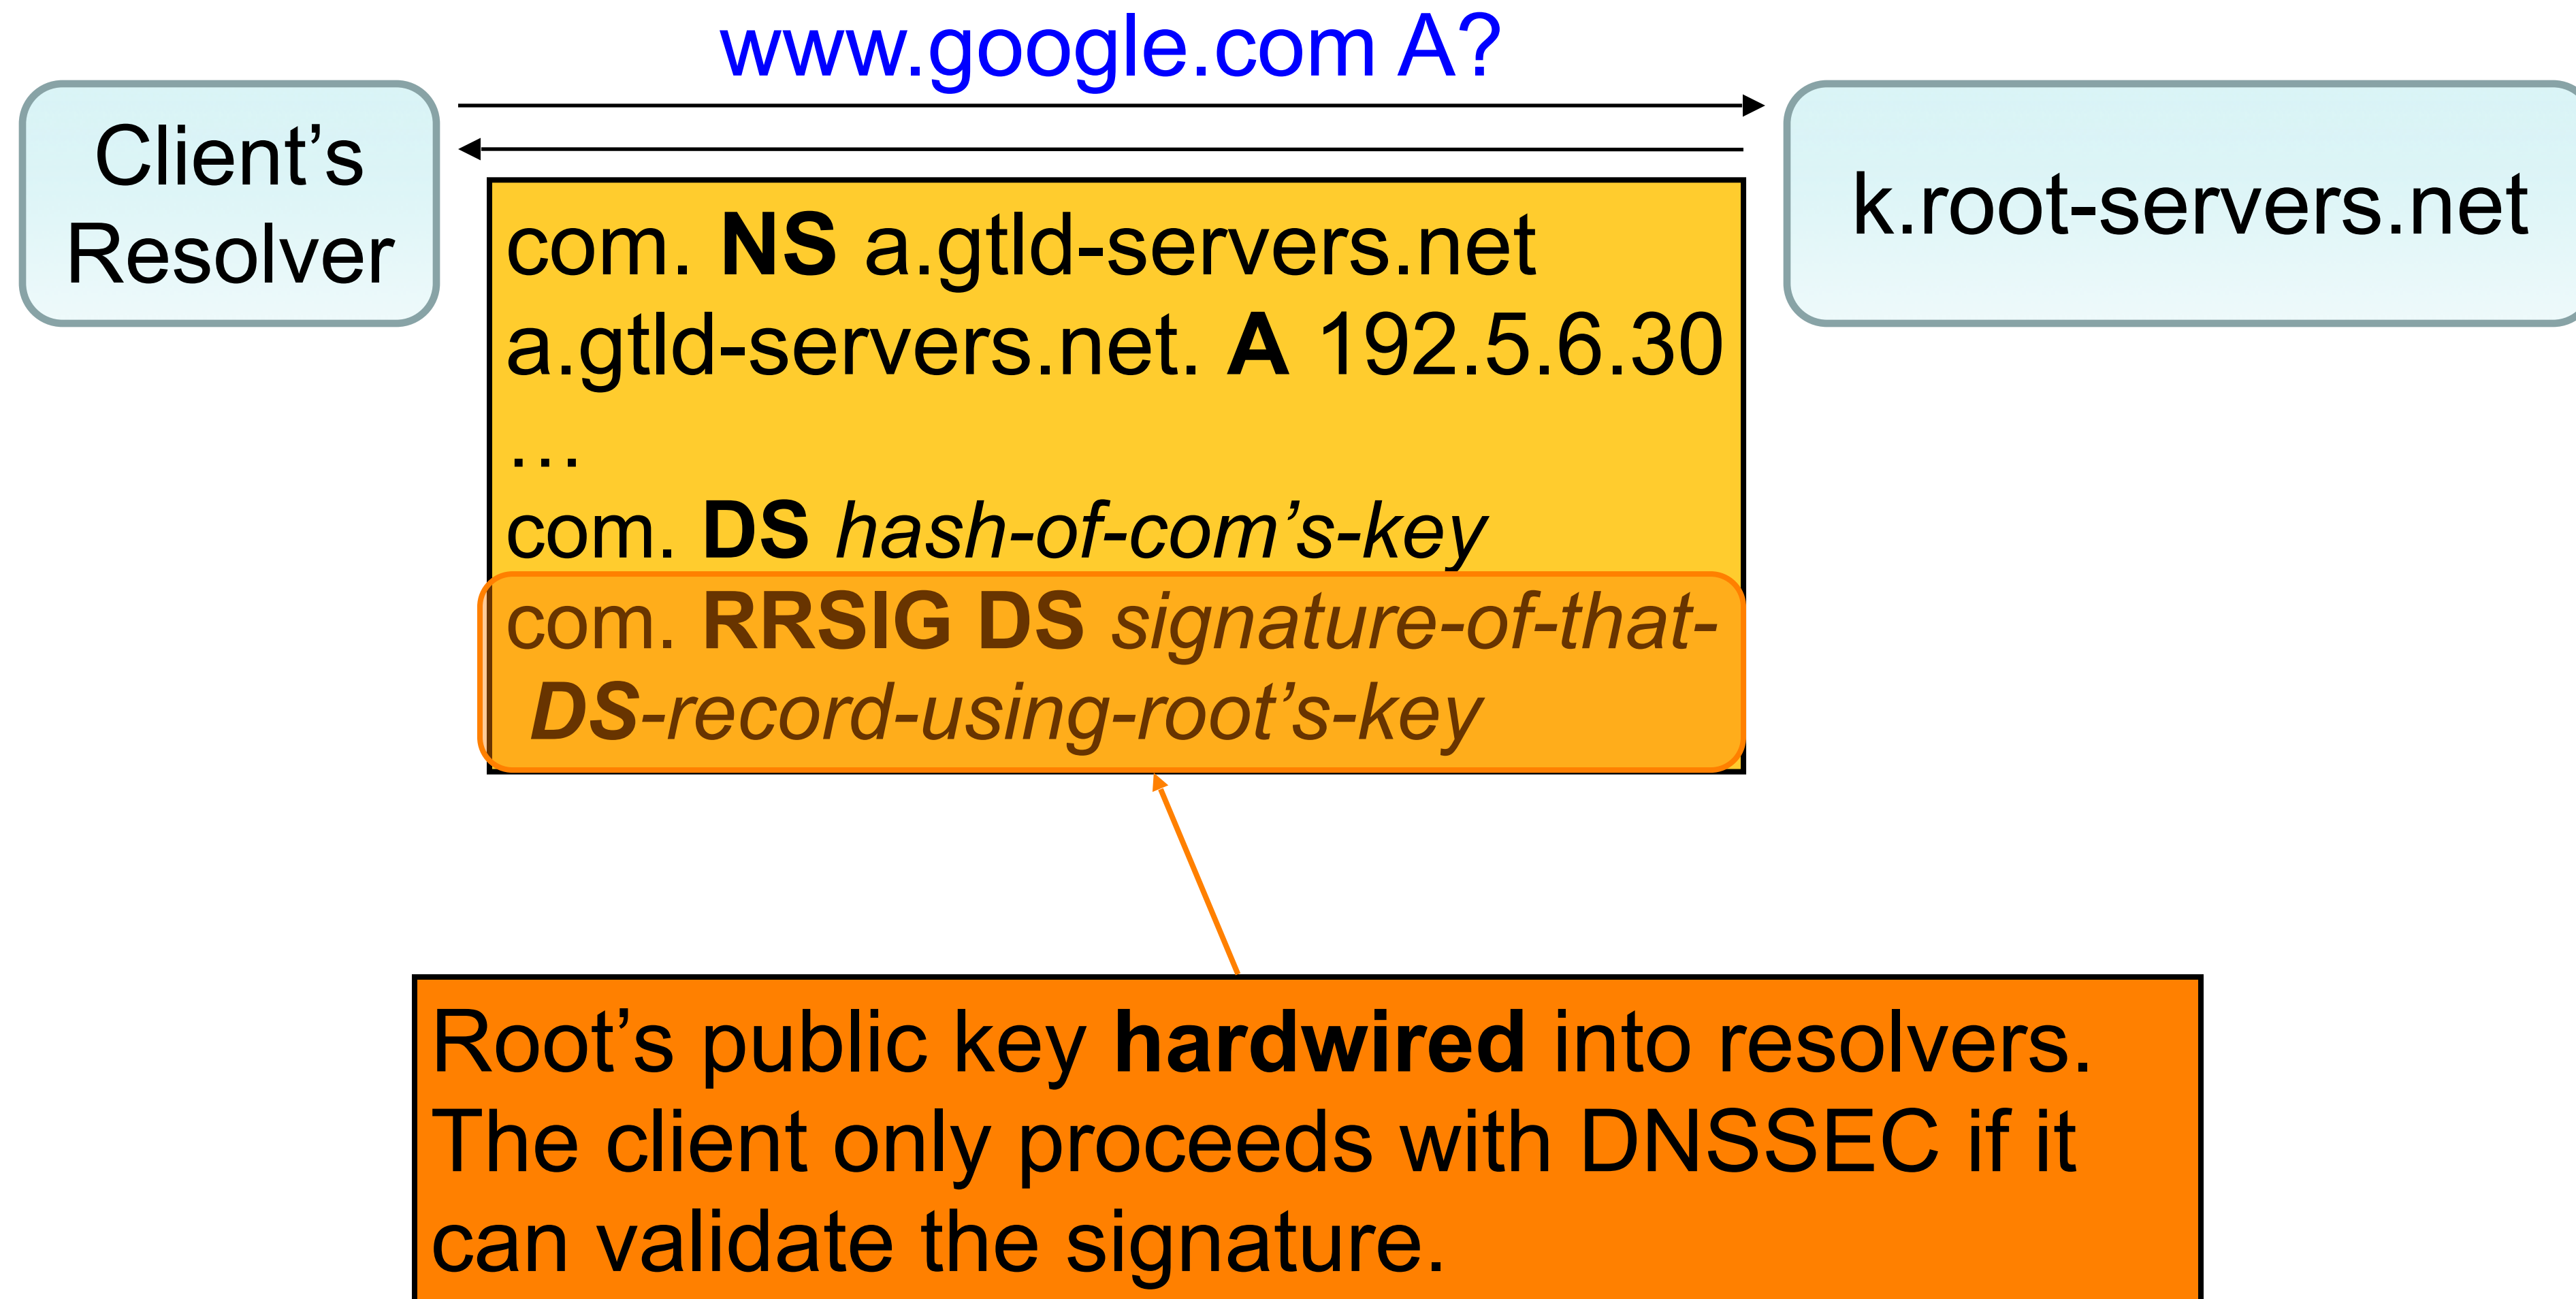


Getting .com NS's key is a bit complicated...we'll talk about in a bit. Assume we can get it for now...

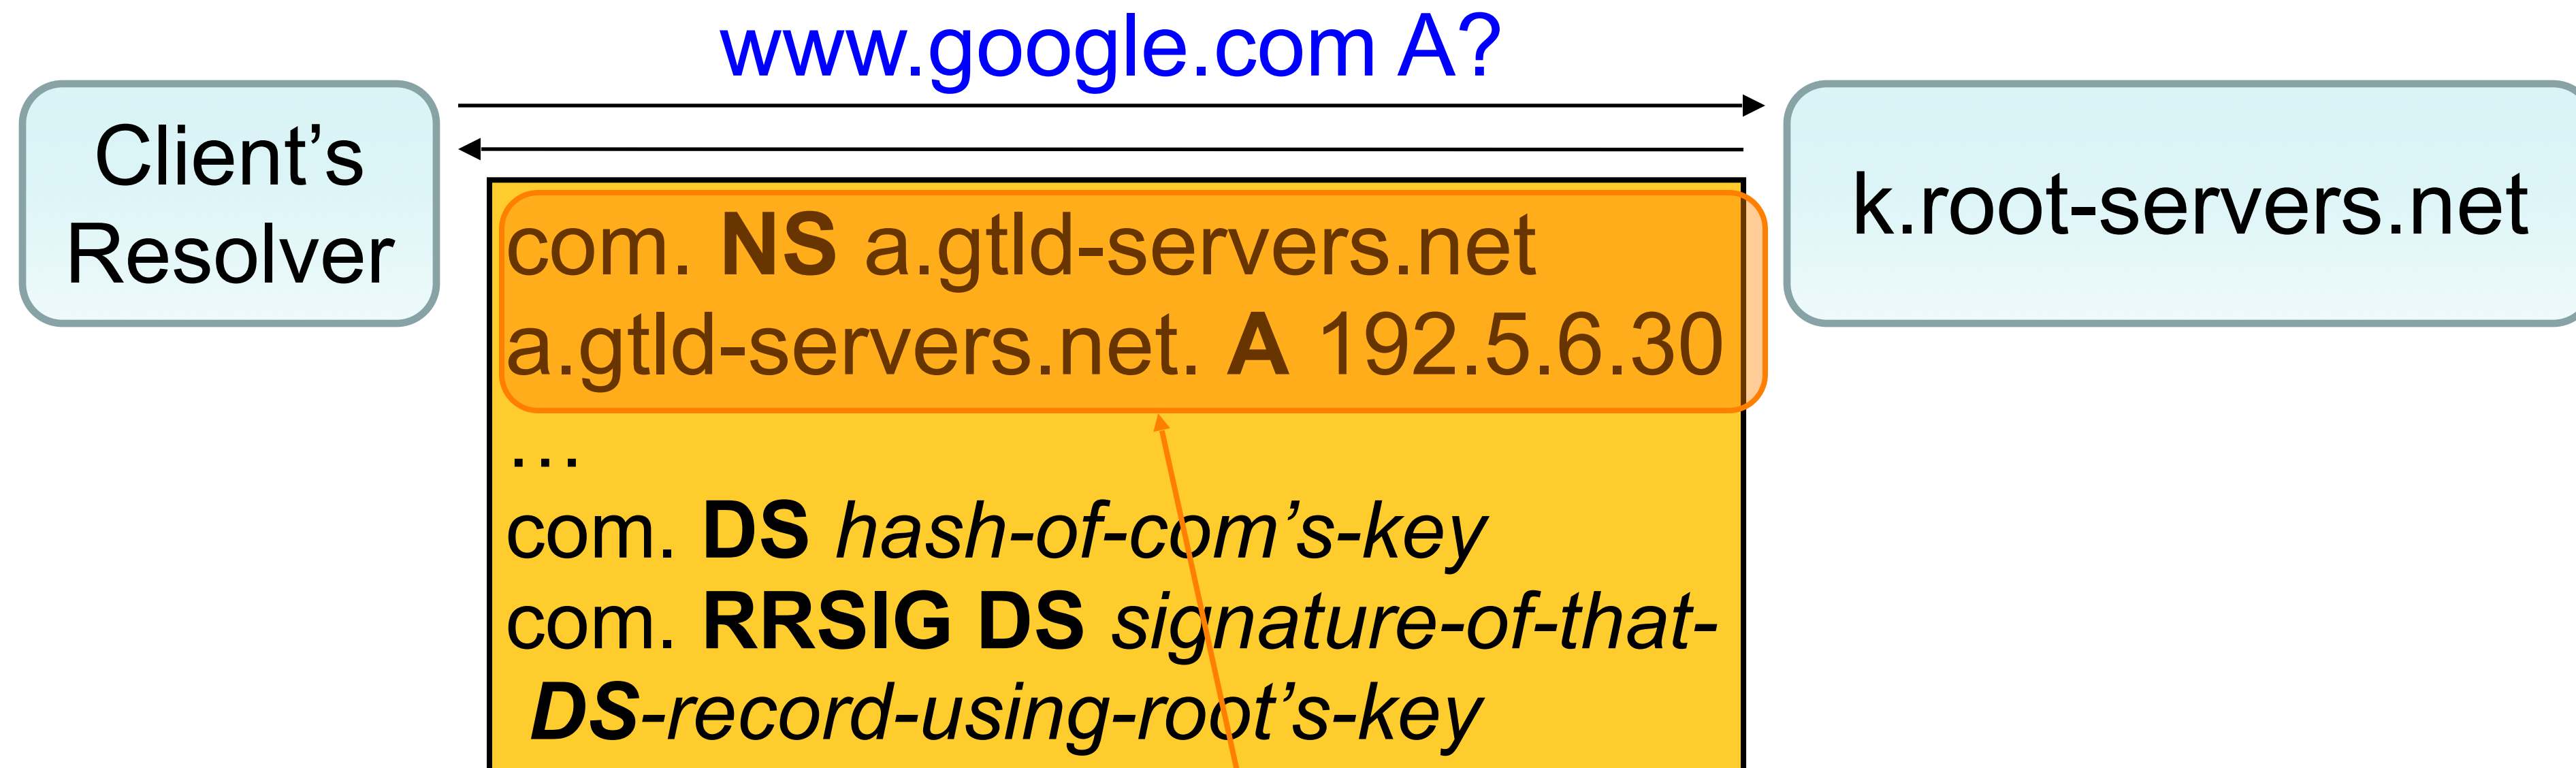
DNSSEC (with simplifications):



DNSSEC (with simplifications):



DNSSEC (with simplifications):

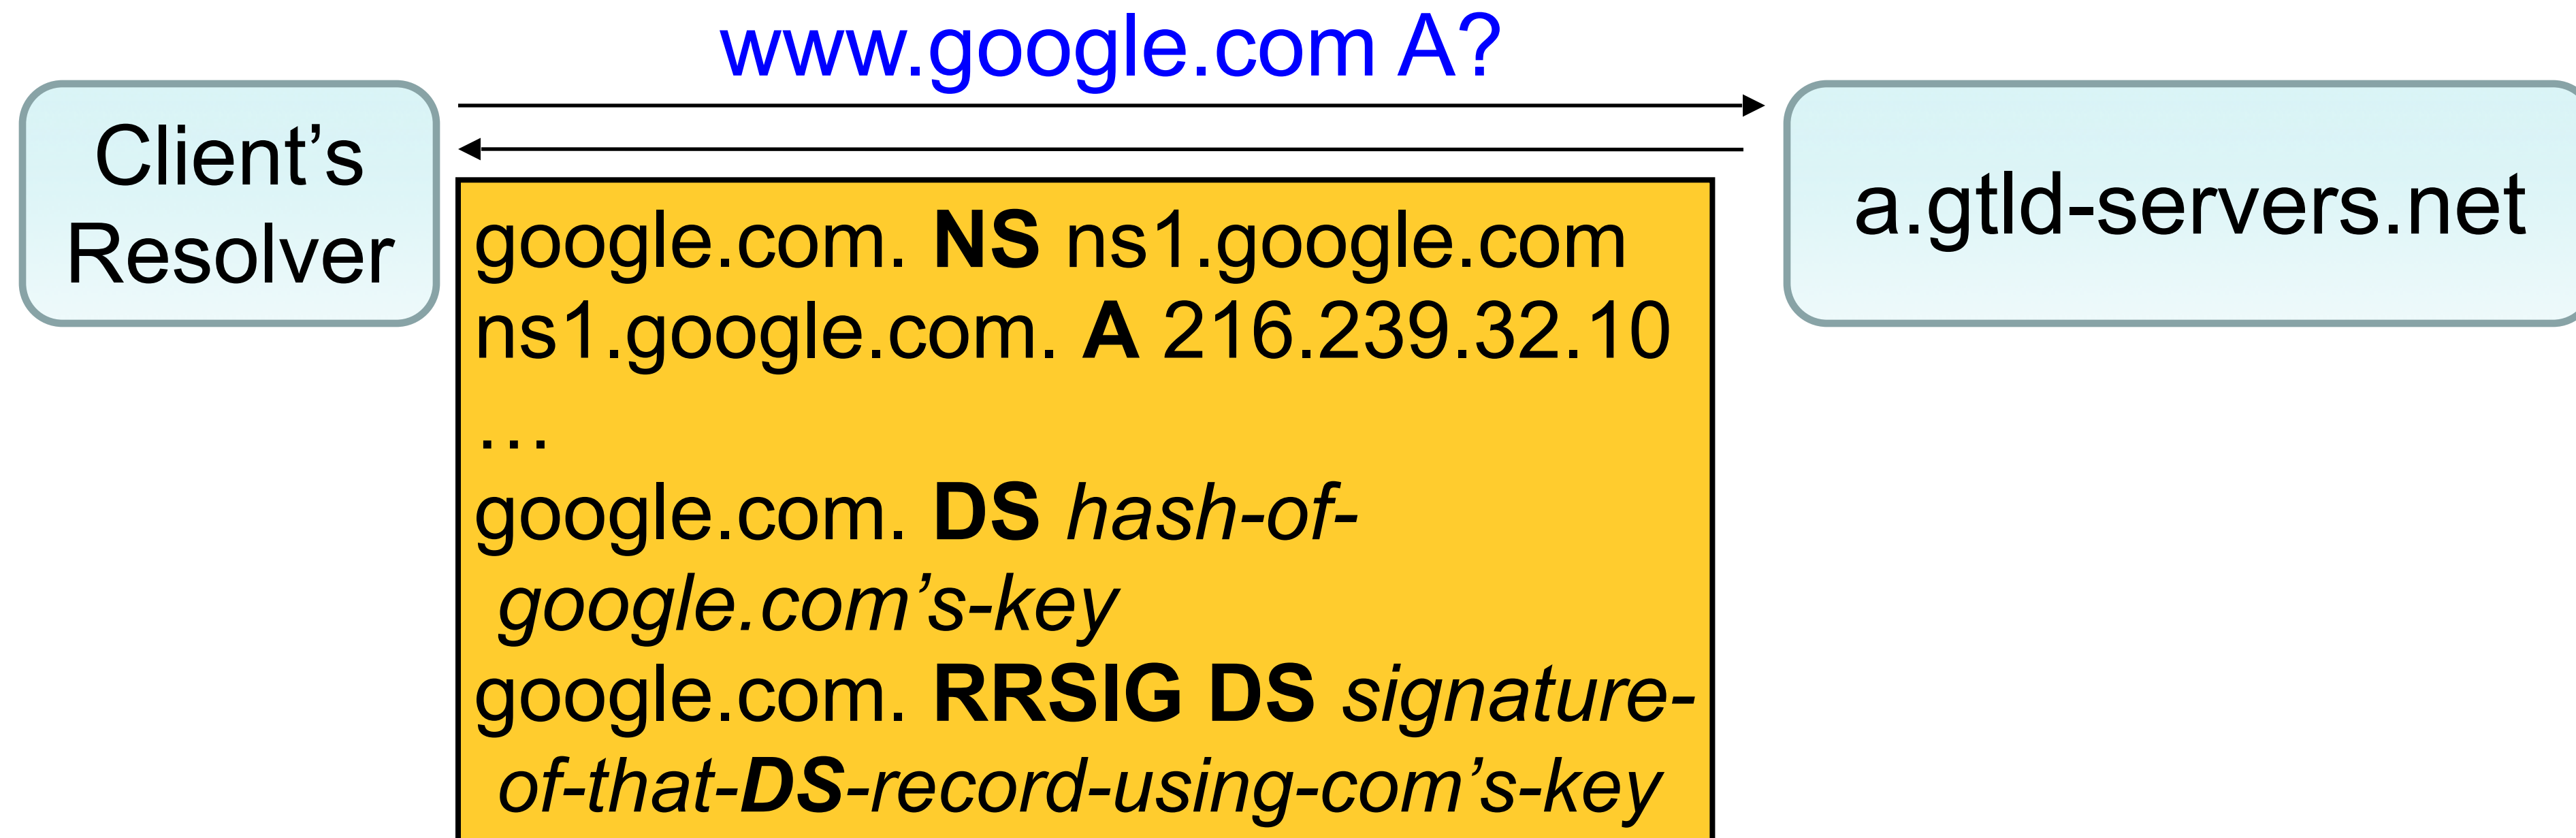


Note: there's no signature over the **NS** or **A** information! If DNS data tampered with, will find out later.

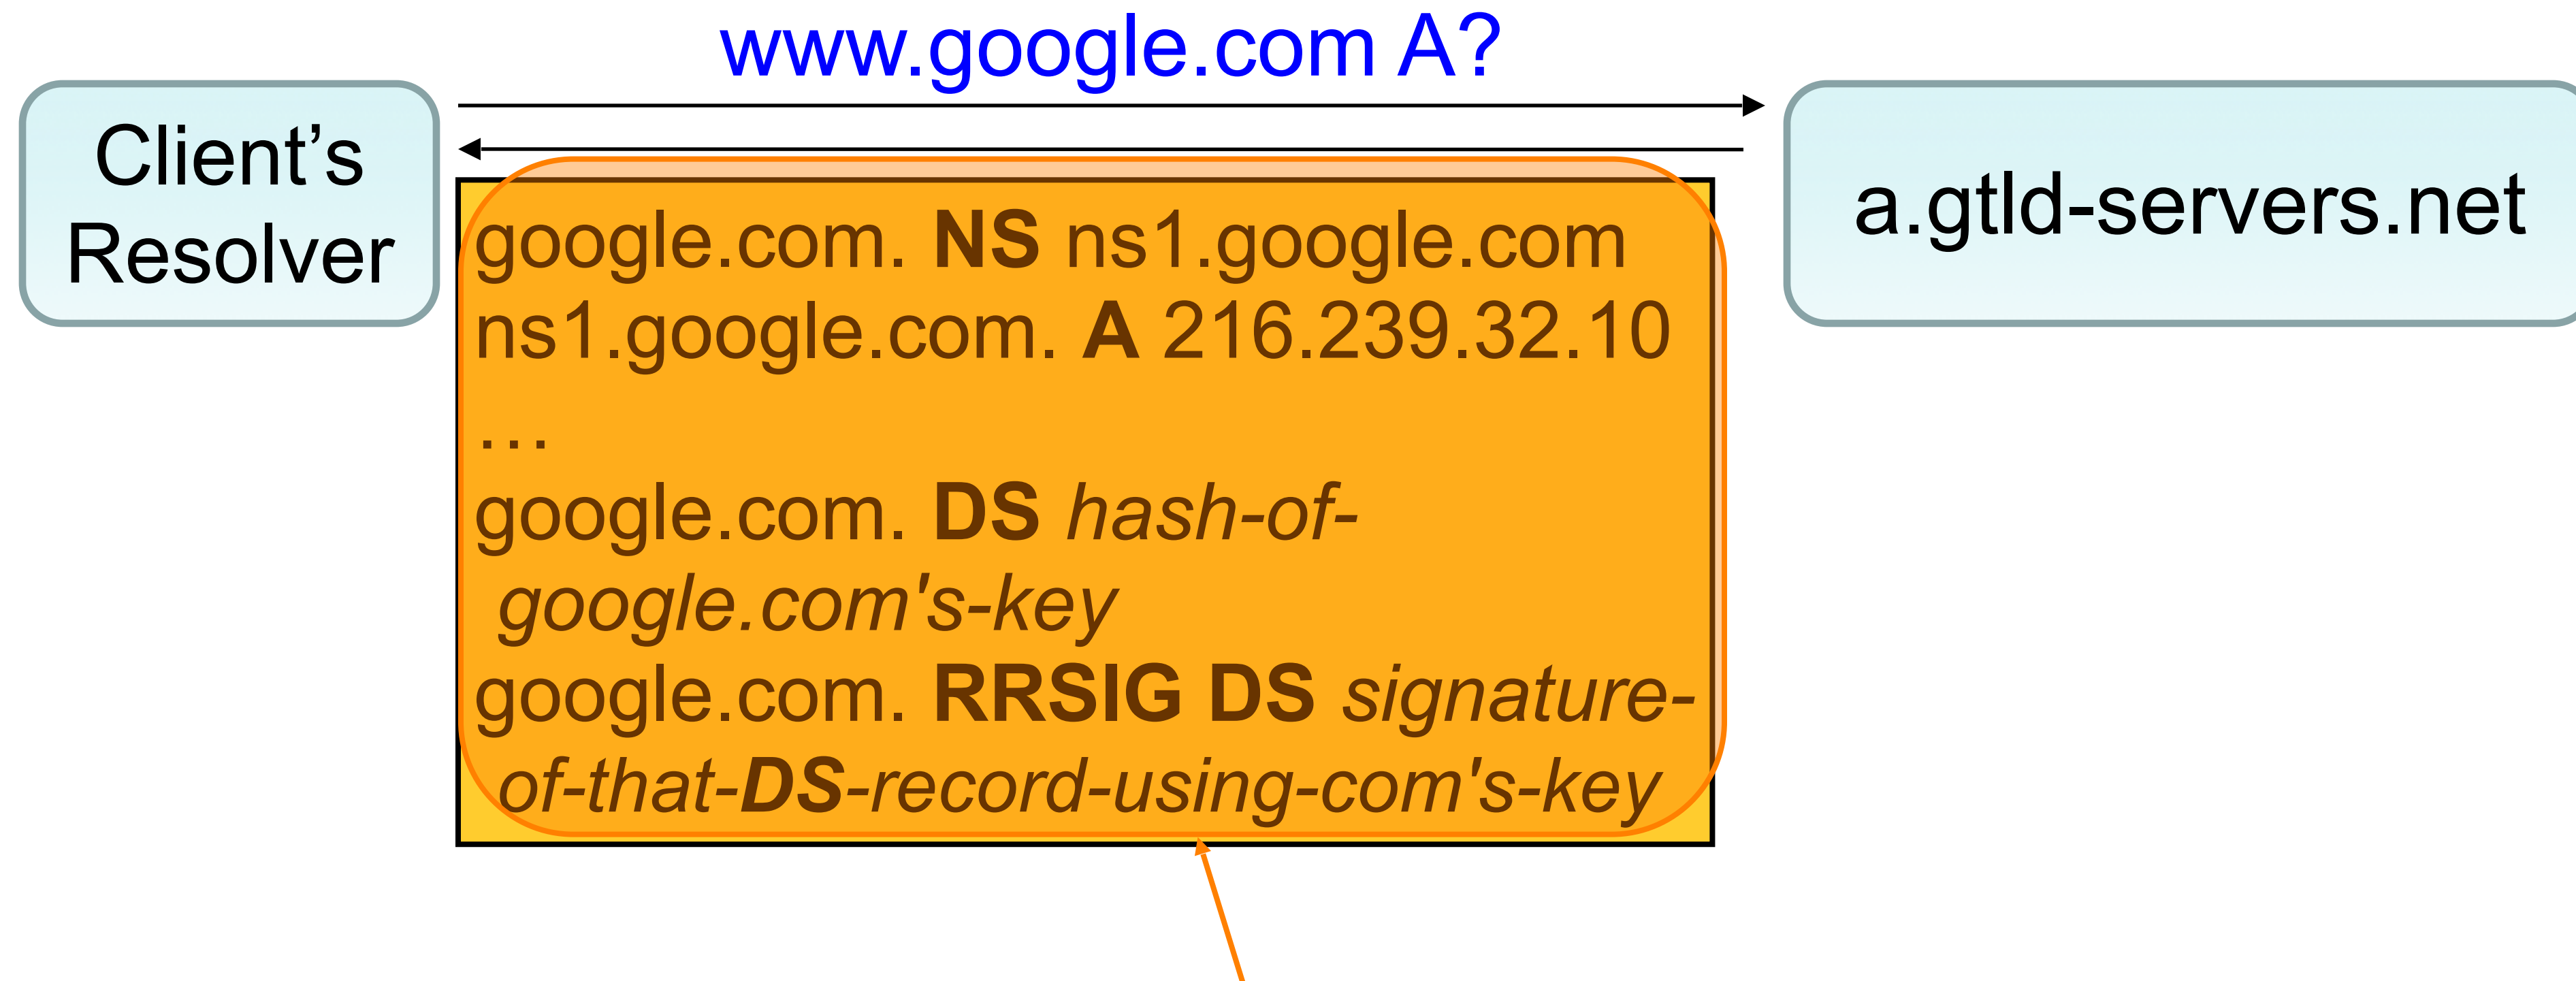
DNSSEC (with simplifications):



DNSSEC (with simplifications):



DNSSEC (with simplifications):



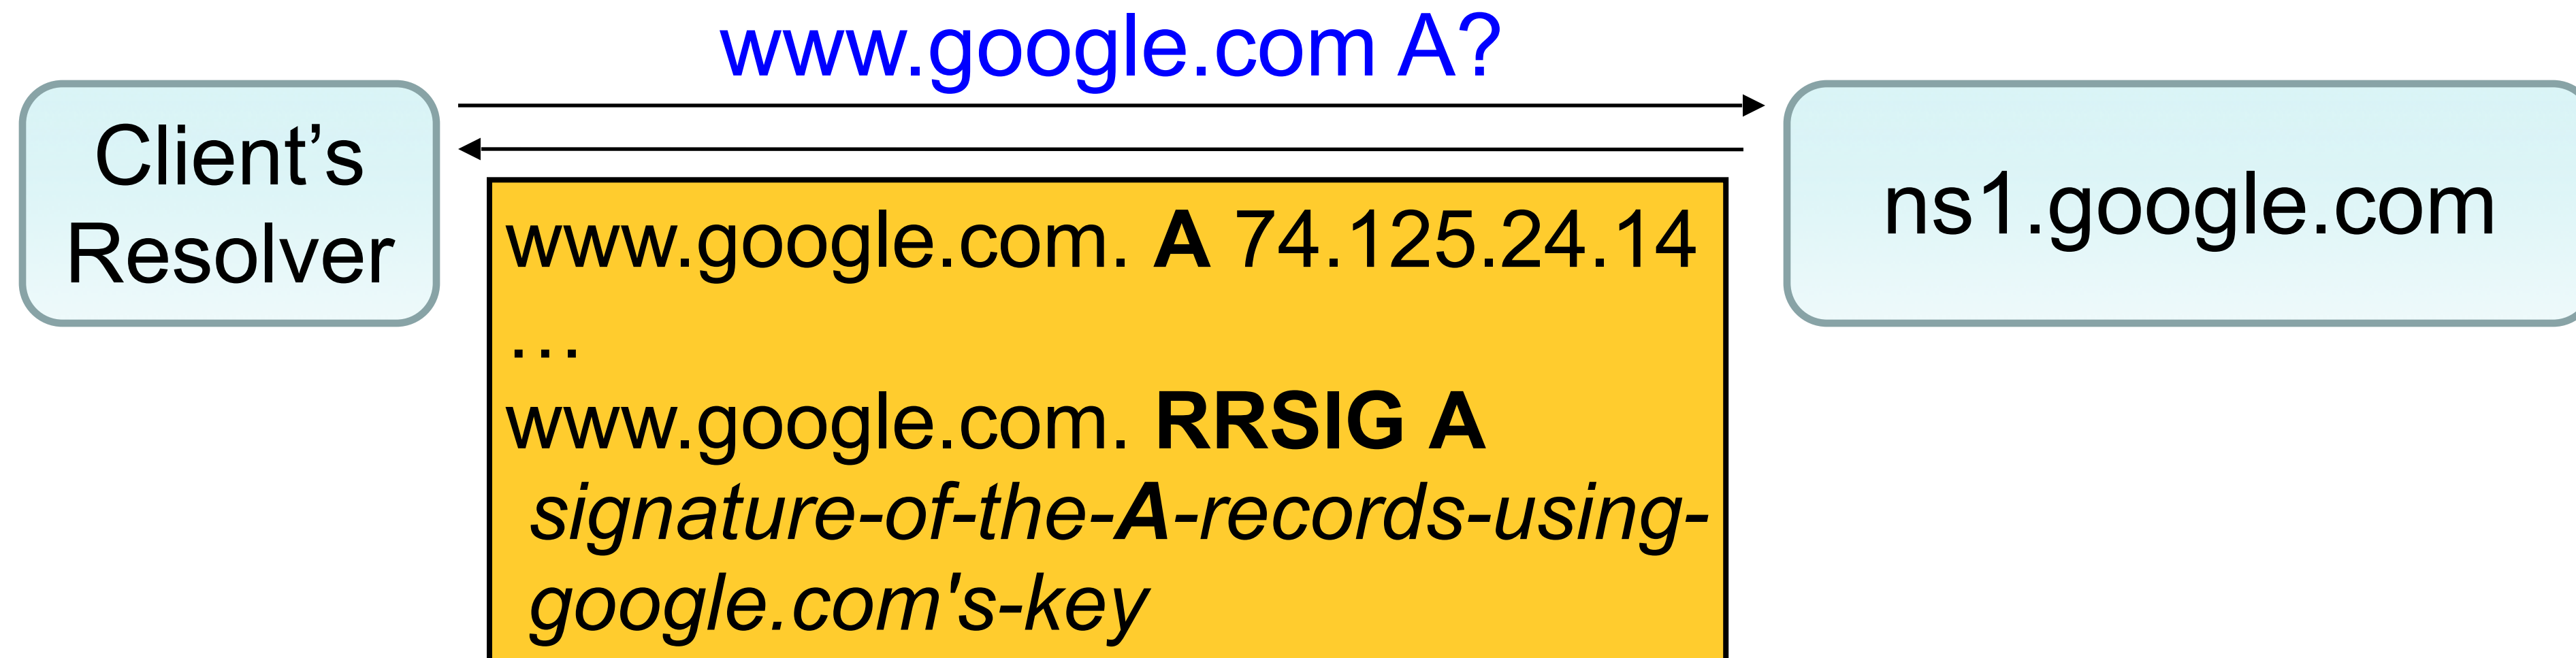
Similar as before:

- Identify google.com's public key
- DS signed by .com NS's key

DNSSEC (with simplifications):



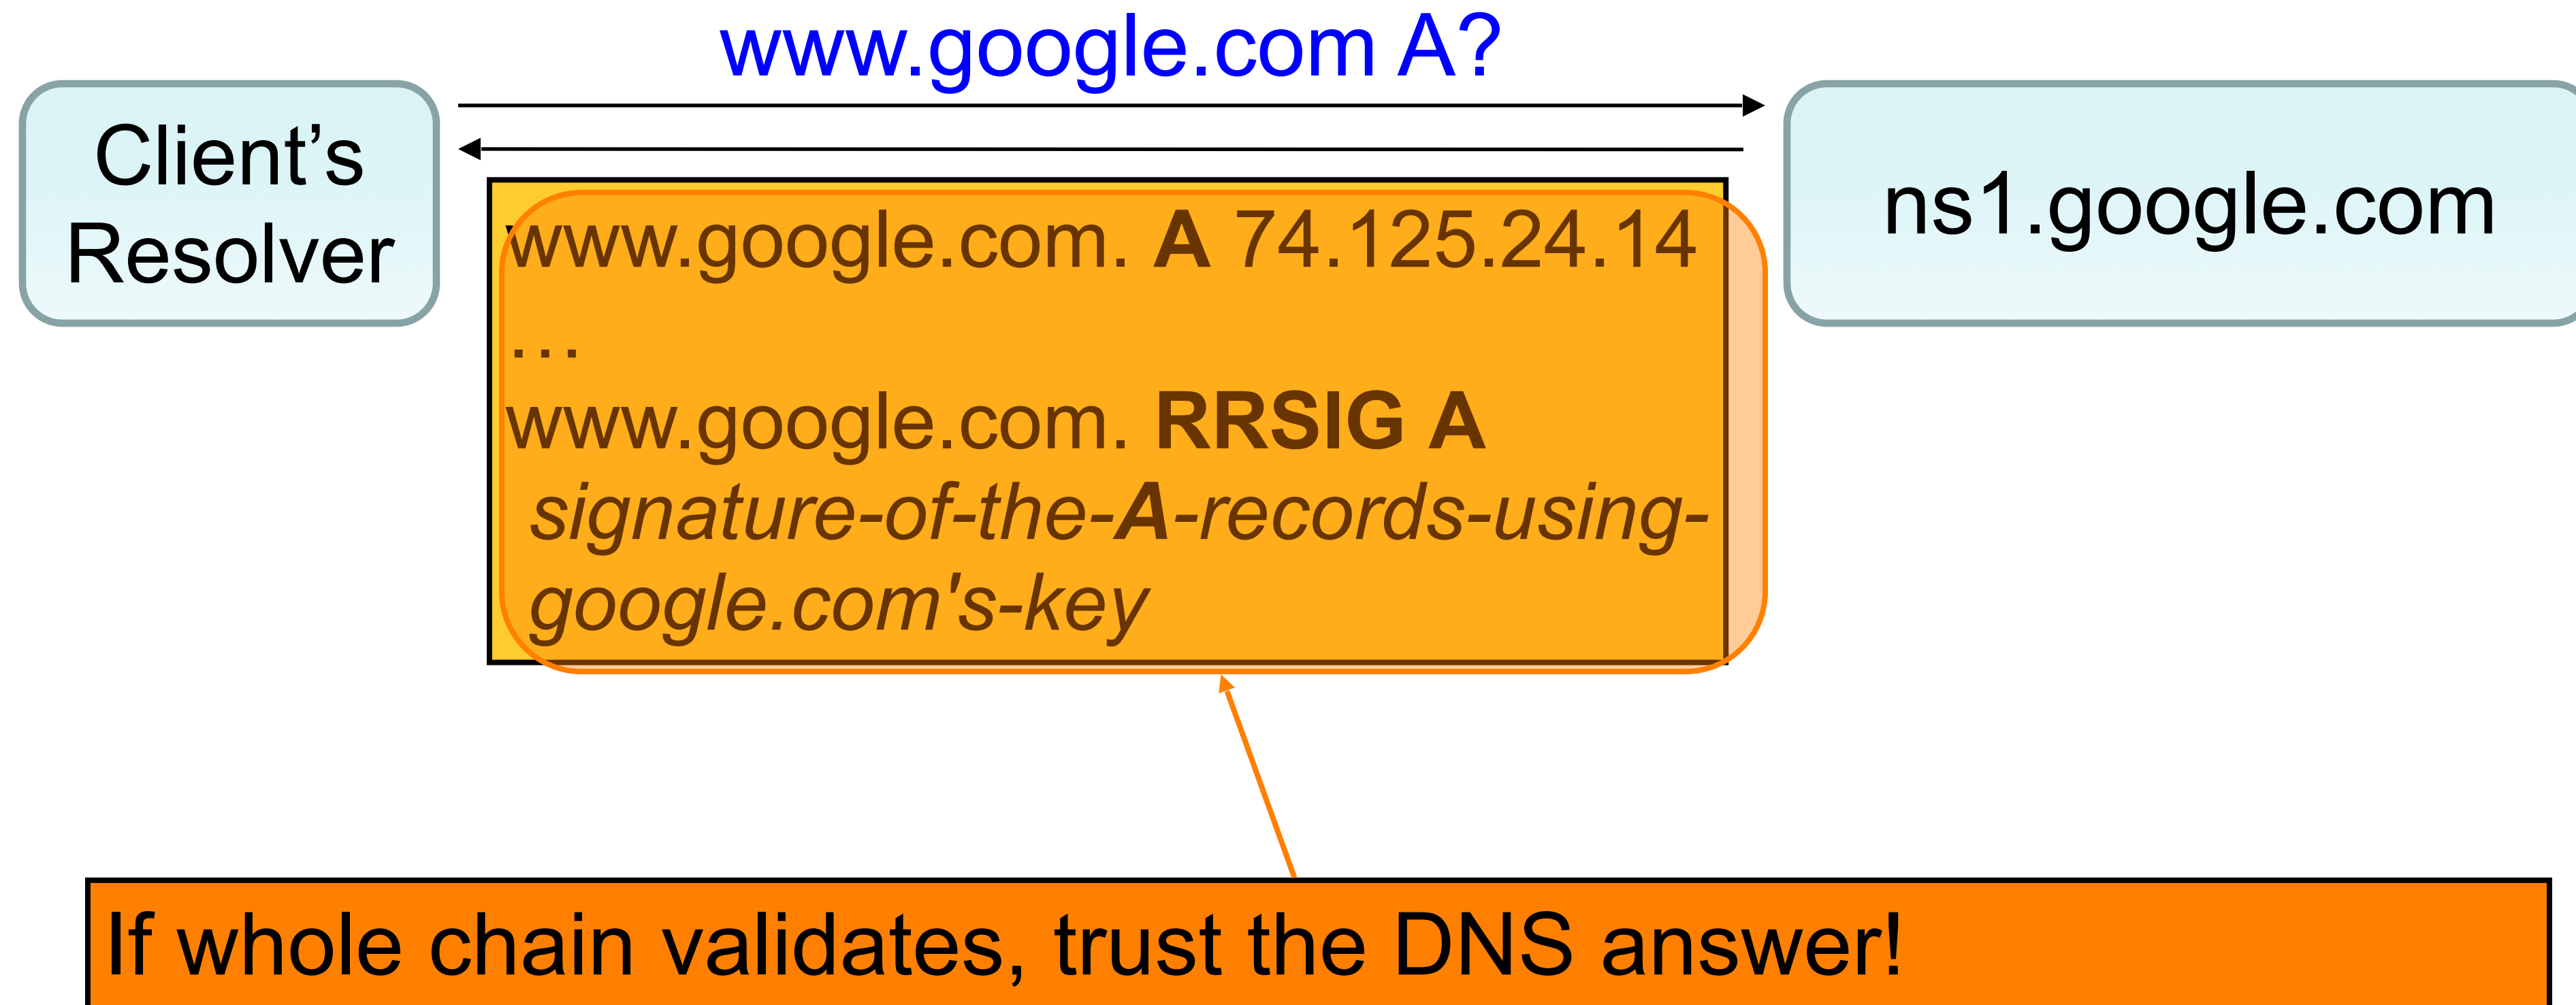
DNSSEC (with simplifications):




DNSSEC (with simplifications):



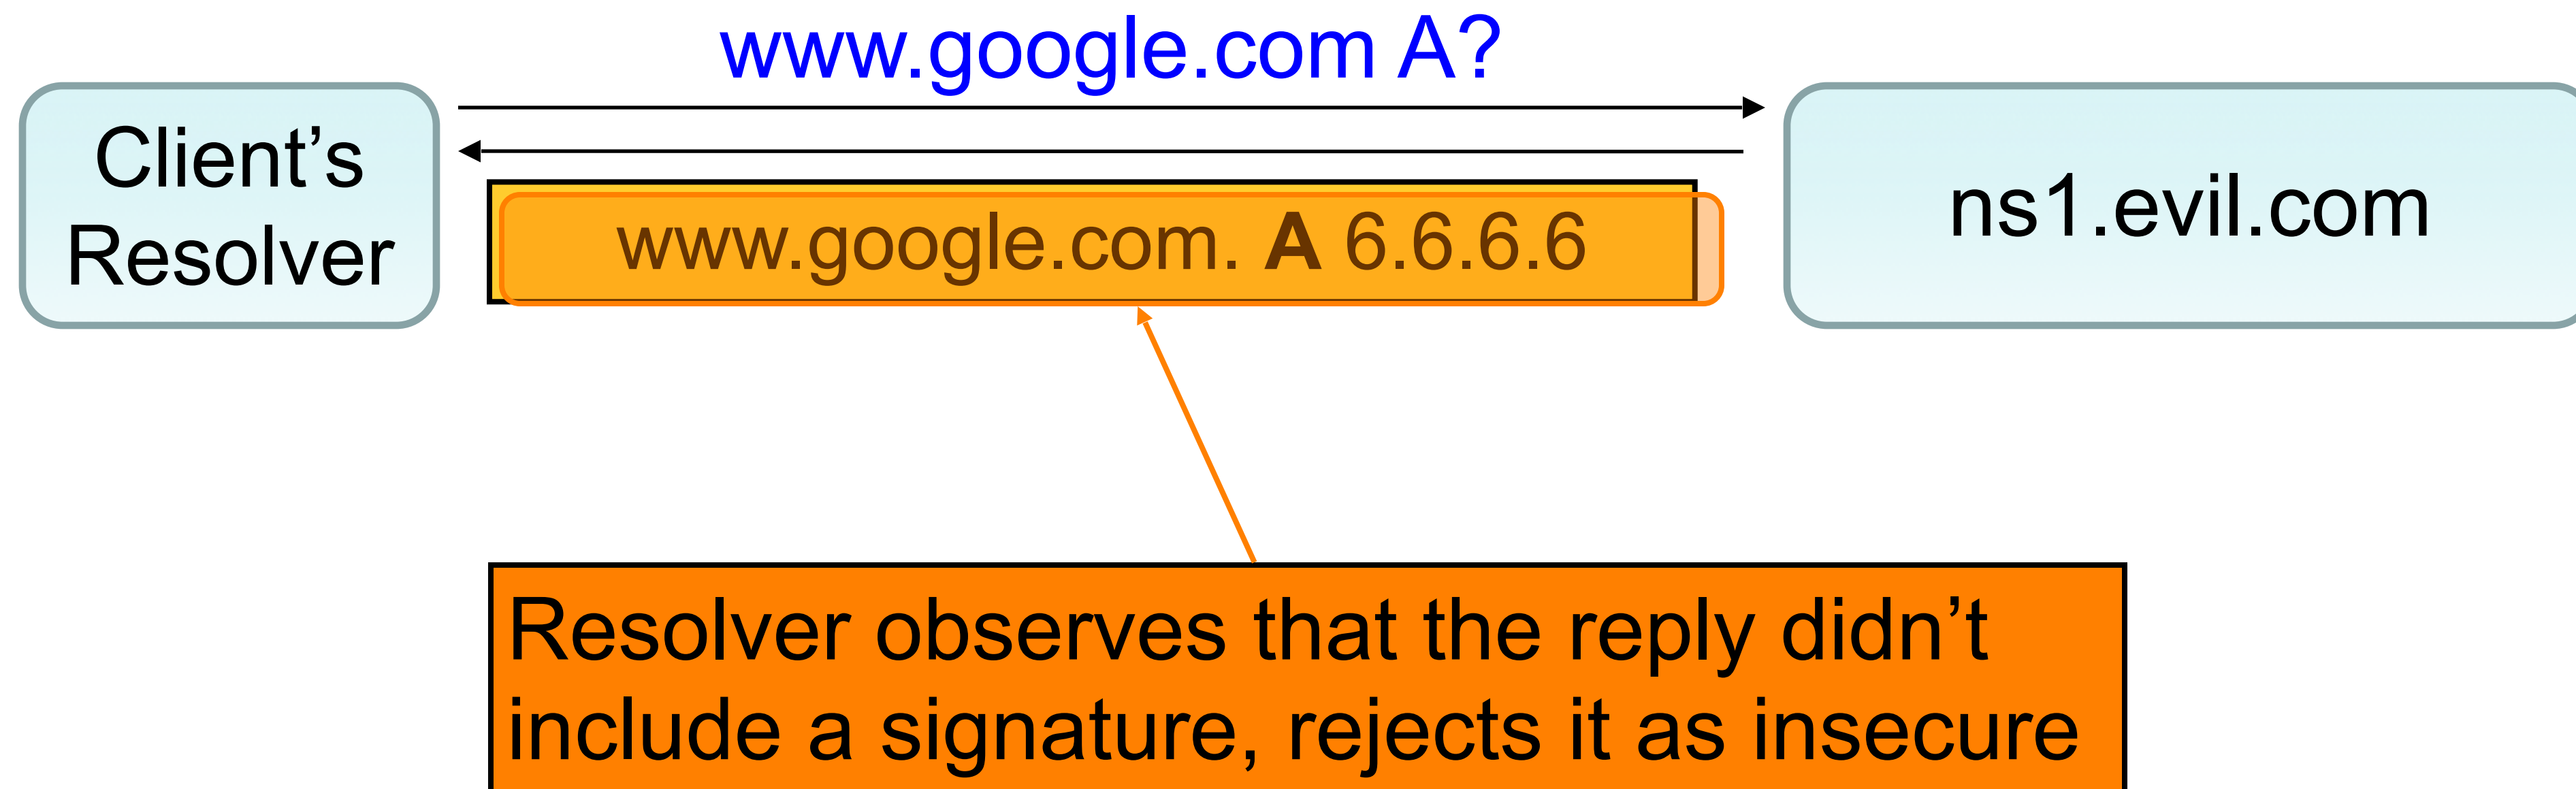
DNSSEC (with simplifications):




DNSSEC - Mallory attacks!



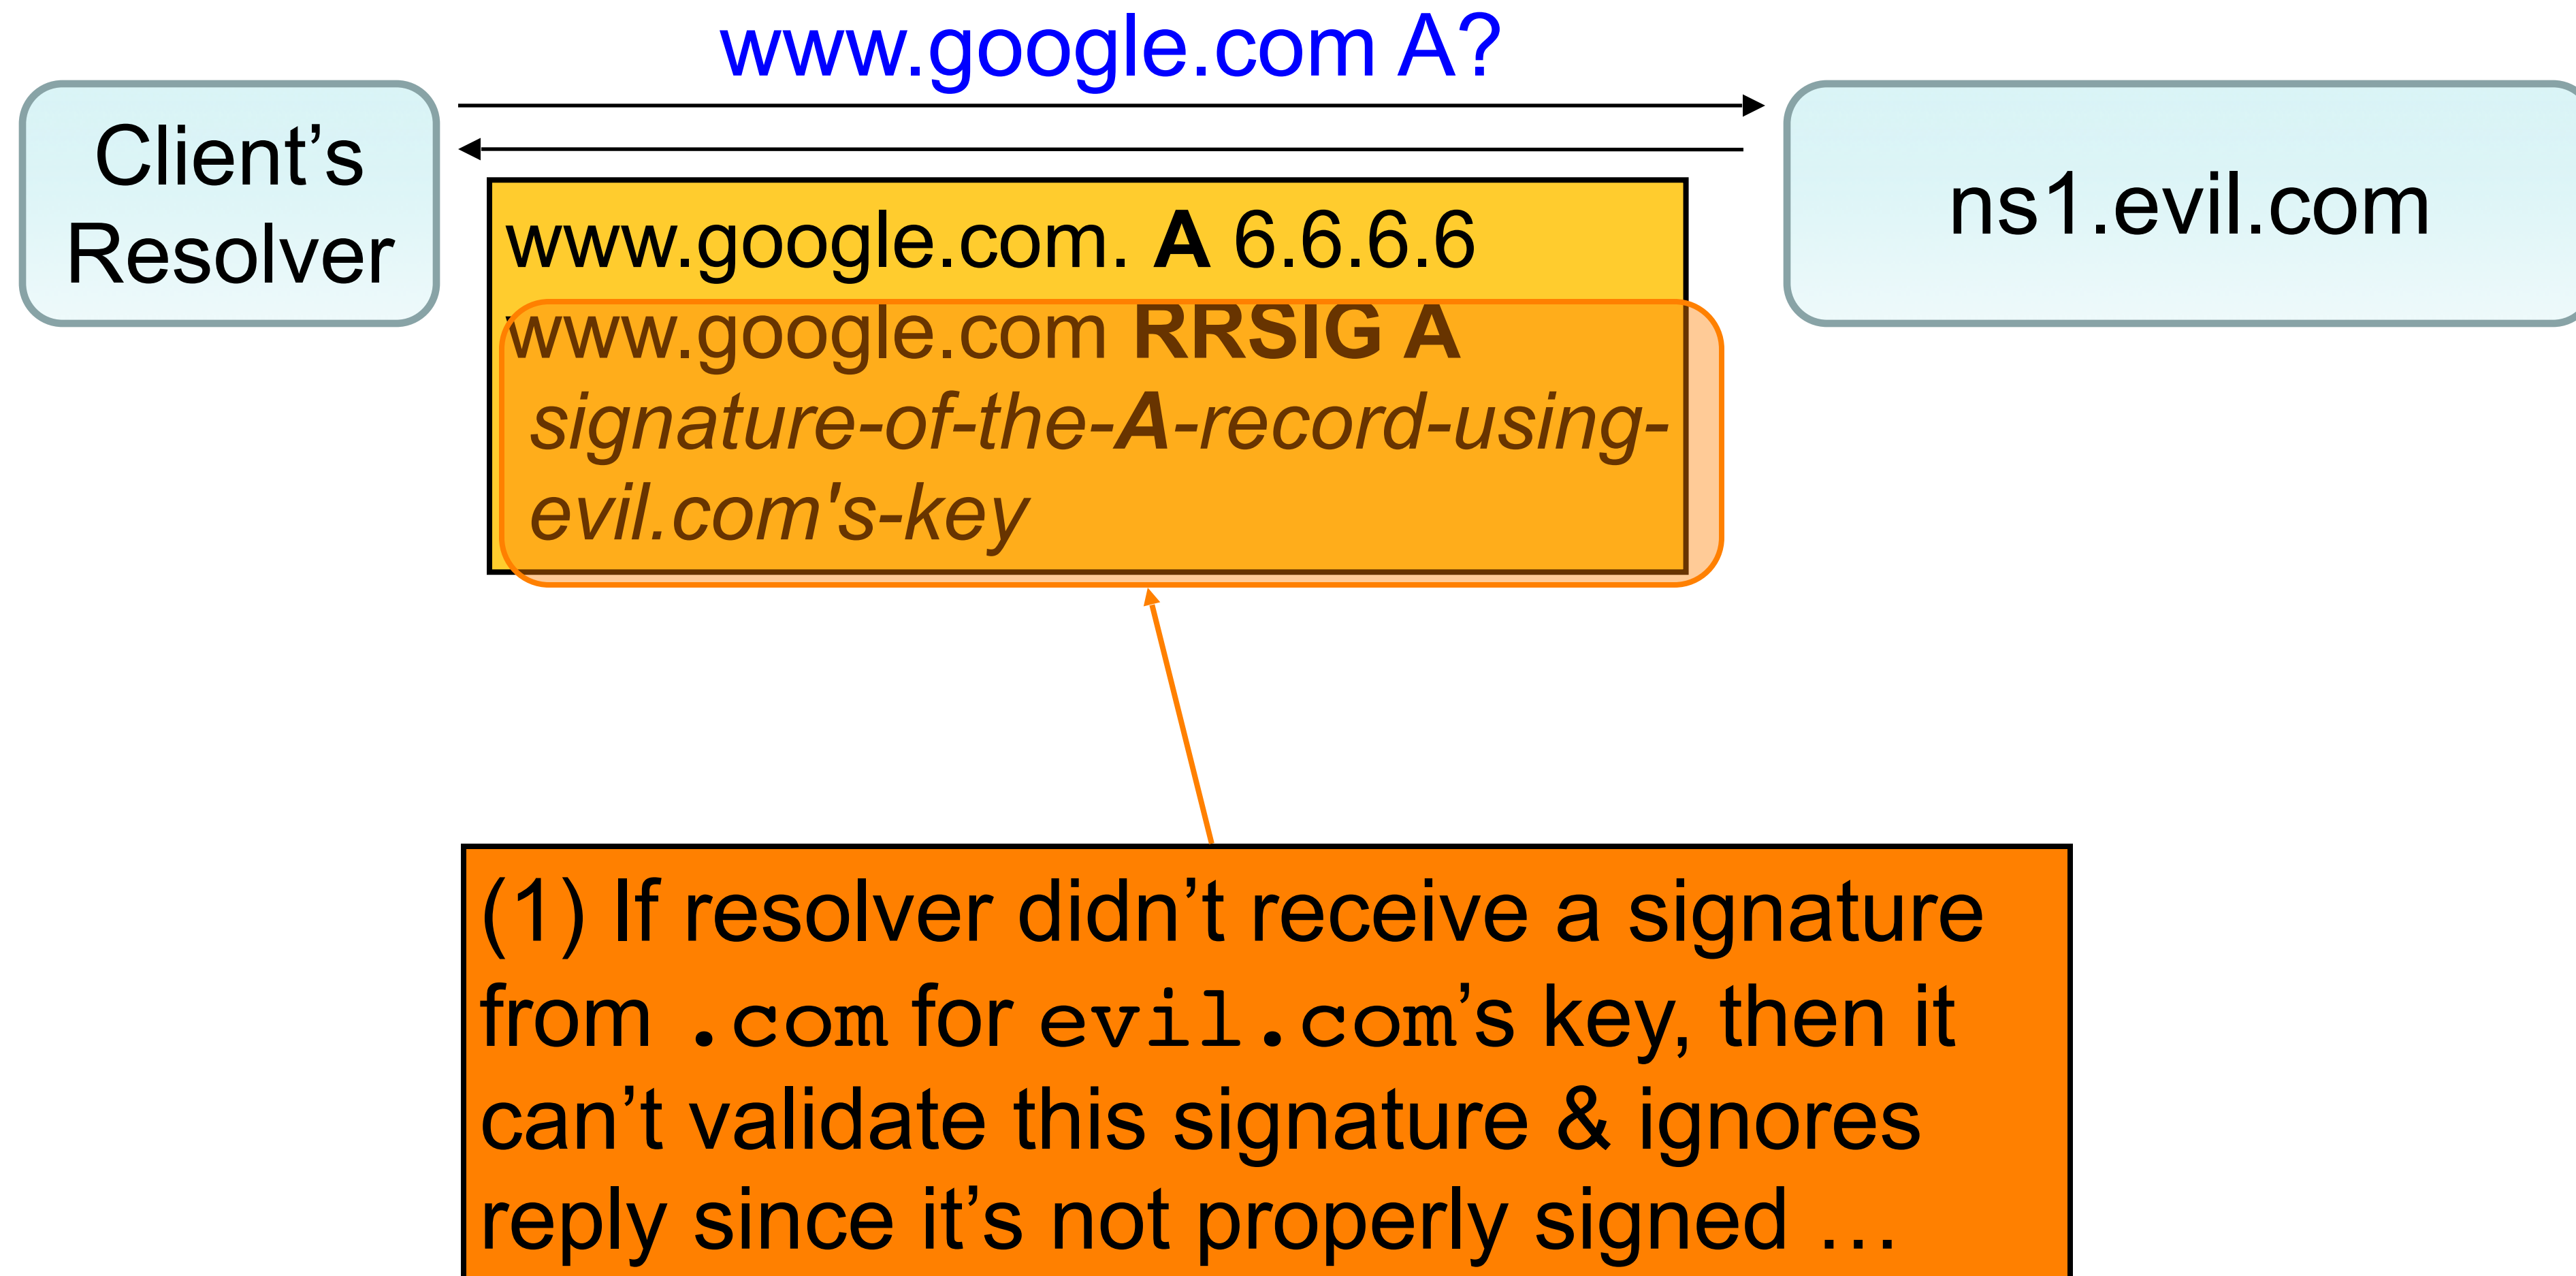
DNSSEC - Mallory attacks!



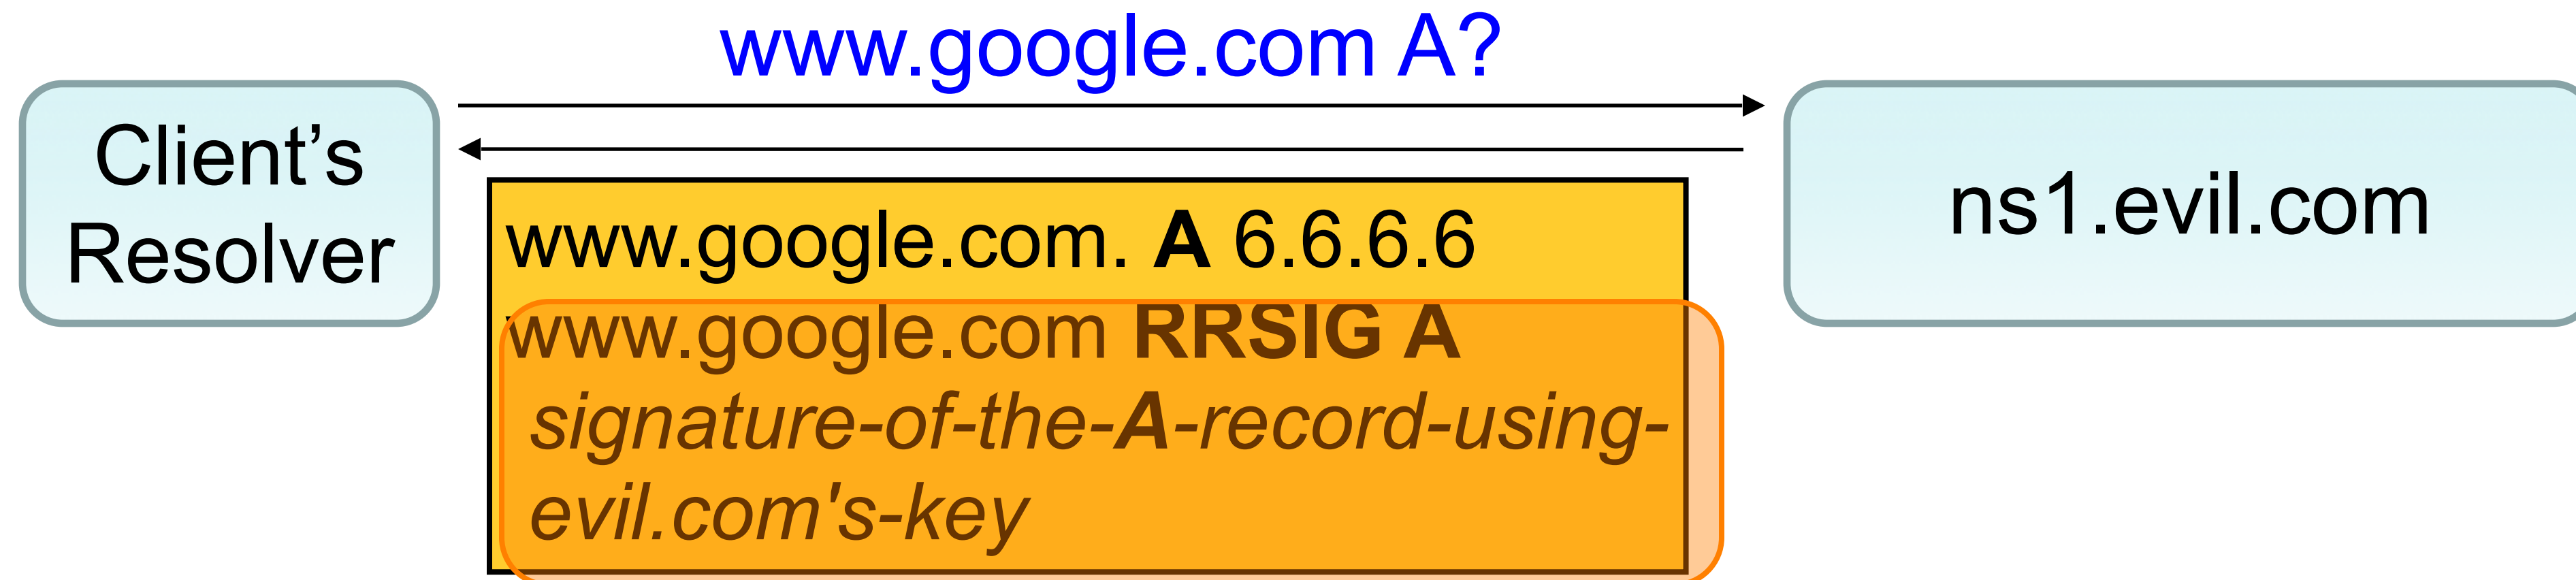
DNSSEC - Mallory attacks!



DNSSEC - Mallory attacks!




DNSSEC - Mallory attacks!

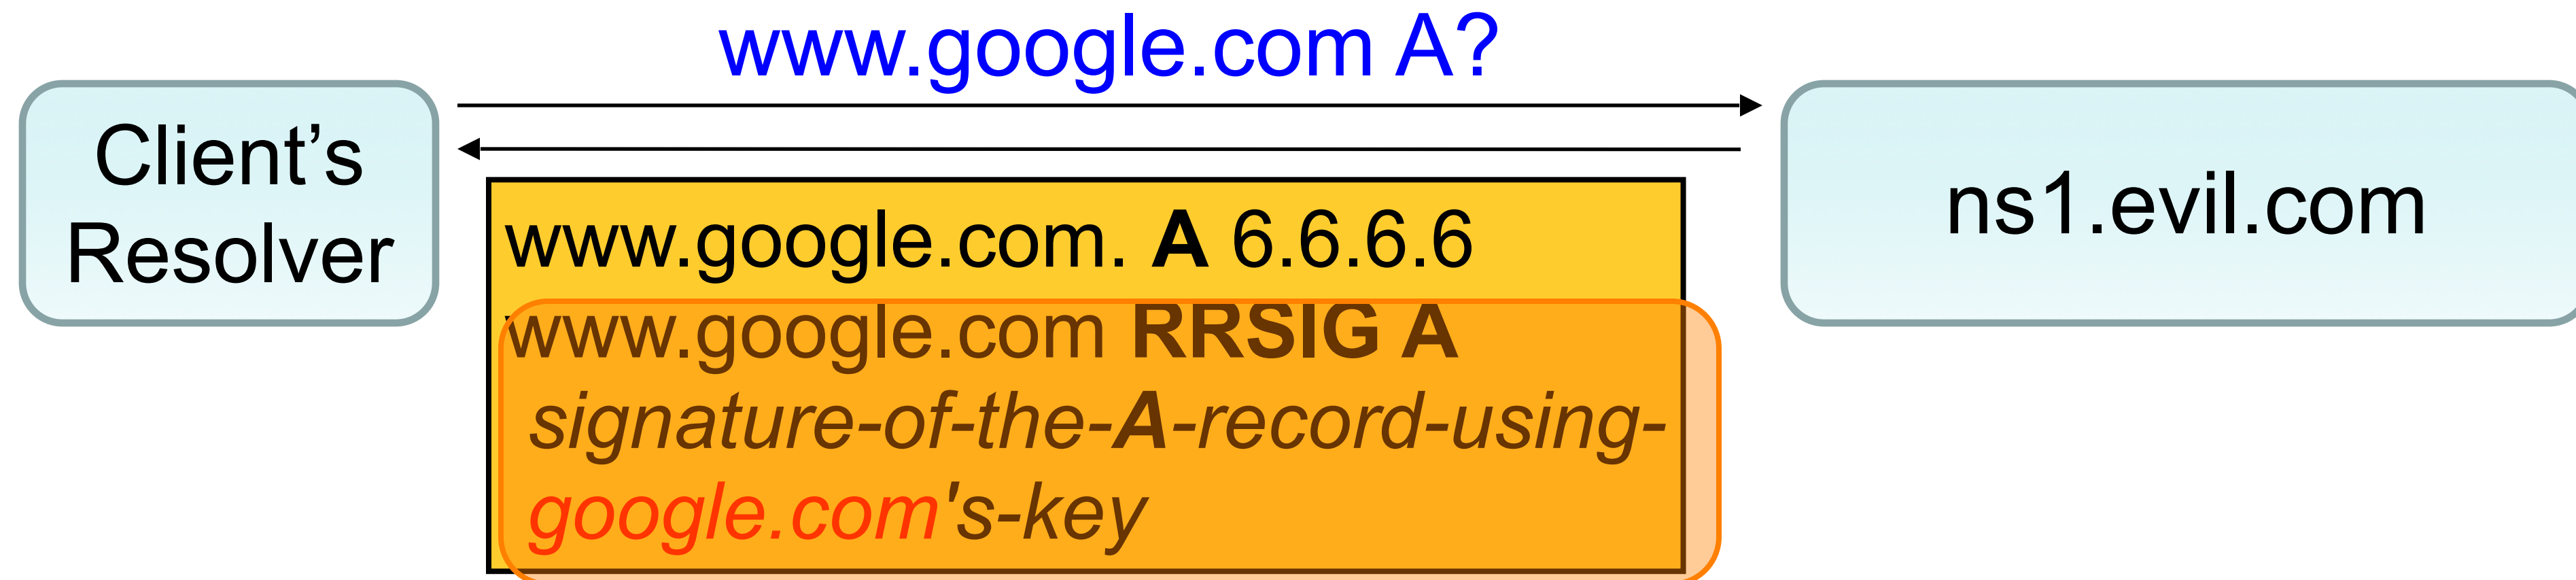


(2) If resolver *did* receive a signature from .com for evil.com's key, then it knows the key is for evil.com and not google.com ... and ignores it

DNSSEC - Mallory attacks!

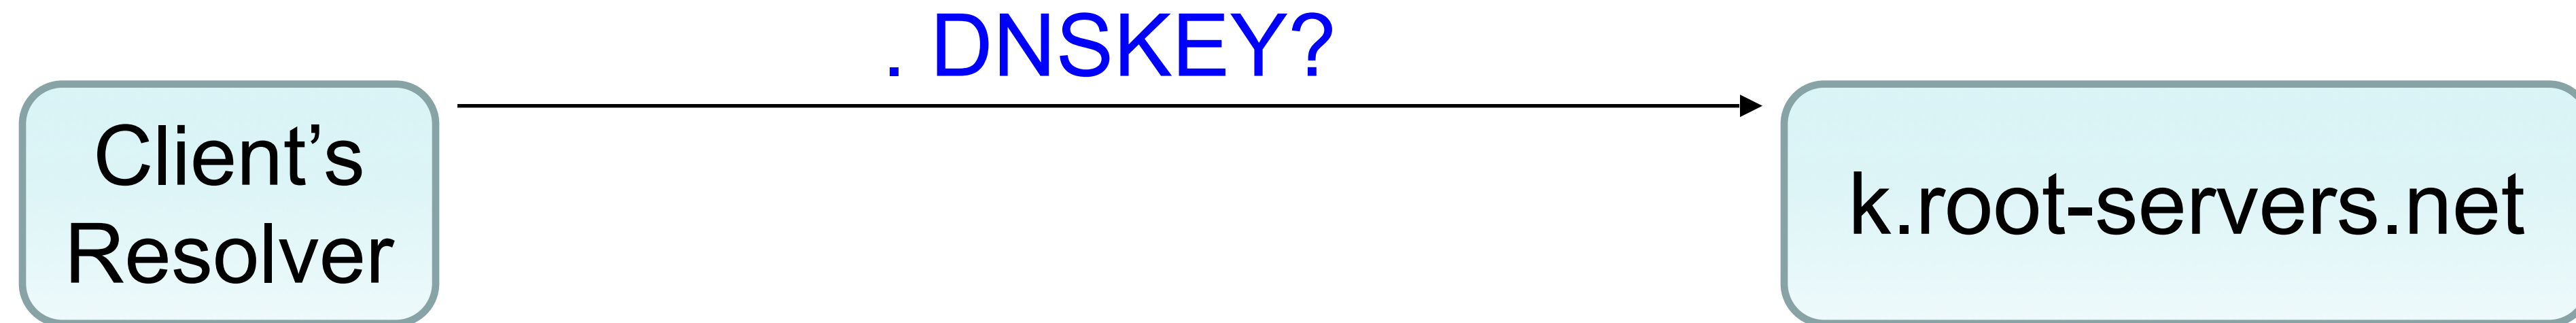


DNSSEC - Mallory attacks!



If signature **actually** comes from google.com's key, resolver will believe it ...
... but no such signature should exist unless either:
(1) google.com *intended* to sign the RR, or
(2) google.com's private key was compromised

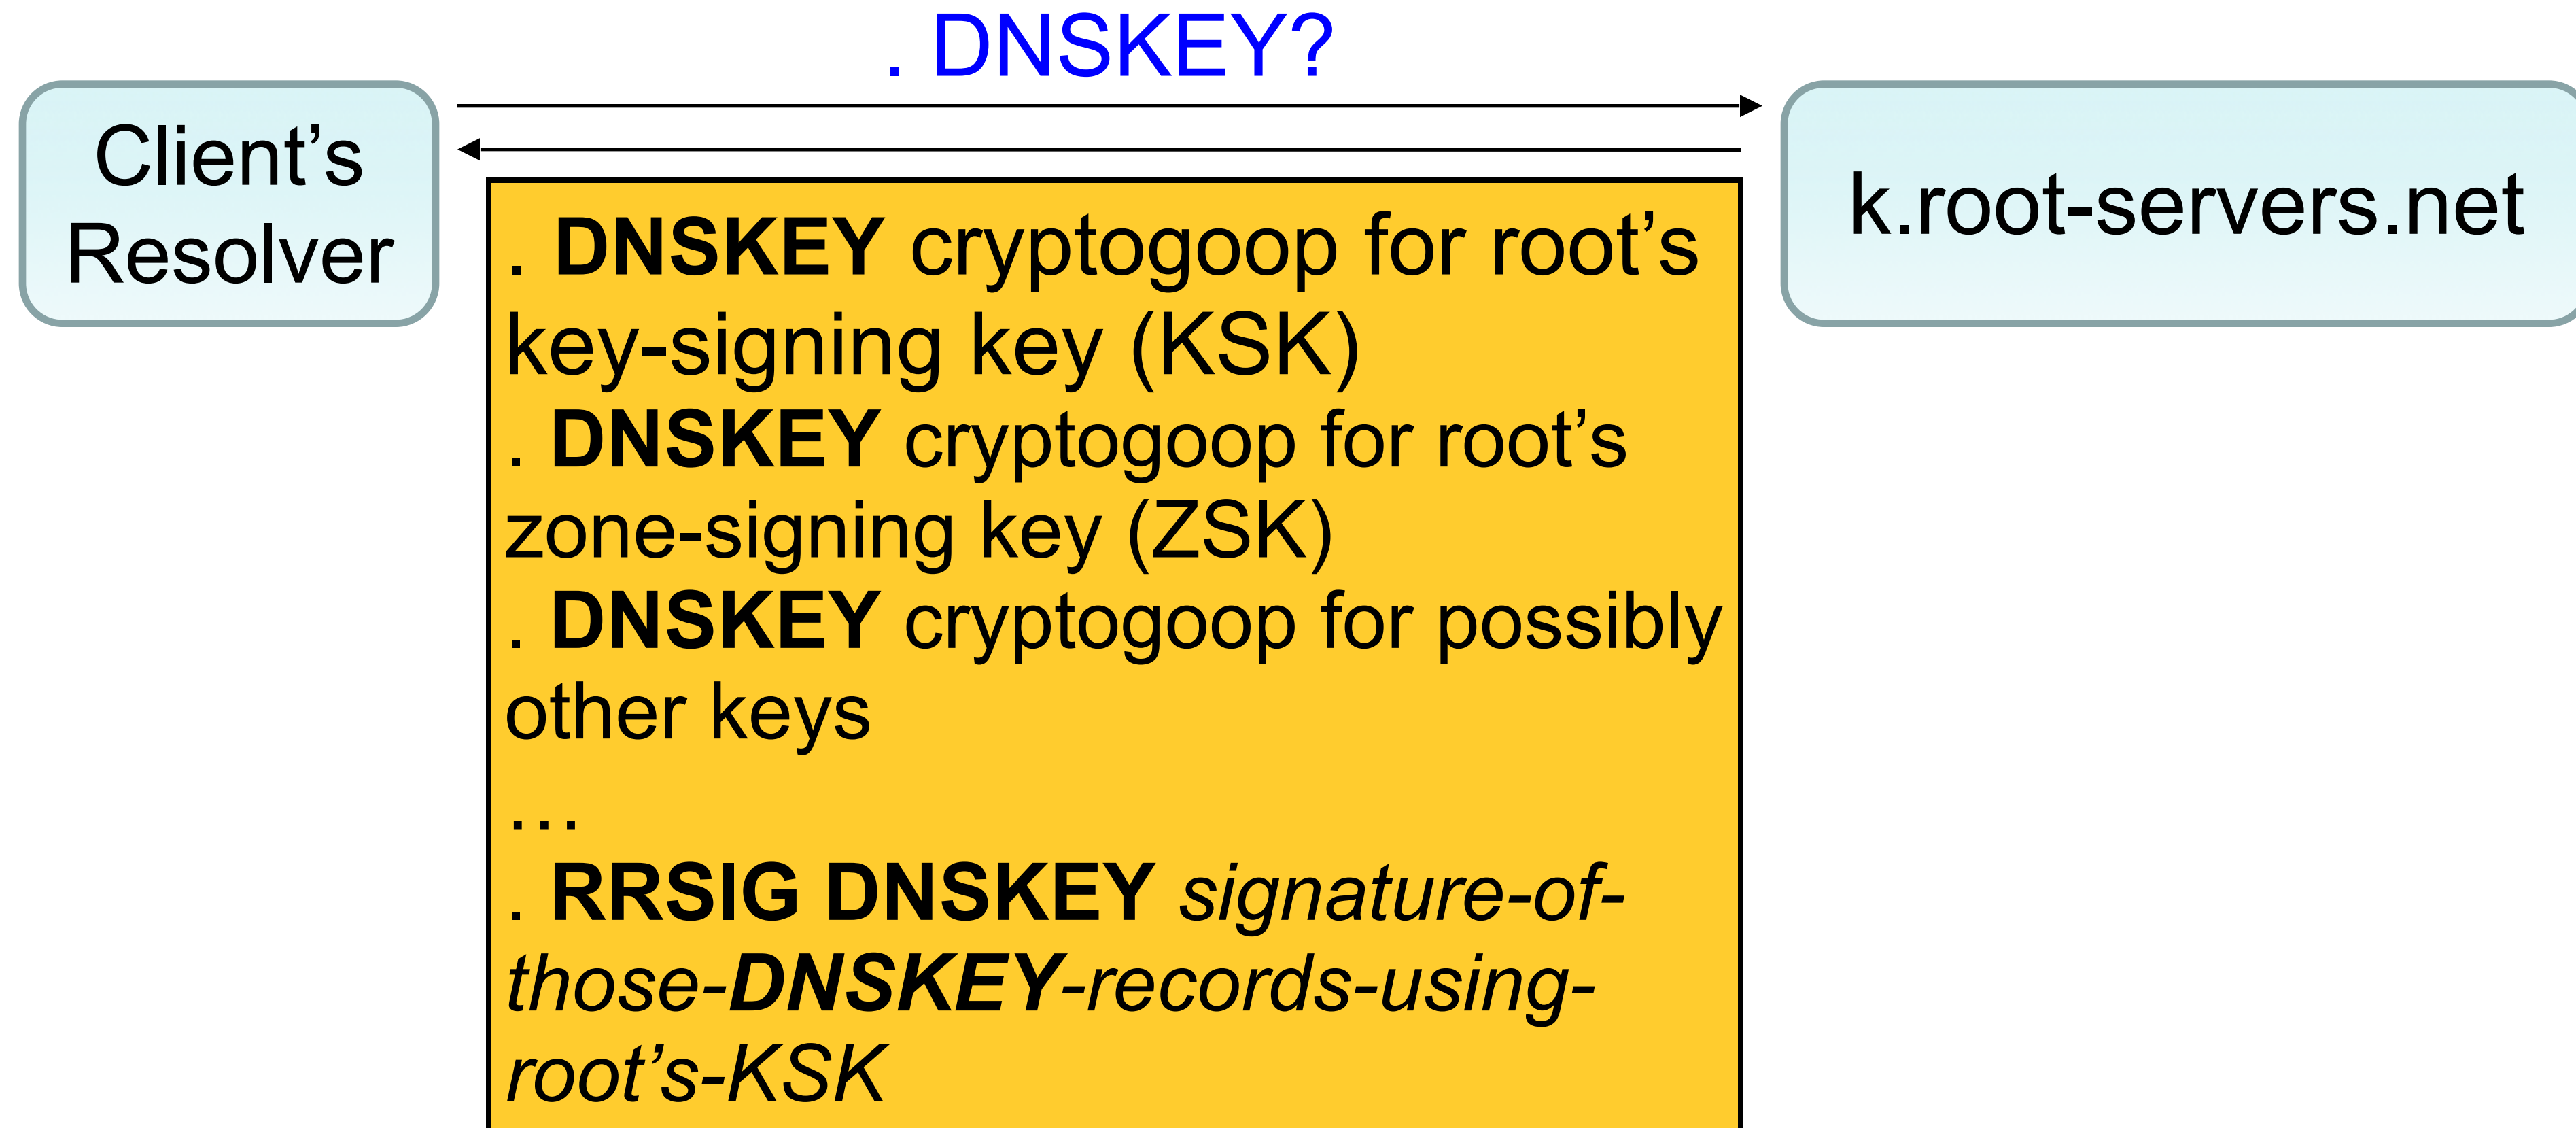
DNSSEC: Accessing keys



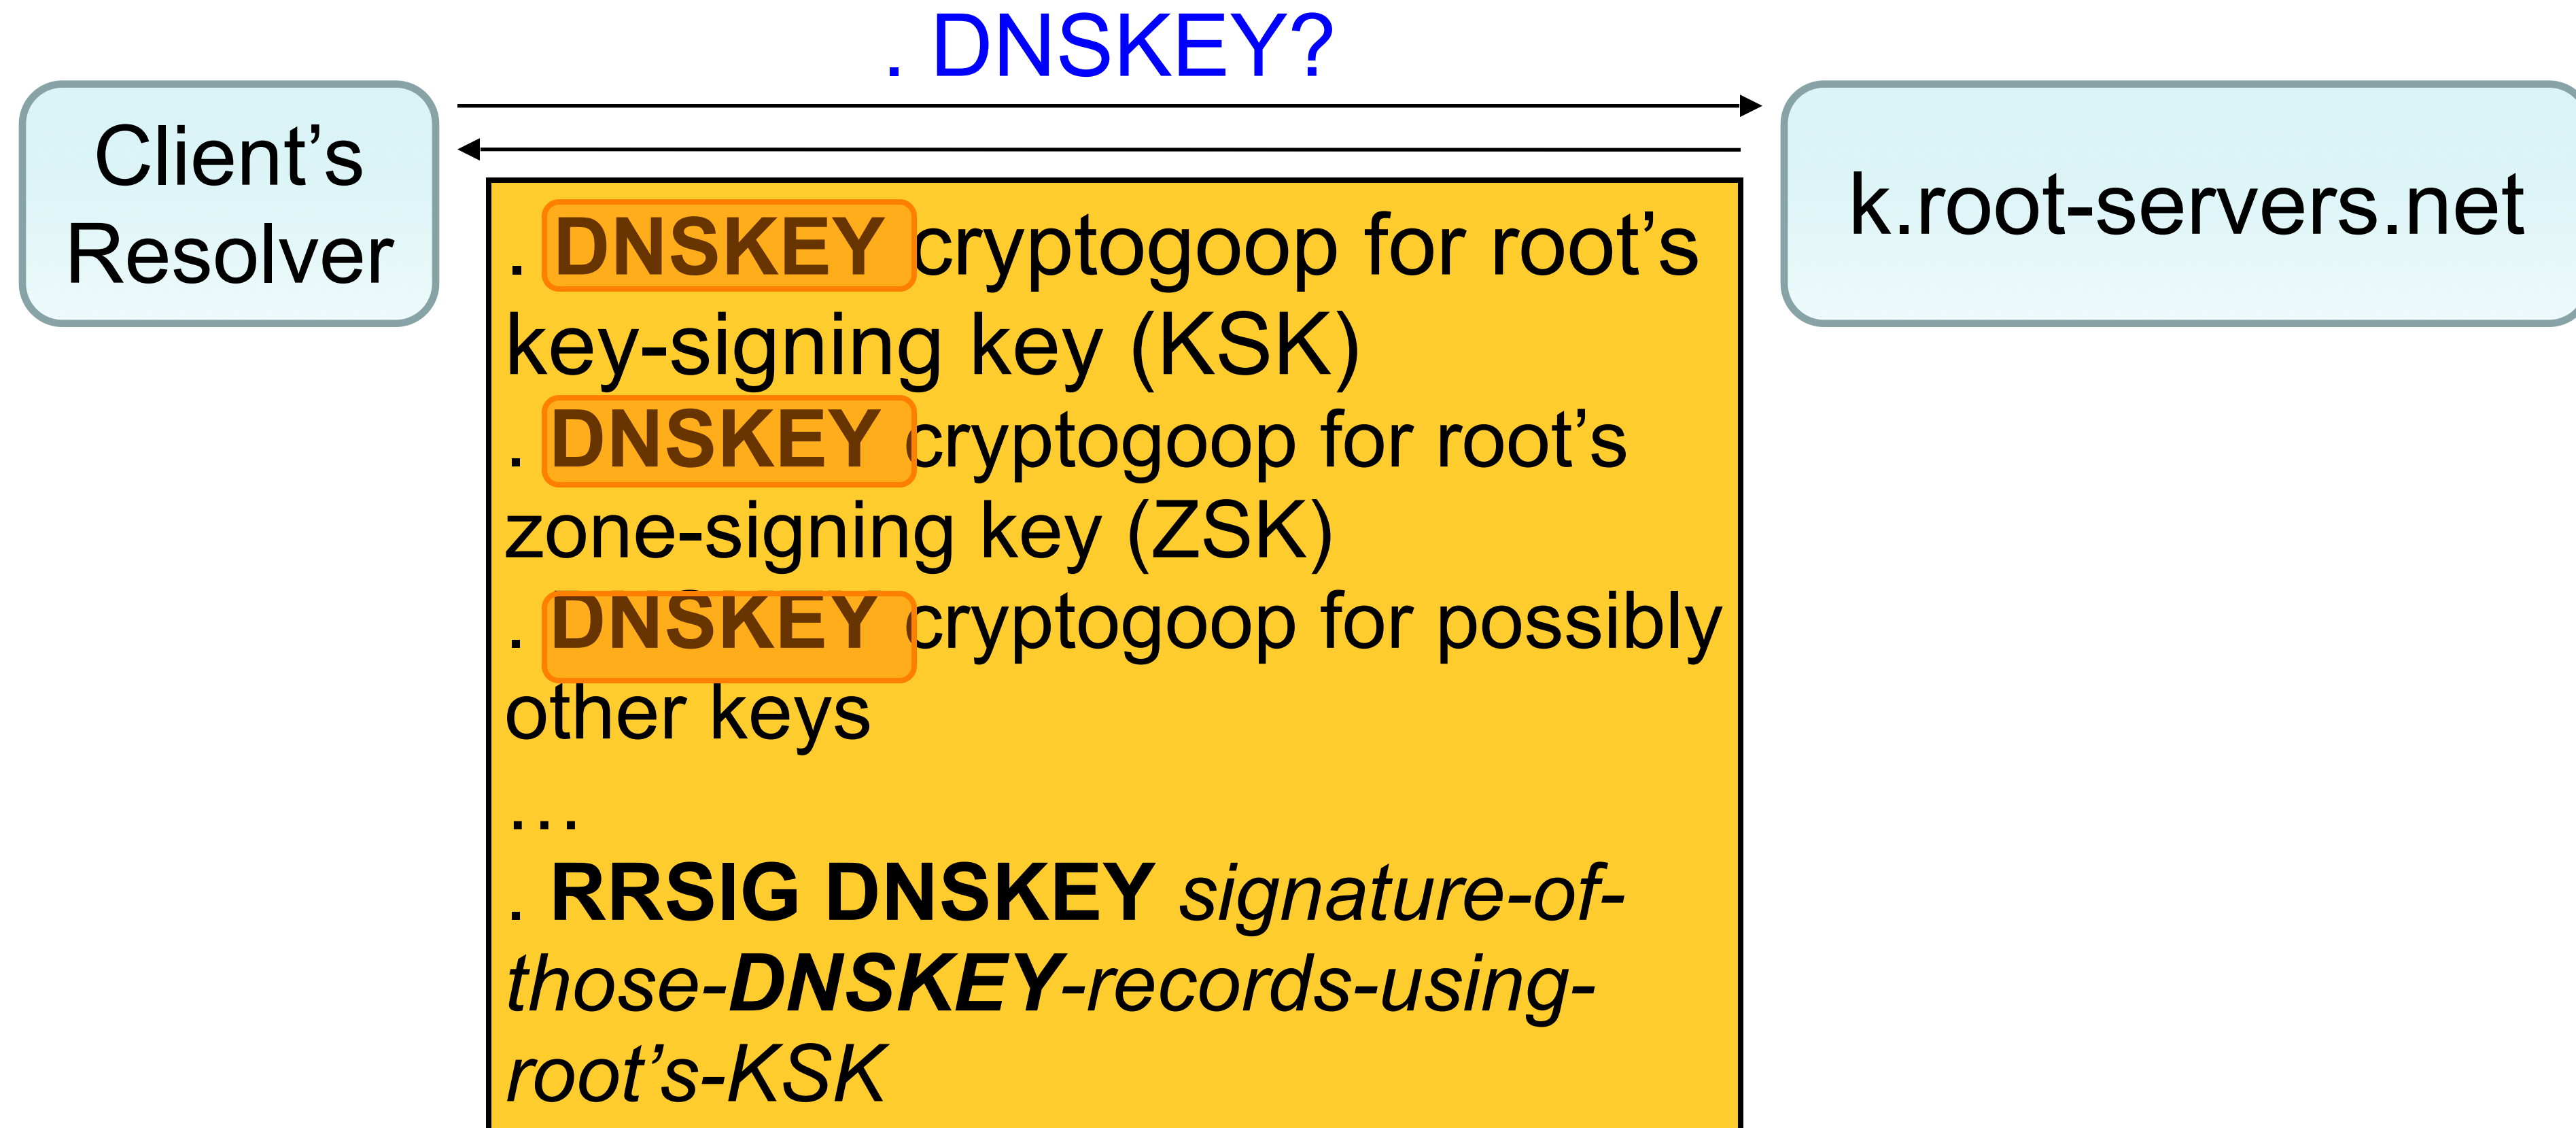
To build up the keys needed for validation, our client contacts each name server in the DNS hierarchy asking it for all of its associated keys.

Here we ask the root for its keys (one of which we already know as our **trust anchor**).

DNSSEC: Accessing keys

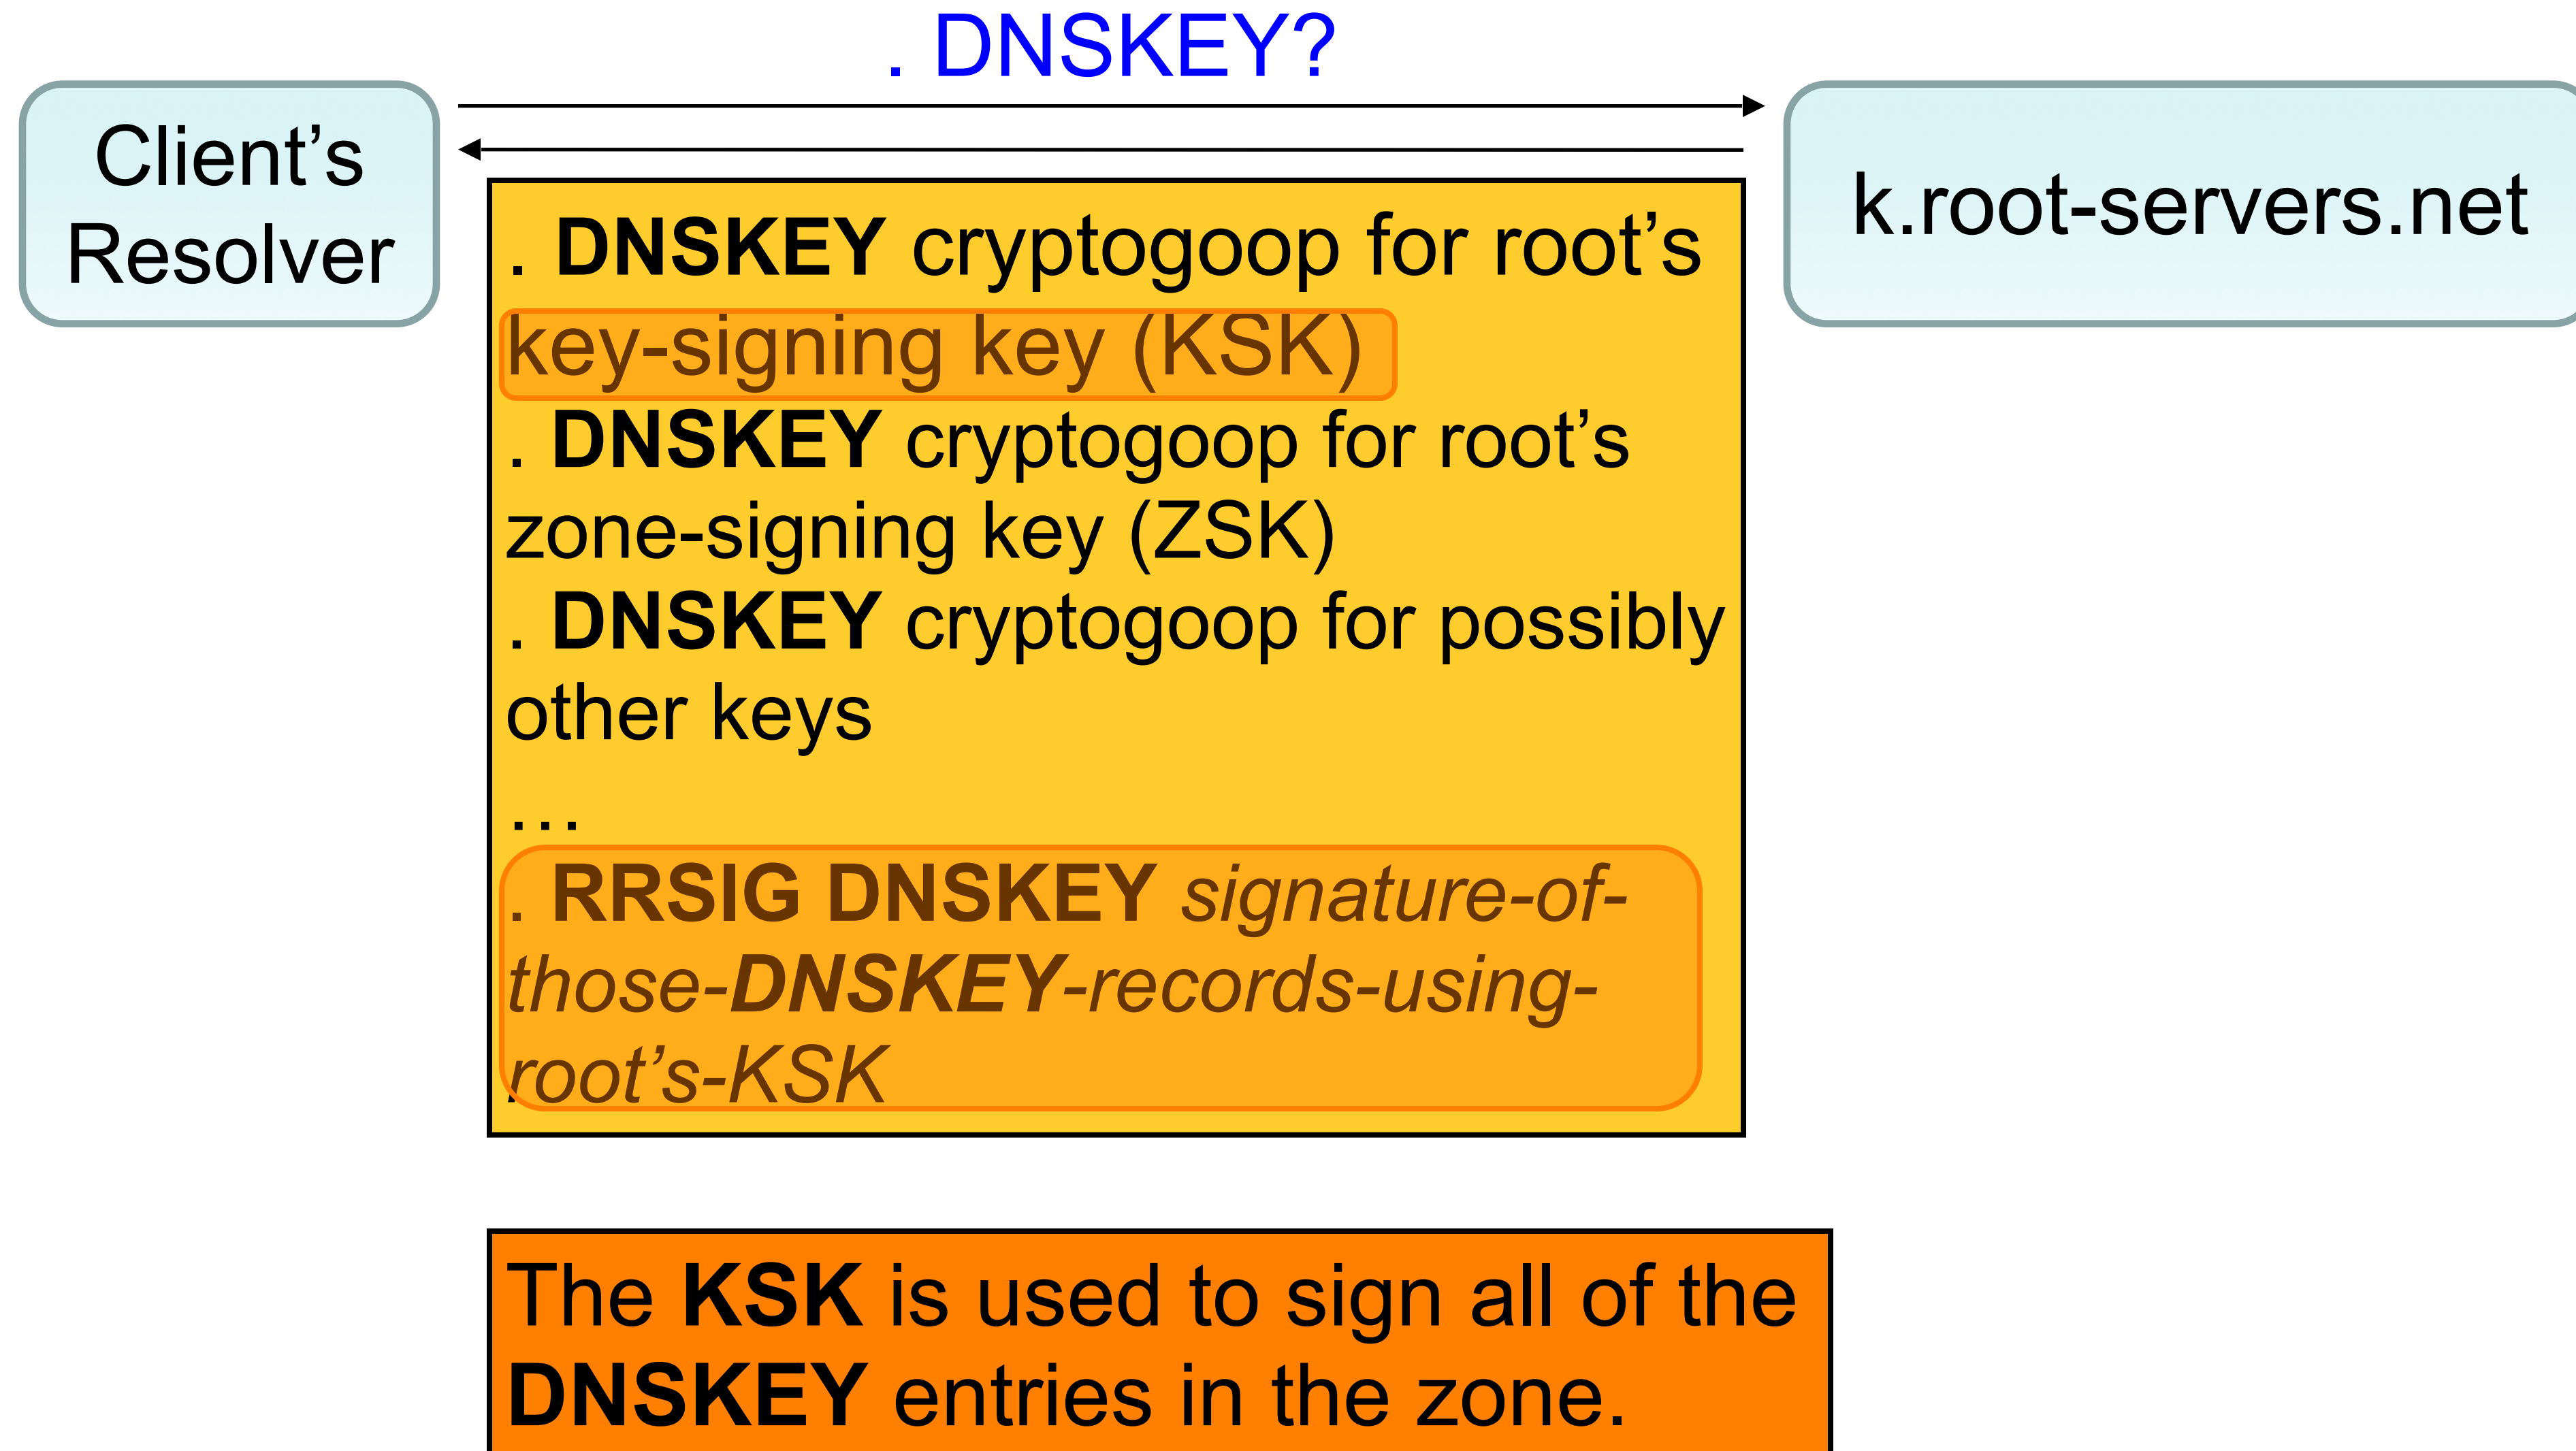


DNSSEC: Accessing keys

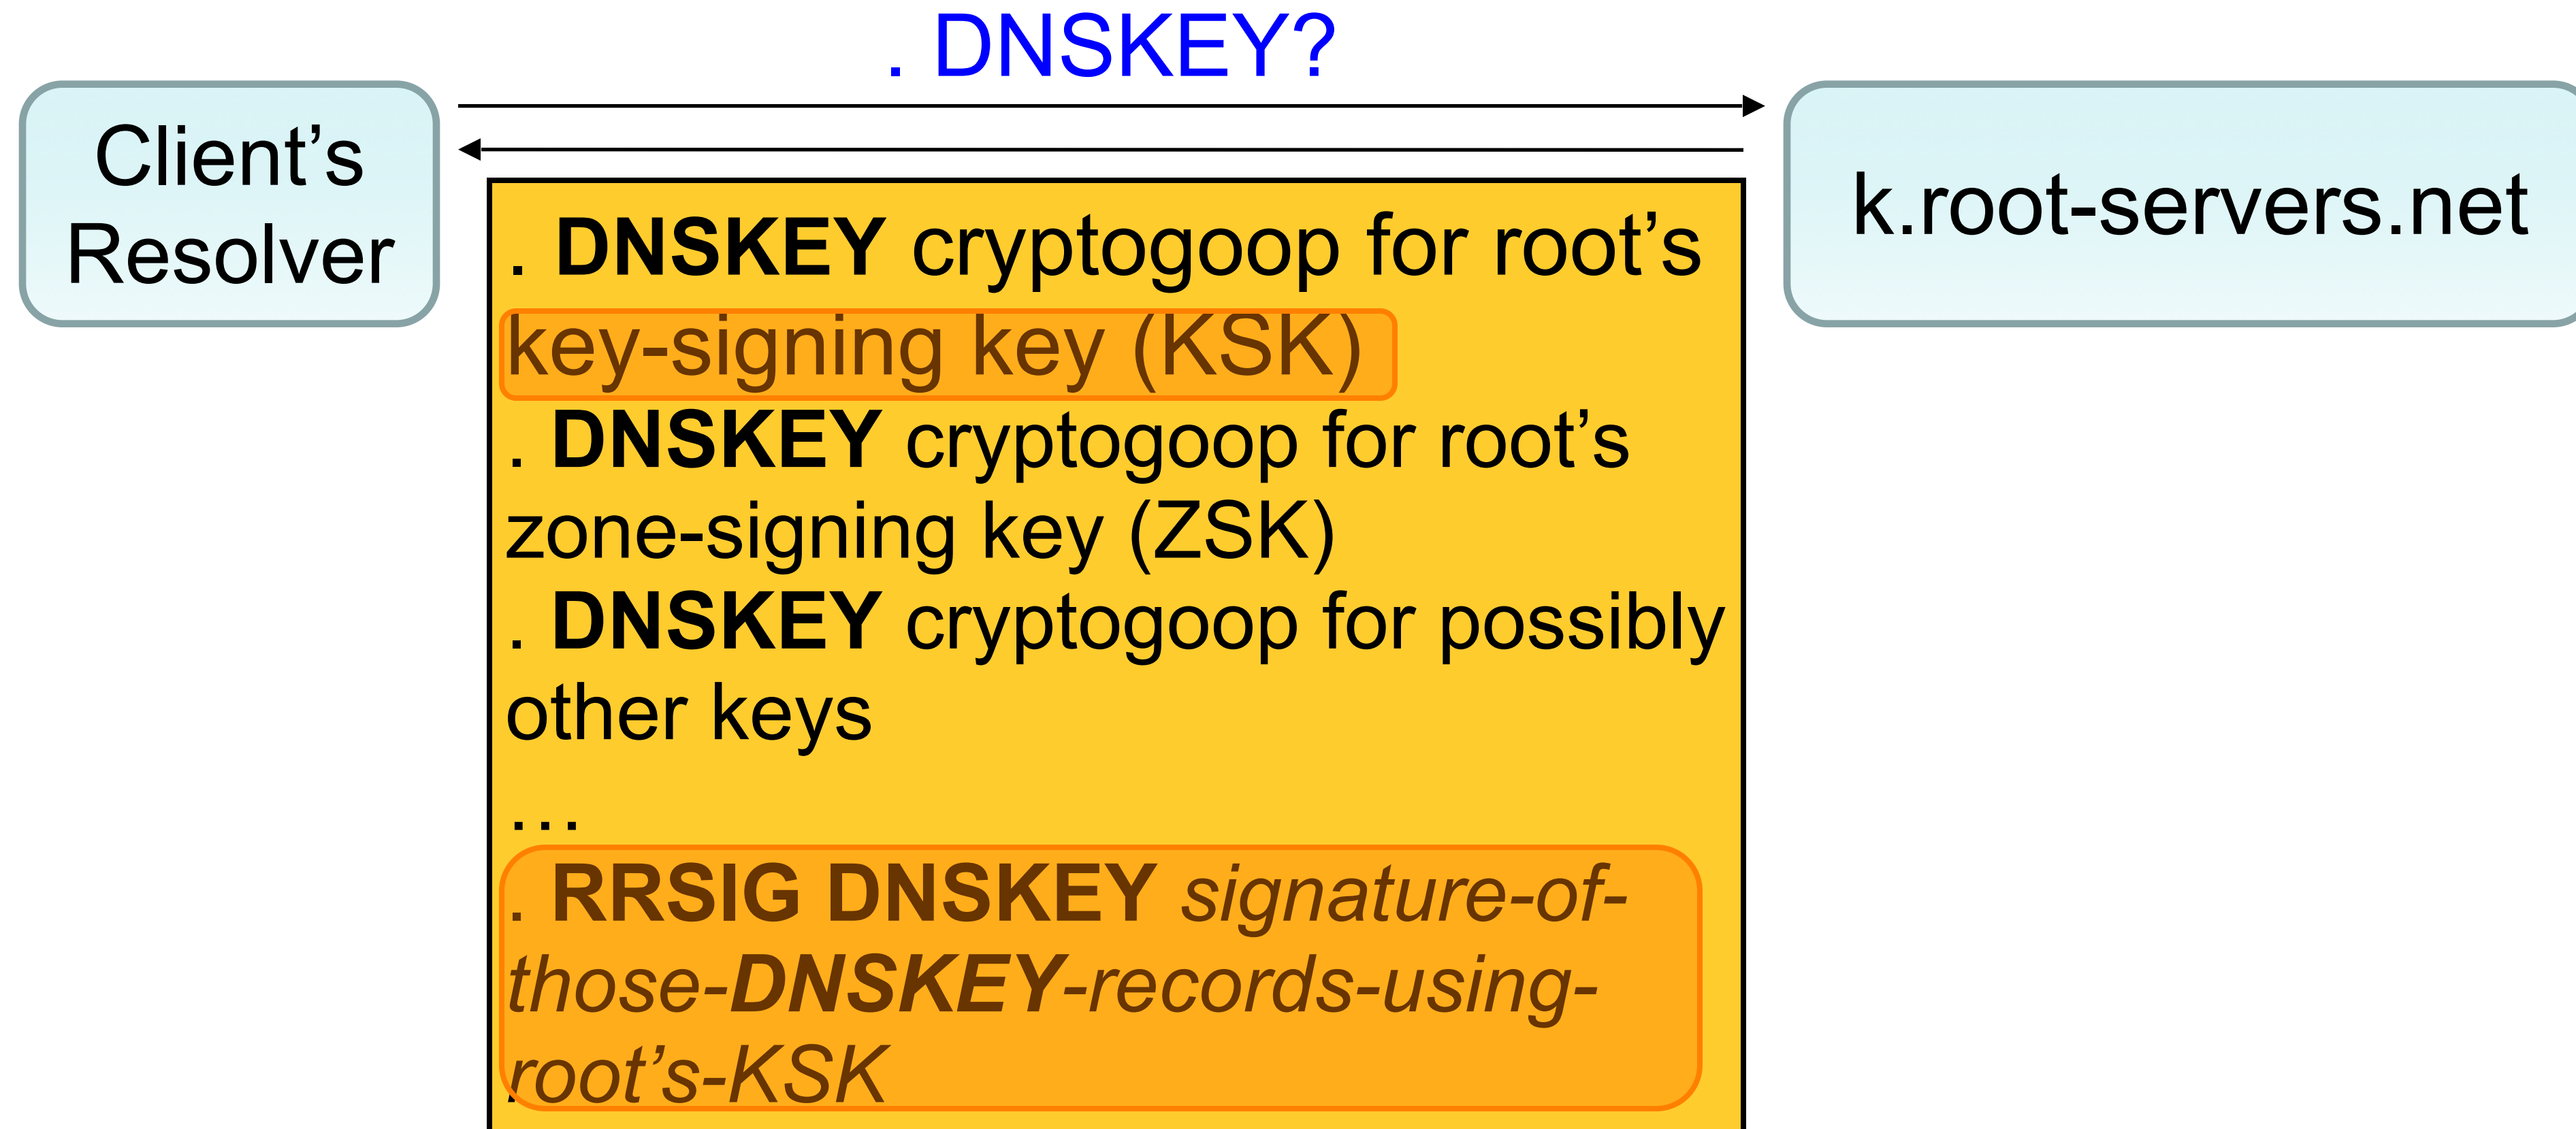


Each **DNSKEY** is a public key plus a description of the algorithms it's associated with (e.g., RSA+SHA256)

DNSSEC: Accessing keys

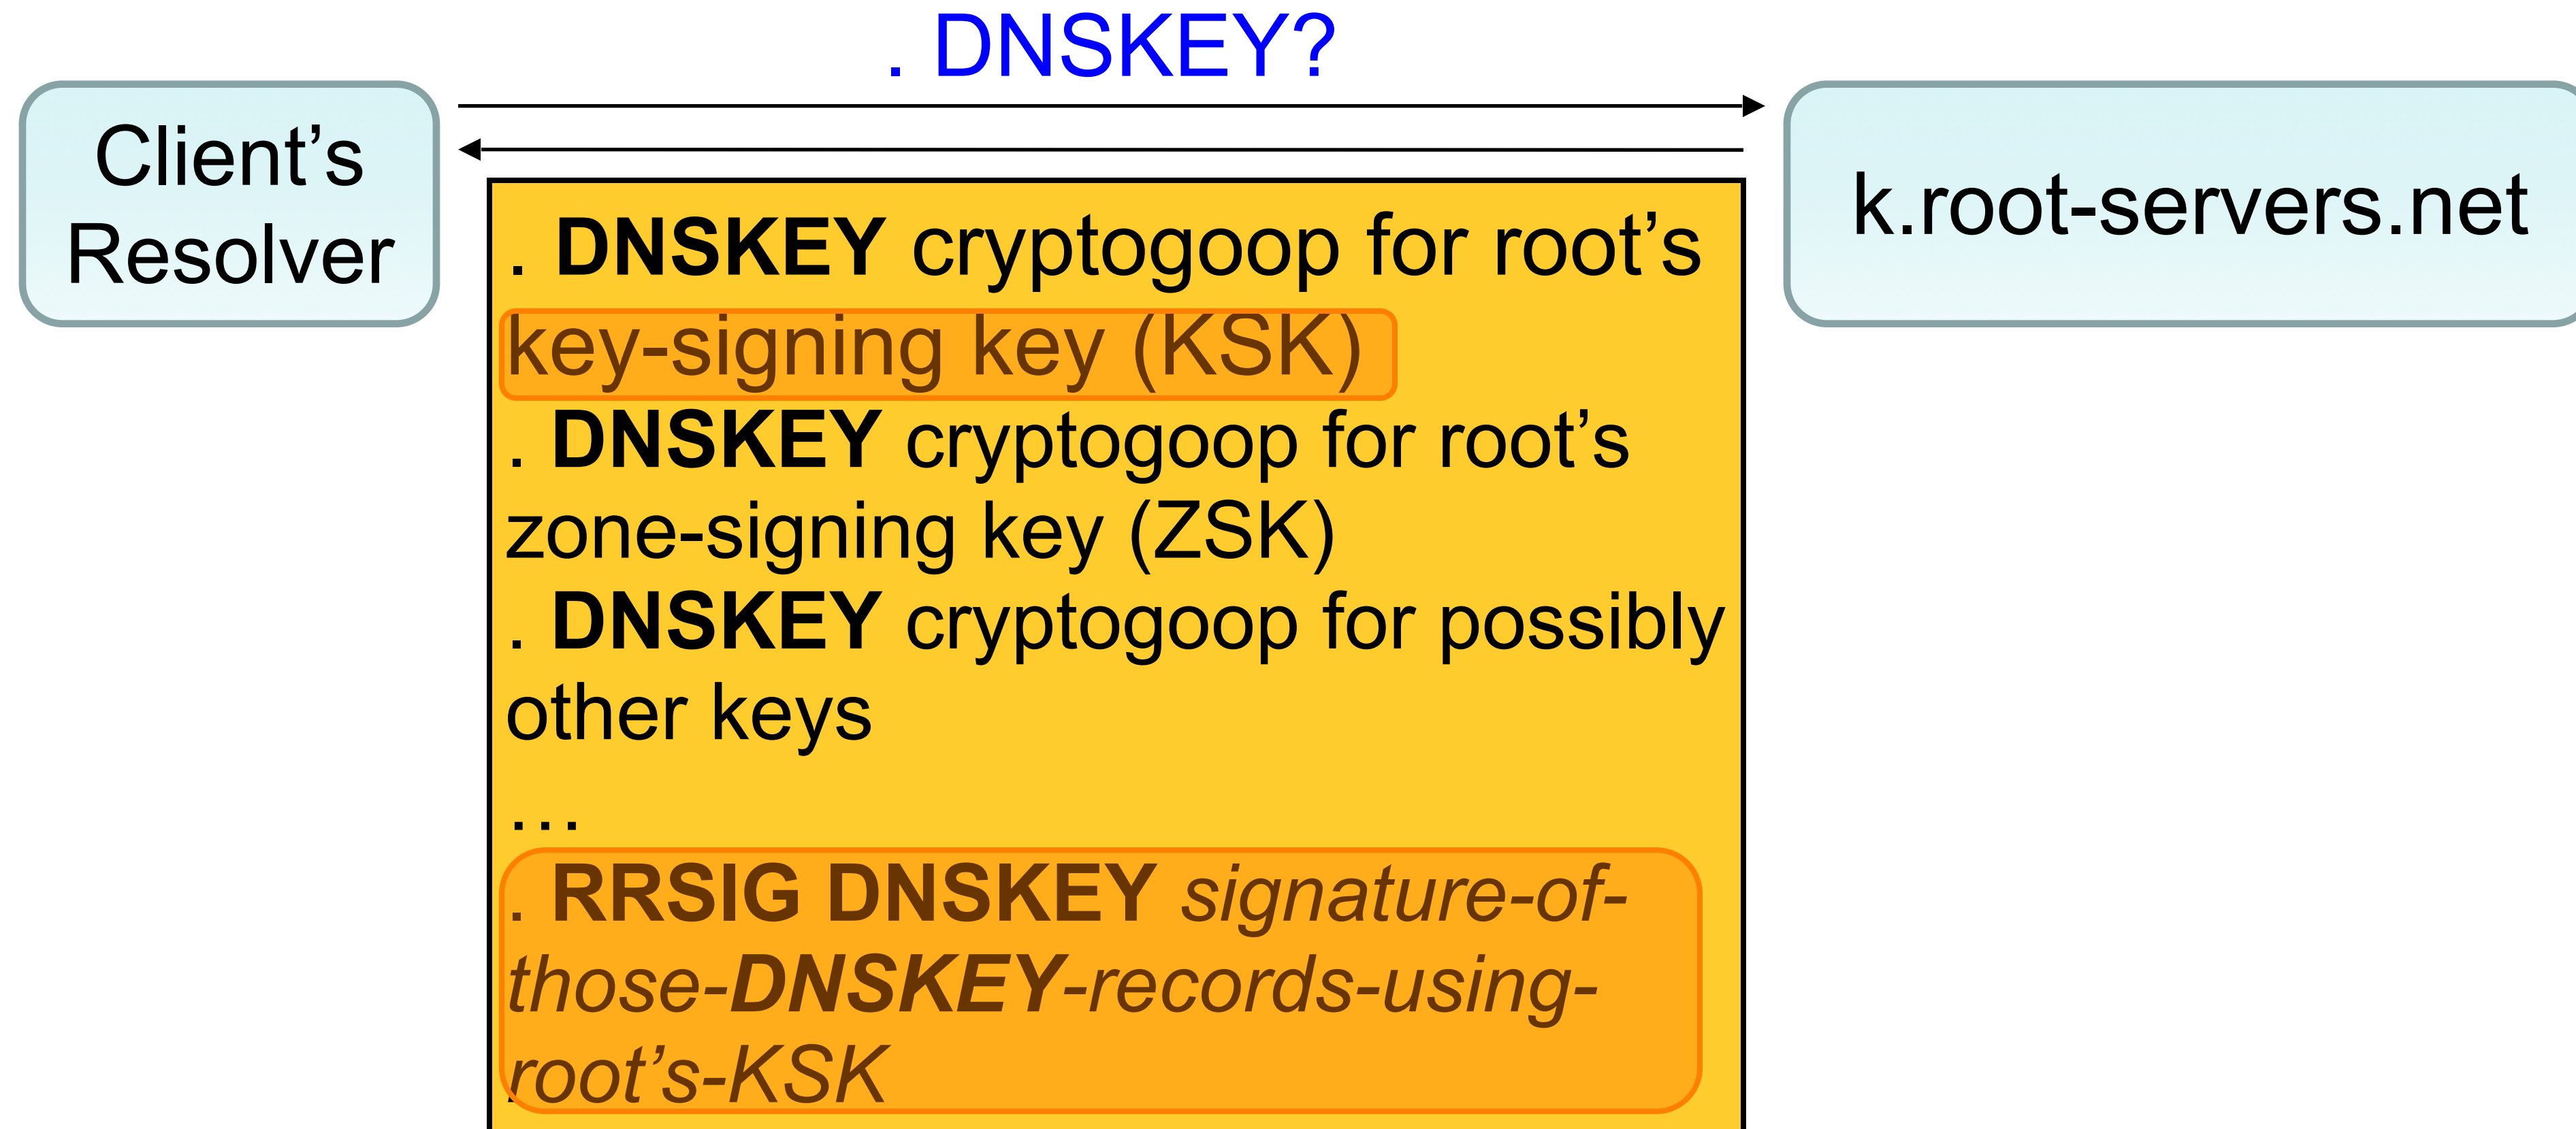


DNSSEC: Accessing keys



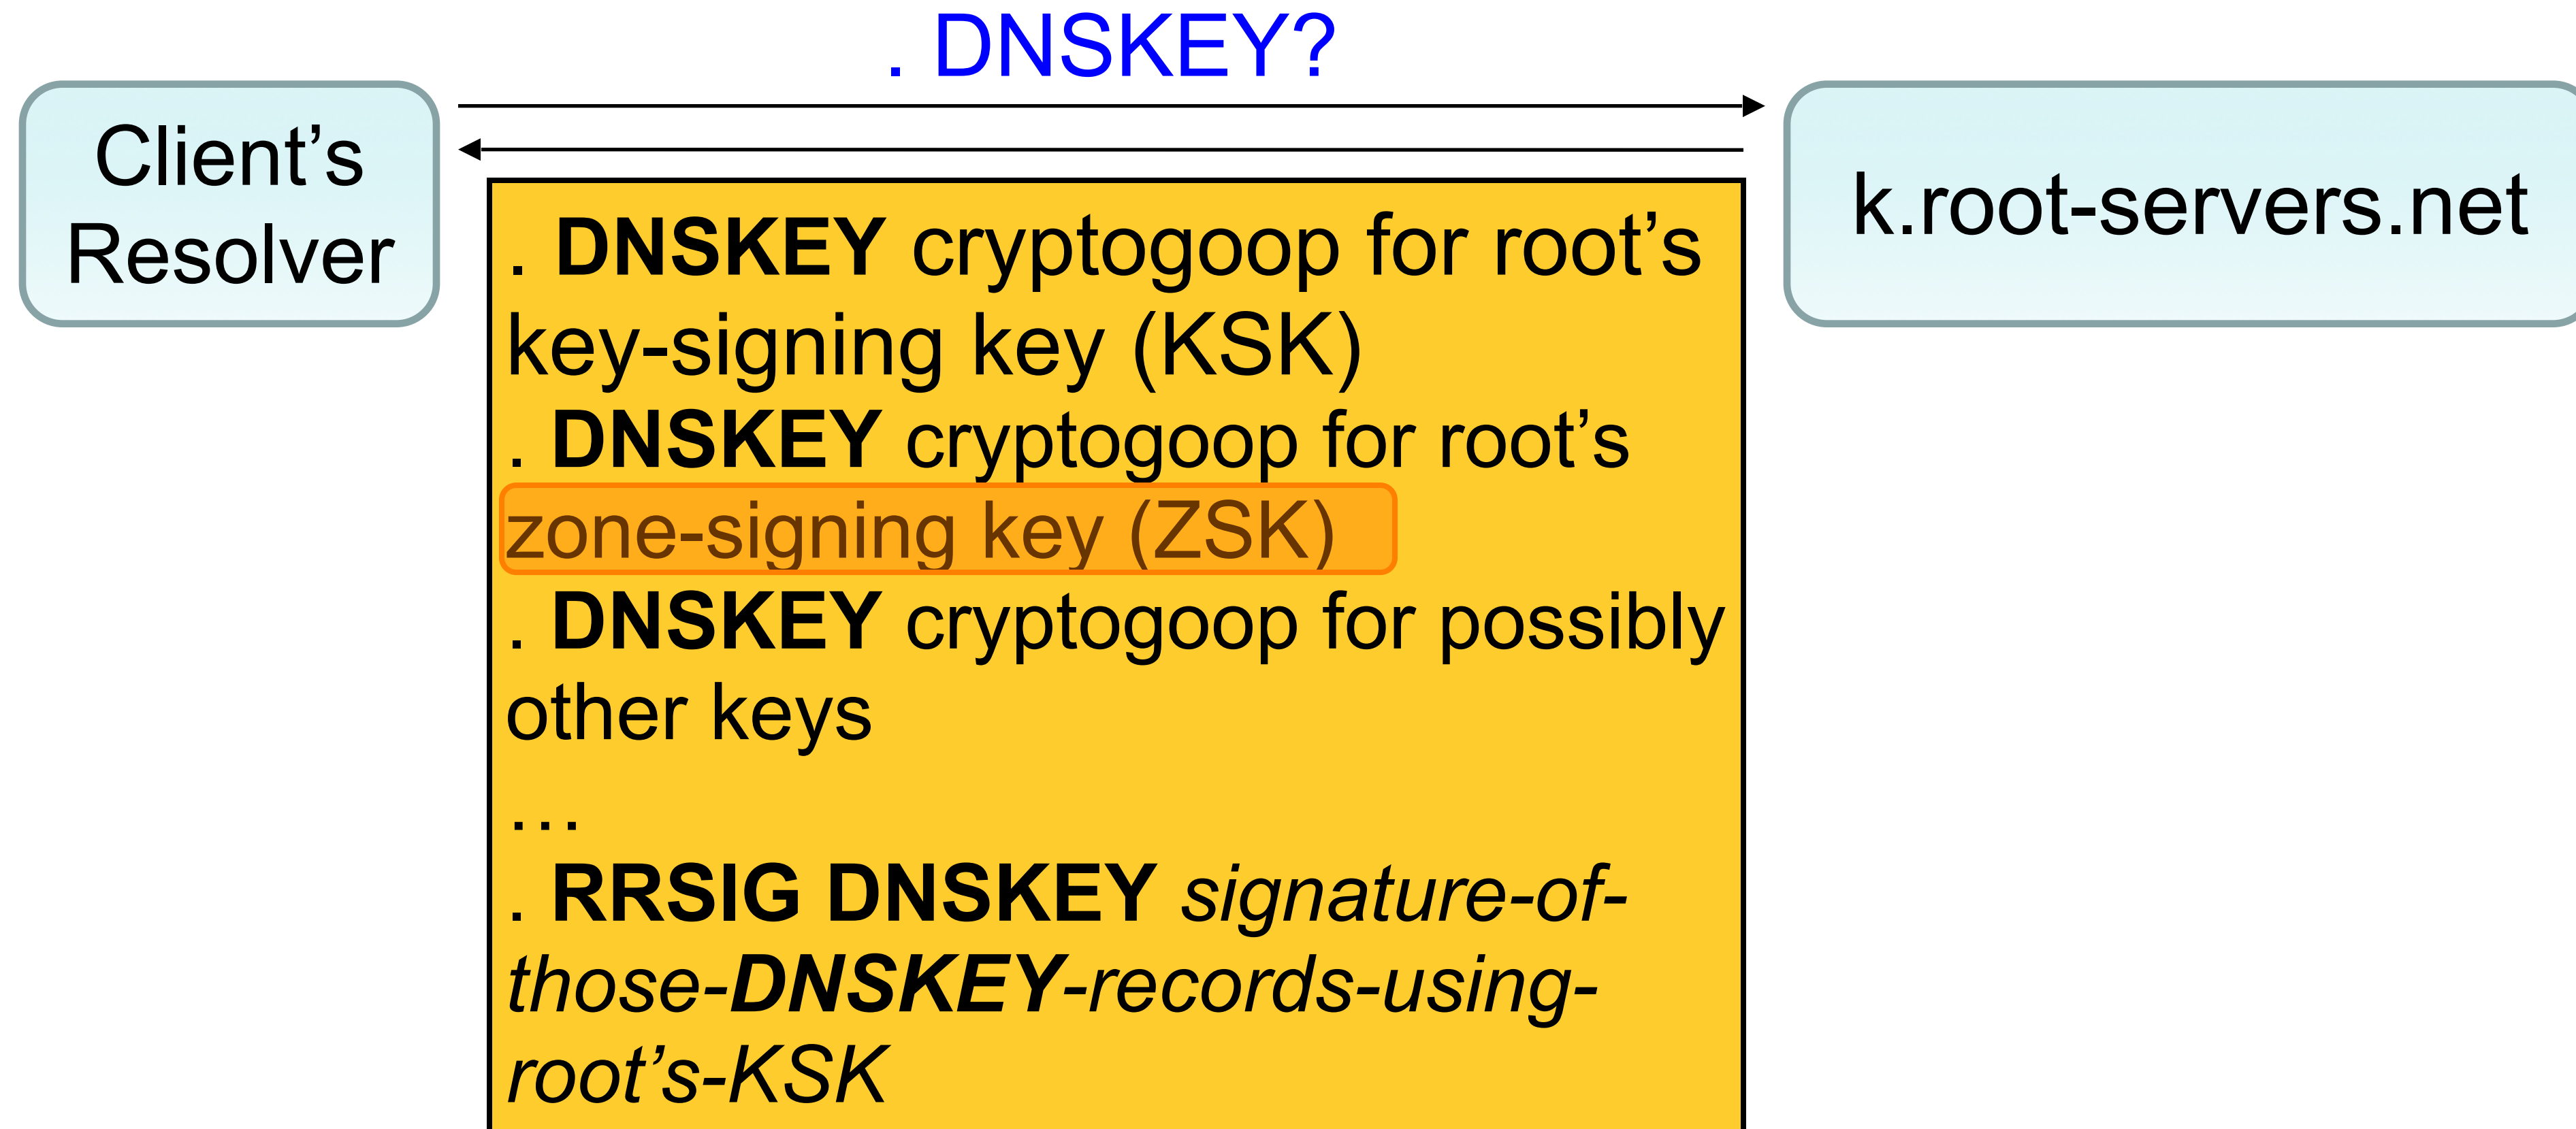
The client has a hash of the root's **KSK** hardwired into its config as a trust anchor.

DNSSEC: Accessing keys



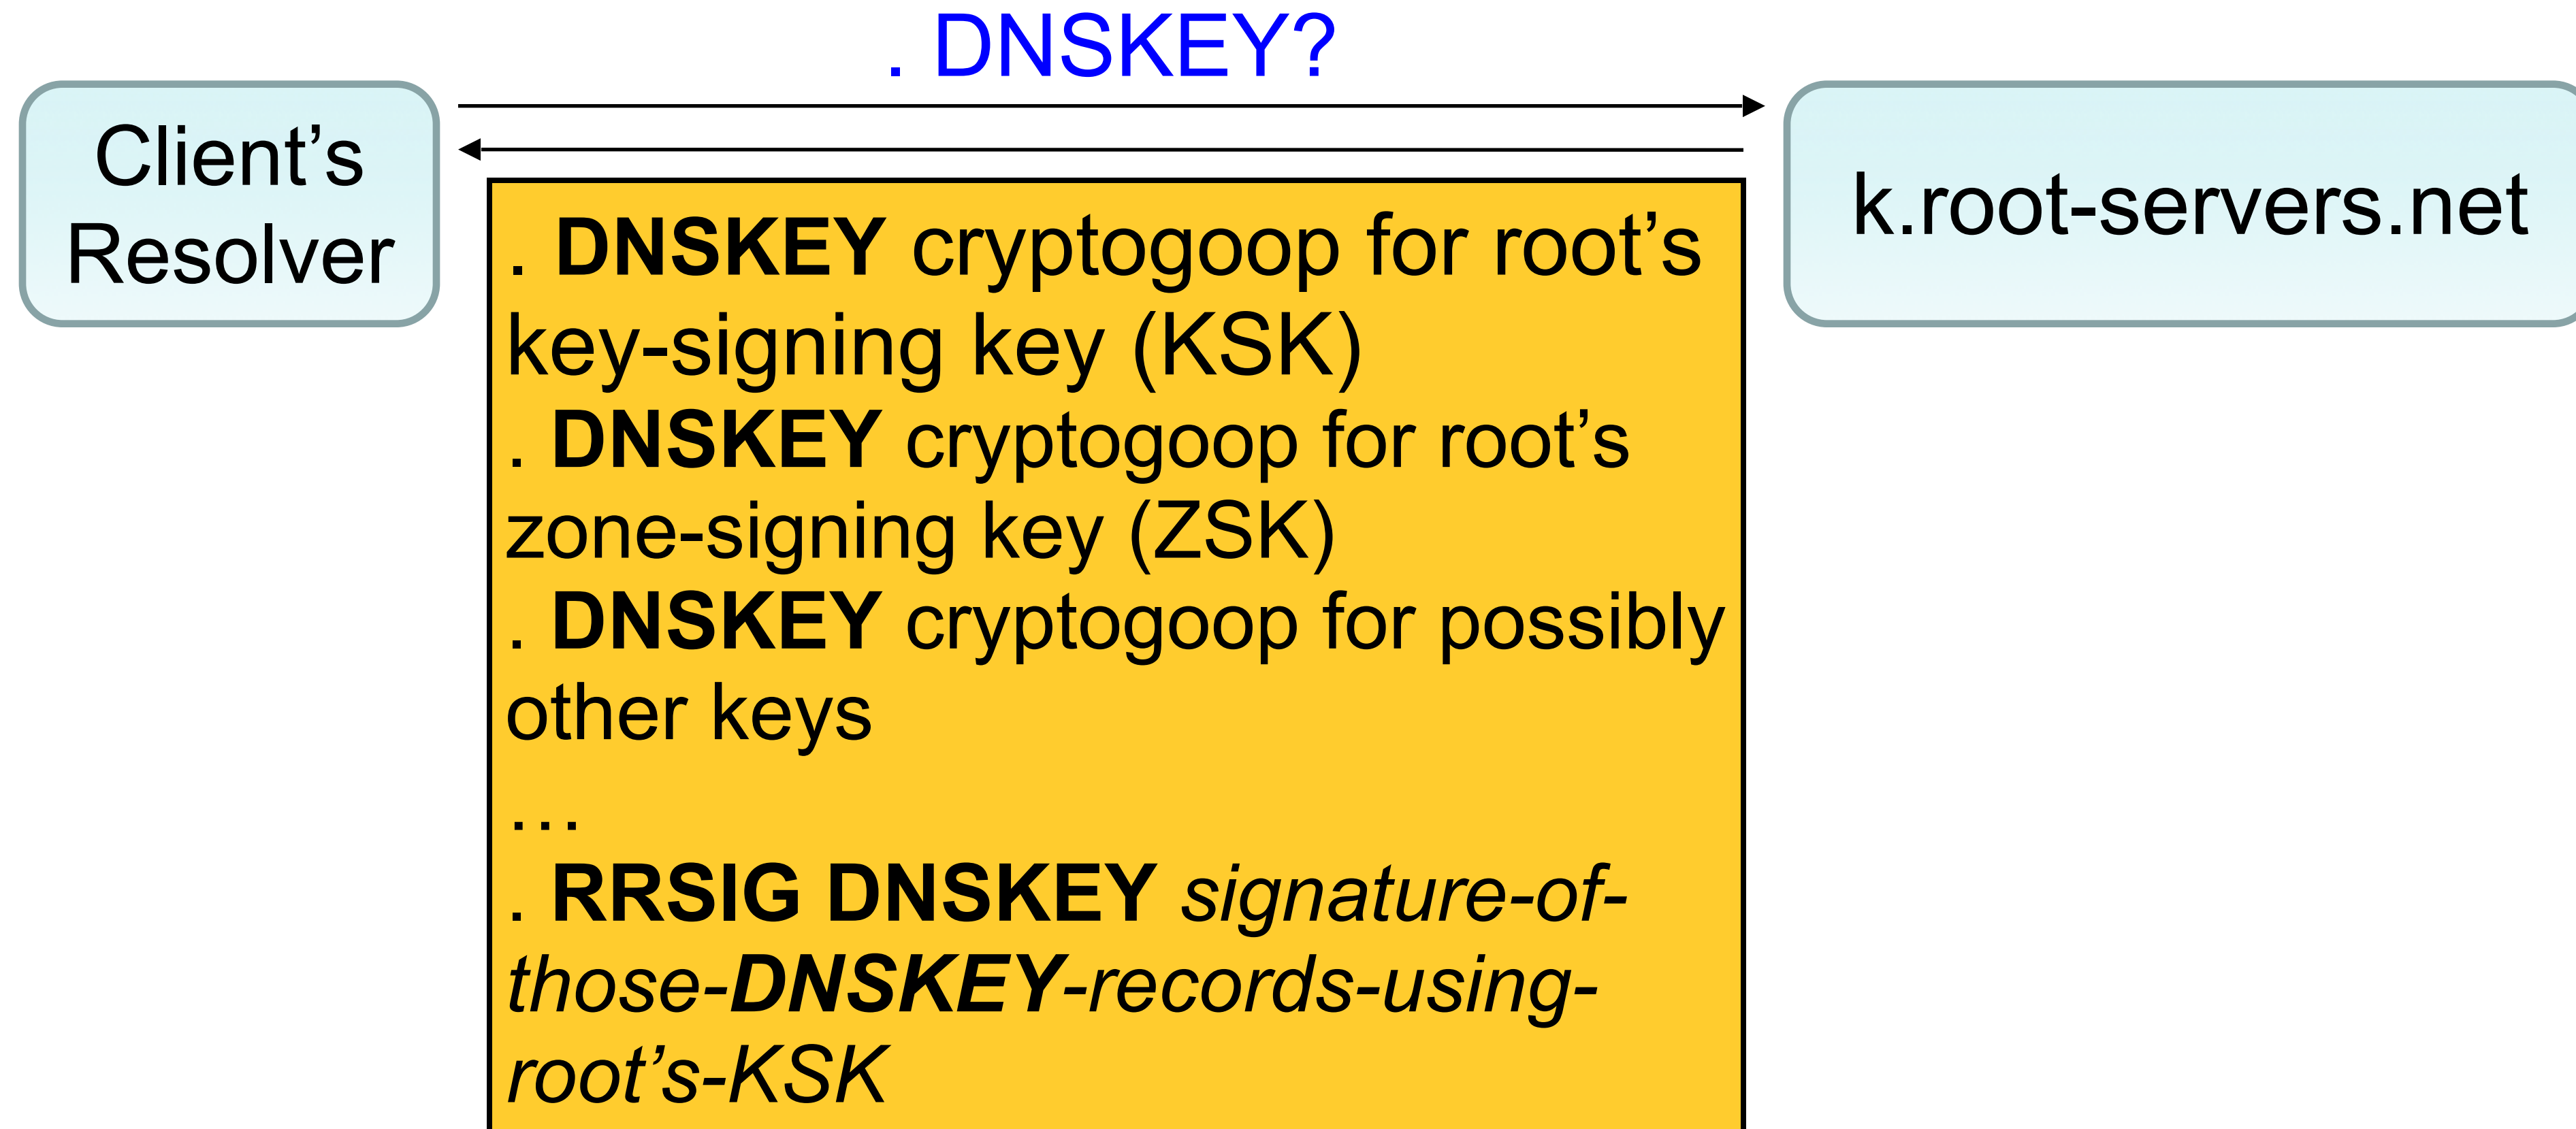
For everything below the root (e.g., .com and google.com) we get a hash of the KSK via a **DS** record, as shown earlier, so we can tell if we get the right KSK in a **DNSKEY** entry.

DNSSEC: Accessing keys




The **ZSK** is used for signing all of the other RRSIG entries in the zone, including DS records for subzones.
(E.g., .com signs its **DS** record for google.com using .com's **ZSK**)

DNSSEC: Accessing keys



Having separate key-signing-keys vs. zone-signing-keys allows a zone to change its **ZSK** without needing to get its parent to re-sign, since parent only signs the **KSK**. Enables frequent *key rollover*.

End-to-end look



Issues With DNSSEC ?

- Issue #1: Replies are Big
 - DoS **amplification**
 - Increased **latency** on low-capacity links
 - Headaches w/ older libraries that assume replies < 512B

Issues With DNSSEC ?

- Issue #1: Replies are Big
 - DoS amplification
 - Increased latency on low-capacity links
 - Headaches w/ older libraries that assume replies < 512B
- Issue #2: *Partial deployment*
 - What do you do with unsigned/unvalidated results?
 - If you trust them, **weakens incentive** to upgrade
 - If you don't trust them, a whole lot of things **break**

Issues With DNSSEC, con't

- Issue #3: *Management headaches*
 - What happens if when updating your site's keys you make a mistake?
 - Suddenly your **Entire Site Breaks**
- Issue #4: Negative results ("no such name")
 - What statement does the nameserver sign?
 - If "gab1uph.google.com" doesn't exist, then have to do dynamic key-signing (expensive) for any bogus request
 - **DoS vulnerability**
 - Instead, sign (off-line) statements about order of names
 - E.g., sign "gabby.google.com followed by gabrunk.google.com"
 - Thus, can see that gab1uph.google.com can't exist
 - But: now attacker can **enumerate** all names that exist :-)

Issues With DNSSEC, con't

- Issue #5: *Who do you really trust?*
 - For your laptop (say), who does all the “grunt work” of fetching keys & validating DNSSEC signatures?
- **Convenient** answer: your laptop's local resolver
 - ... which you acquire via DHCP in your local coffeeshop
 - I.e., exactly the most-feared potentially **untrustworthy** part of the DNS resolution process!
- Alternatives?
 - ⇒ Your laptop needs to do all the validation work itself :-)

Summary of DNSSEC

- DNSSEC: provides **object security** for DNS results
 - Just **integrity** & **authentication**, not confidentiality
 - No client/server setup “dialog”
 - Tailored to be **caching-friendly**
 - Underlying security dependent on trust in Root Name Server’s key ...
 - ... plus support provided by every level of DNS hierarchy from Root to final name server... **and local resolver!**